

Chapter 7


Memory Management



Roadmap

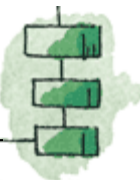
- Basic requirements of Memory Management
- Memory Partitioning
- Basic blocks of memory management
 - Paging
 - Segmentation





Uni-Programming

- In a uni-programming system, main memory is divided into two parts:
 - One part for the **operating system** (resident monitor, kernel) and
 - One part for the **program** currently being executed.






Multi-Programming

- In a multiprogramming system, the “user” part of memory must be further subdivided to accommodate multiple processes.
- *“Memory management is vital in a multiprogramming system.”*
 - If only a few processes are in memory, then for much of the time all of the processes will be waiting for I/O and the processor will be idle.
 - Thus memory needs to be allocated to ensure a reasonable supply of ready processes to consume available processor time.






The need for memory management

- Memory is cheap today, and getting cheaper
 - But applications are **demanding more** and more memory, there is never enough!
- Memory Management, involves **swapping** blocks of data from secondary storage.
- Memory I/O is **slow** compared to a CPU
 - The OS must cleverly time the swapping to maximise the CPU's efficiency






Memory Management

“Memory needs to be allocated to ensure a reasonable supply of ready processes to consume available processor time”






Memory Management Requirements

- Relocation
- Protection
- Sharing
- Logical organisation
- Physical organisation





Requirements: Relocation

- Main memory is shared by multiple processes
- The programmer does not know **where the program will be placed** in memory,
 - it may be swapped to disk and return to main memory at a different location (**relocated**)
- Memory references must be **translated** to the actual physical memory address



Challenge for Relocation: Addressing

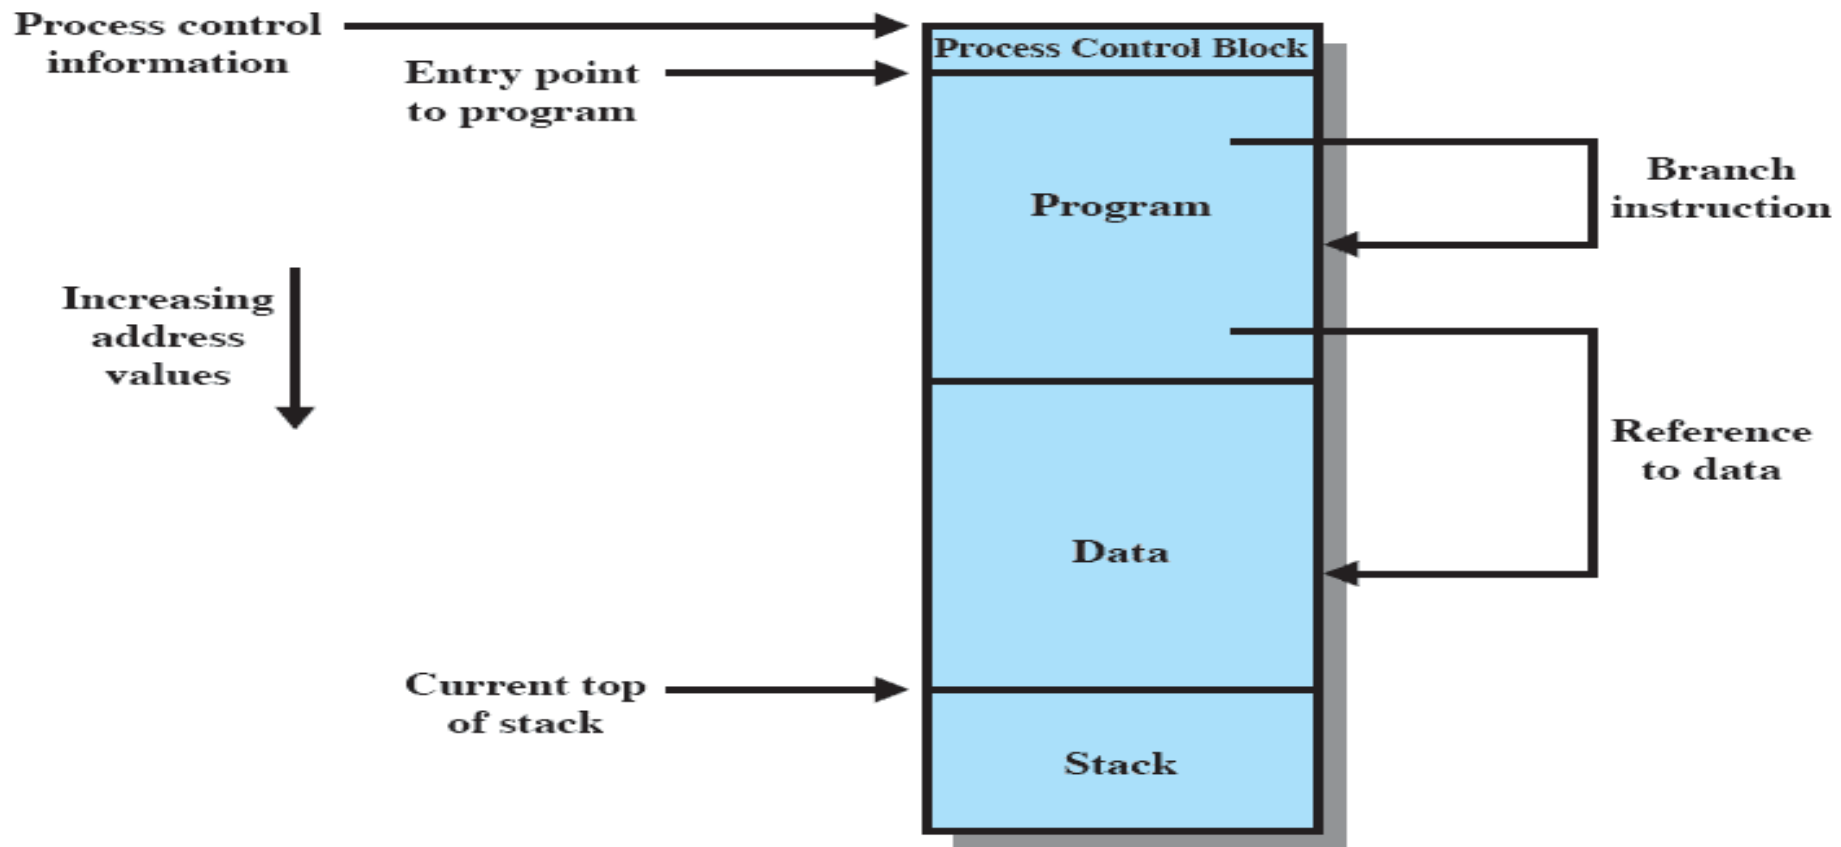




Figure 7.1 Addressing Requirements for a Process



Challenge for Relocation: Addressing

- OS needs to know addresses of
 - PCB
 - Program Area
 - Stack Area
- OS loads process in main memory at specific location so these addresses are known to it
- Mapping between actual program references and physical memory locations needs to be done
- This is done by processor hardware and OS together






Requirements: Protection

- Normally, a user process cannot access any portion of the operating system, neither program nor data. – Why?
- Usually a program in one process cannot branch to an instruction in another process or access the data area of another process. – Why?






Requirements: Protection

- “Satisfaction of Relocation requirement increases the difficulty of satisfying protection”
 - Location of a process in main memory is not fixed
 - Run time checking of memory references is needed





Requirements: Protection

- Memory protection must be satisfied by processor rather than OS
 - OS cannot anticipate the memory references that a program will make
 - It will be a time consuming task if assigned to OS






Requirements: Sharing

- Protection mechanism must have **flexibility** to allow several processes to access the same portion of memory
- Better to allow each process access to the **same copy** of the program rather than have their own separate copy
- Processes that are cooperating on some task may need to share access to the **same data structure**.





Requirements: Sharing

- Hence memory management system must allow **controlled access** to shared area without compromising the protection requirement.





Requirements: Logical Organization

- Memory is organized linearly (usually)
- Programs are written in modules
 - Modules can be written and compiled independently
 - Different degrees of protection given to modules (read-only, execute-only)
 - Share modules among processes
- Hence this mapping needs to be managed





Requirements: Physical Organization

- **Memory** is organized in **two levels**: main memory and secondary memory
- **Main memory** has the advantage of faster access but it is volatile
- **Secondary memory** has the advantage of bulk and permanent storage but it is slow in access
- The **flow** between main and secondary memory needs proper management





Requirements: Physical Organization

- Cannot leave the programmer with the responsibility to manage memory
- Memory available for a program plus its data may be insufficient
 - **Overlaying** allows various modules to be assigned the same region of memory but is time consuming to program
- Programmer does not know how much space will be available

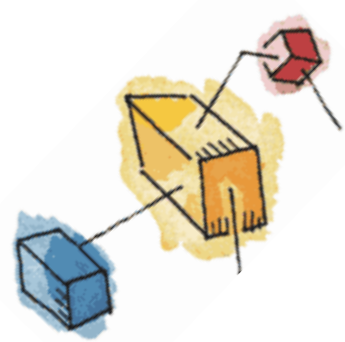




Partitioning

- **Basic operation** of memory management unit is
 - Bring a process into main memory
- This requires the memory to be properly partitioned
 - Hence **partitioning** is one of the important memory management techniques





Types of Partitioning

- Fixed Partitioning
- Dynamic Partitioning
- Simple Paging
- Simple Segmentation
- Virtual Memory Paging
- Virtual Memory Segmentation





Fixed Partitioning

- If we assume that **OS** occupies some **fixed portion** of main memory
 - **Rest** of the space is available for **processes**
- “Partition the available memory into regions with fixed boundaries”
- *Fixed Partitioning*
 - *Equal size*
 - *Unequal size*





Fixed Partitioning

- ***Equal-size partitions***
 - Any process whose size is less than or equal to the partition size can be loaded into an available partition
- The operating system can **swap** a process out of a partition
 - If none are in a ready or running state



(a) Equal-size partitions






Fixed Partitioning Problems


- A program may not fit in a partition.
 - The programmer must design the program with overlays
- Main memory use is inefficient.
 - Any program, no matter how small, occupies an entire partition.
 - This results in *internal fragmentation*.
 - ***Solution??***





Solution – Unequal Size Partitions

- Lessens both problems
 - but doesn't solve completely
- In Figure,
 - Programs up to 16M can be accommodated without overlay
 - Smaller programs can be placed in smaller partitions, reducing internal fragmentation



(b) Unequal-size partitions





Placement Algorithm

- **Equal-size**

- Placement is trivial (no options)
- As long as a **partition is available**, process can be loaded into it
- Does not matter which partition to choose
- If all partitions are occupied and no ready or running processes, **swapping** can be done

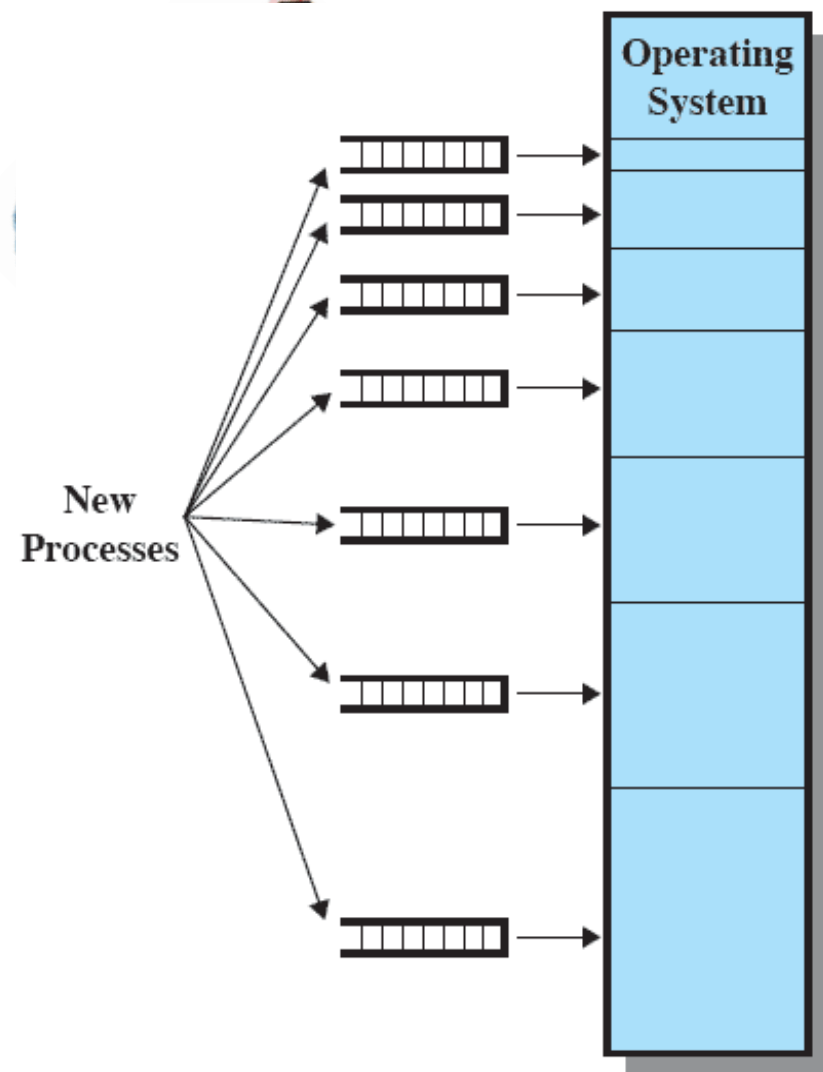





Placement Algorithm

- Unequal-size
 - Can assign each process to the **smallest partition** within which it will fit
 - Processes are assigned in such a way as to minimize wasted memory within a partition
 - *Queue models:*
 - One queue per partition
 - Single queue



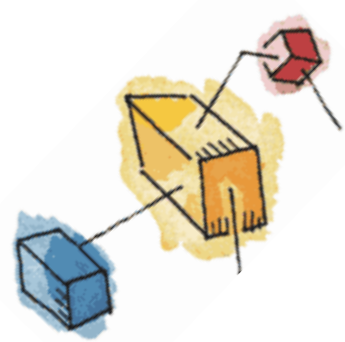


(a) One process queue per partition



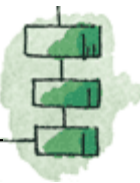
(b) Single queue


Figure 7.3 Memory Assignment for Fixed Partitioning



Placement Algorithm

- Problem with one queue per partition is
 - It might be optimum from a partition point of view, but not optimum from system's point of view
 - If there are no processes which can fit a particular partition, then it will remain unused
 - Hence, using a single queue is a better alternative





Advantages: Fixed Partitioning

- Fixed partitioning schemes are **relatively simple** and require **minimum** operating system software and processing **overhead**
- Unequal size partitioning provides **flexibility** to the basic approach





Remaining Problems with Fixed Partitions

- The number of active processes is limited by the system
 - i.e. limited by the pre-determined number of partitions
- A large number of very small processes will not use the space efficiently
 - In either fixed or variable length partition methods







Dynamic Partitioning

- Partitions are of **variable length** and **number**
- Process is allocated exactly as much memory as required

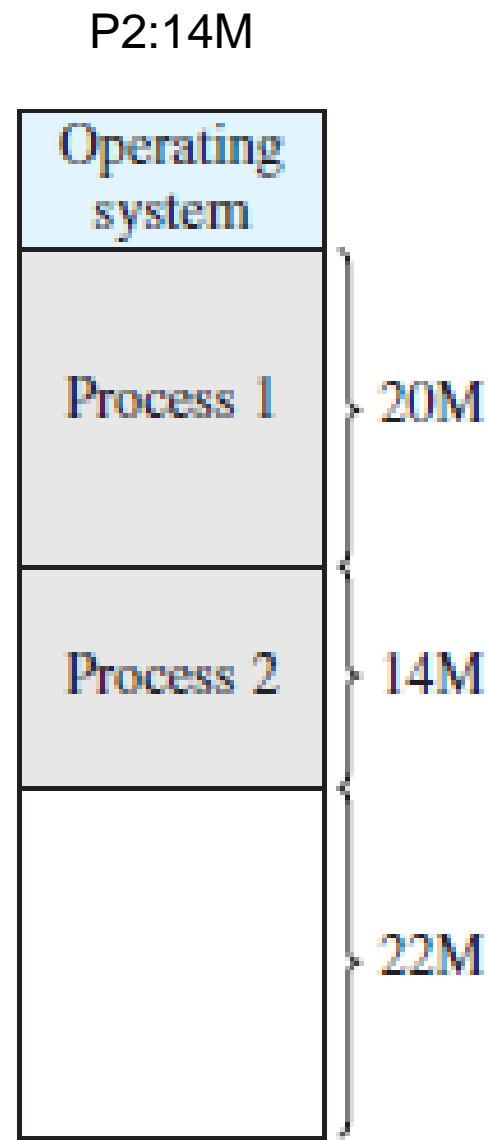




(a)



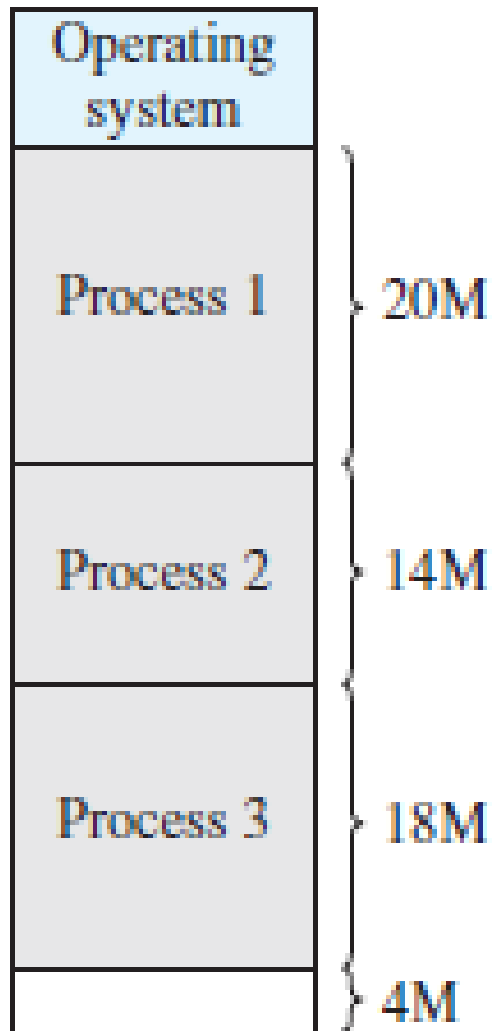
(b)



(c)





P3:18M



(d)


None Ready, new process arriving, Swap P2






None Ready,
P2 Ready-Suspend,
Swap P1

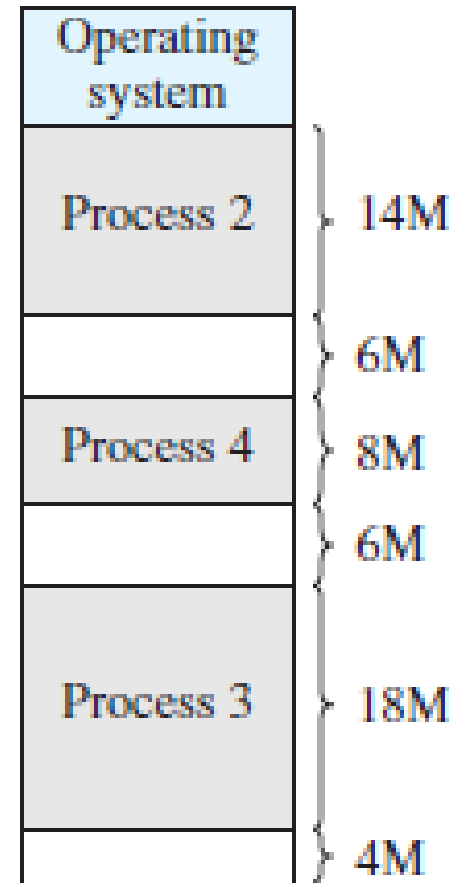
P2 swapped in



(f)




(g)



(h)





Dynamic Partitioning: Issue

- **External Fragmentation**
 - Memory external to all processes is fragmented
- Can resolve using **compaction**
 - OS moves processes so that they are contiguous
 - Needs dynamic relocation
 - Time consuming and wastes CPU time





Placement Algorithms

- Compaction is costly
- Hence operating system must make **intelligent decision** when a process is brought in and multiple options are available
- Three placement algorithms:
 - Best Fit
 - Next Fit
 - First Fit





Placement Algorithms

- **Best-fit algorithm**

- Chooses the block that is closest in size to the request
- **Worst performer** overall
 - Since smallest block is found for process, the smallest amount of fragmentation is left
 - Memory compaction must be done more often





Placement Algorithms

- *First-fit algorithm*

- Scans memory from the beginning and chooses the first available block that is large enough
- Fastest
- May have many process loaded in the front end of memory that must be searched over when trying to find a free block
- Small holes in front






Placement Algorithms

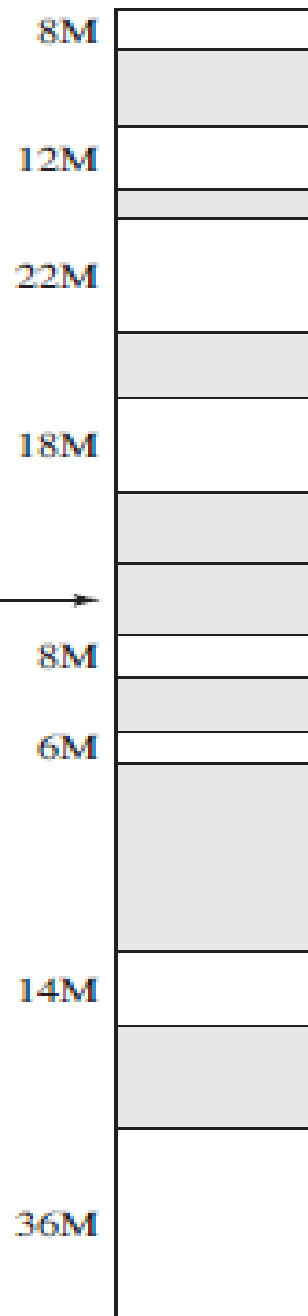
- *Next-fit algorithm*

- Scans memory from the location of the last placement
- More often allocate a block of memory at the **end of memory** where the largest block is found
- The largest block of memory is broken up into smaller blocks
- Compaction is required to obtain a large block at the end of memory





Last
allocated
block (14K)




Allocated block



Free block



(a) Before



16M Allocation

First fit

Best fit

Next fit

8M

12M

6M

2M

8M

6M

14M

20M



Allocated block





Free block



Possible new allocation

(b) After





Replacement Algorithms

- In dynamic partitioning, when **all processes** in main memory are **blocked** and there is **insufficient memory**, even after compaction
 - OS must **swap** one of the process to bring a new process or a ready suspended process
- This needs **replacement algorithm**





Common Drawbacks

- Fixed Partitioning
 - Internal Fragmentation
 - Limits the number of active processes
- Dynamic Partitioning
 - External Fragmentation
 - Compaction increases overhead
- Another Approach
 - Buddy System






Buddy System

- Entire space available is treated as a single block of 2^U
- If a request of size s is made satisfying $2^{U-1} < s \leq 2^U$
 - entire block is allocated
- Otherwise block is split into two equal buddies
- Process continues until smallest block greater than or equal to s is generated





Example of Buddy System

- Consider block of 512 bytes
 - Request $S=512$, $256 < 512 \leq 512$
 - Condition satisfied, Allocate entire block
 - Request $S=400$, $256 < 400 \leq 512$
 - Condition satisfied, Allocate entire block
 - Request $S=200$, Condition not satisfied
 - so split in two equal buddies (256,256)
 - Check condition again, $128 < 200 \leq 256$
 - Allocate 256 byte block





Available Space

1M



Request A:100K

A = 128K

128K

256K

512K

Request B:240K

A = 128K

128K

B = 256K

512K

Request C:64K

A = 128K

C = 64K

64K

B = 256K

512K

Request D:256K

A = 128K

C = 64K

64K

B = 256K

D = 256K

256K





Release B



A = 128K	C = 64K	64K	256K	D = 256K	256K
----------	---------	-----	------	----------	------

Release A

128K	C = 64K	64K	256K	D = 256K	256K
------	---------	-----	------	----------	------

Request E:75K

E = 128K	C = 64K	64K	256K	D = 256K	256K
----------	---------	-----	------	----------	------

Release C

E = 128K	128K	256K	D = 256K	256K
----------	------	------	----------	------

Release E

512K	D = 256K	256K
------	----------	------

Release D

1M



Tree Representation of Buddy System






Figure 7.7 Tree Representation of Buddy System




Tree Representation of Buddy System

- If two buddies are leaf nodes, then one of them must be an allocated node, otherwise they should be merged
- At any point of time, the leaf nodes define current allocation



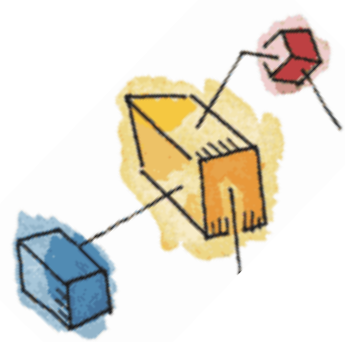
Relocation

- When program loaded into memory the actual (absolute) memory locations are determined
- A process may occupy different partitions which means different absolute memory locations during the lifetime
 - Reasons?
 - Swapping
 - Compaction



Relocation

- The schemes where relocation is needed:
 - Equal size fixed partitioning
 - Unequal size fixed partitioning (single queue)
 - Dynamic partitioning
 - Compaction
- Hence physical memory references of a process are not fixed
 - This issue is the focus of addressing





Types of Addresses

- Physical or Absolute
 - The absolute address or actual location in main memory.
- Logical
 - Reference to a memory location independent of the current assignment, generated by CPU
- Relative
 - Logical address, where address expressed as a location relative to some known point (i.e. beginning).





Address Translation

- For system with relative addressing, address translation is done in two steps:
 - Load
 - Load process in main memory with all references in relative form
 - Calculate
 - Calculate physical address when instructions or data with relative address is encountered





Registers Used during Execution


- Base register
 - Starting address for the process
- Bounds register
 - Ending location of the process
- These values are set when the process is loaded or when the process is swapped in





Address Translation

- When a process is assigned running state, **base register** is loaded with **starting physical address** of the process
- **Bound register** is loaded with **ending physical address**
- When **relative address** is encountered
 - The value of the base register is added to a relative address to produce an absolute address





Address Translation

- The **resulting address** is **compared** with the value in the **bounds register**
- If the address **is not within bounds**, an **interrupt** is generated to the operating system
- This mechanism supports protection



Hardware Support

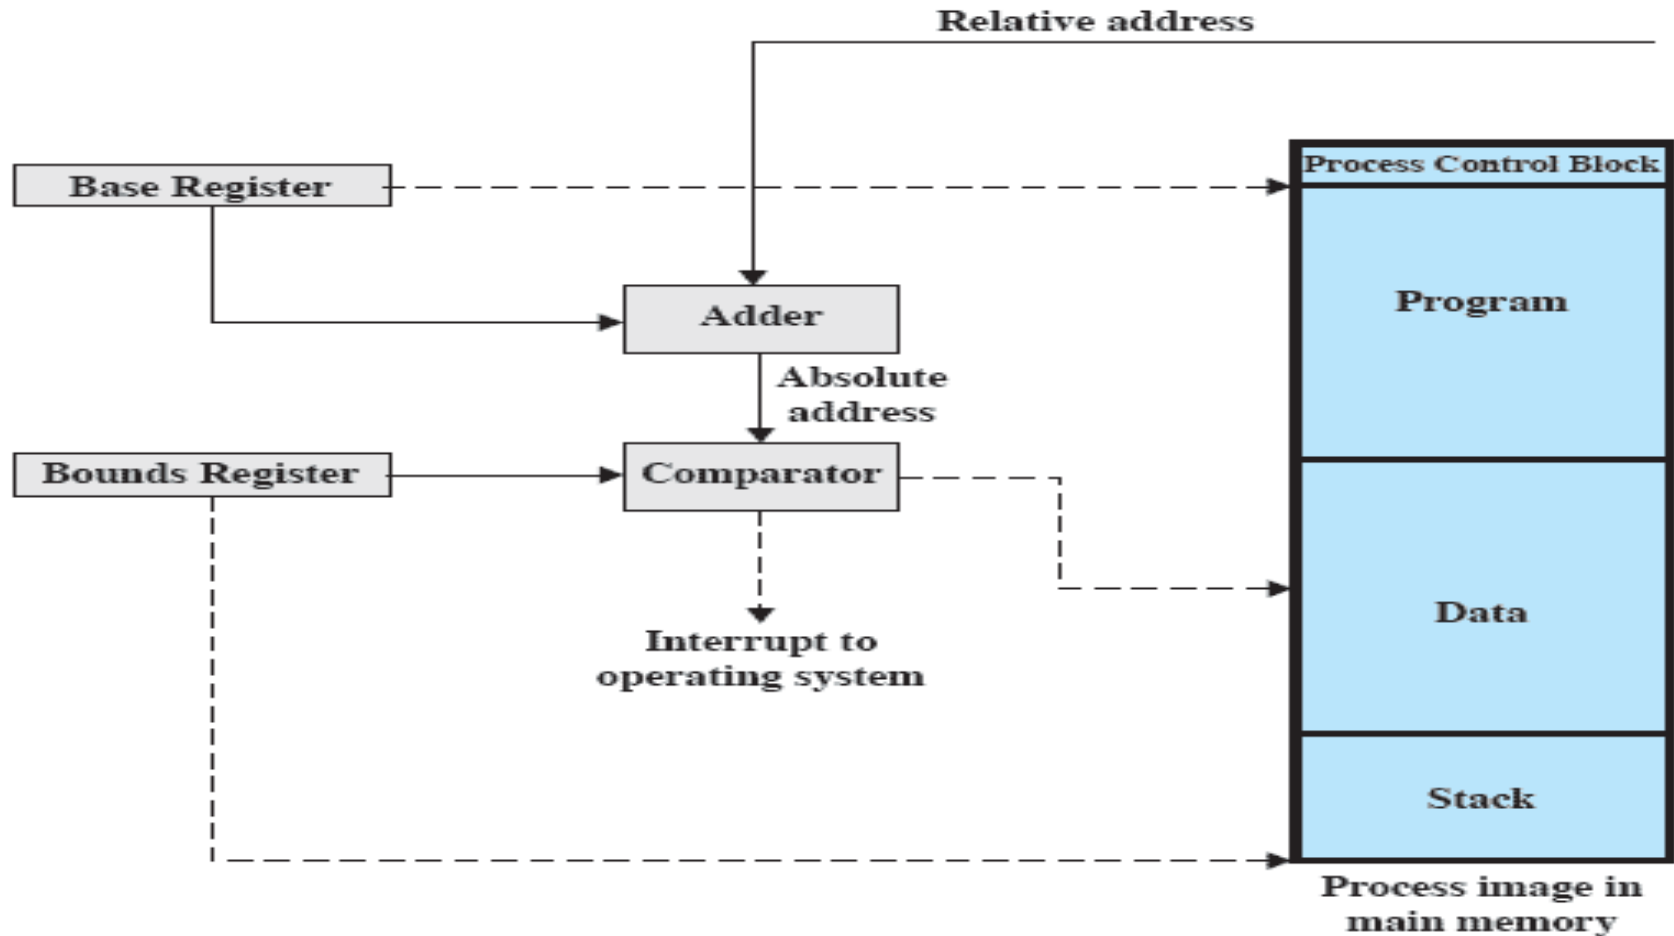


Figure 7.8 Hardware Support for Relocation



Paging

- In paging scheme,
 - Main memory is partitioned into small, fixed-size chunks
 - Each process is divided into the same size chunks
- The chunks of memory are called *frames*
- The chunks of a process are called *pages*





Paging

- When a process is to be loaded in main memory,
 - Free frames are found
 - Pages are loaded into those frames
- *List of free frames is maintained by OS*
- If contiguous frames are not available,
 - Then also new process can be loaded
 - This requires a data structure to keep track of current allocation





Paging- Page Table

- Operating system maintains a **page table** for this purpose
 - *1 page table per process*
 - Page Table Entry contains the frame location where the page is stored



Paging Example



Frame number	Main memory
0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(a) Fifteen available frames

A:4 Pages

	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(b) Load process A

B:3 Pages

	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	
8	
9	
10	
11	
12	
13	
14	

(c) Load process B



Paging Example

C:4 Pages

All blocked,
bring new D:5 pages

D:5 Pages

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(d) Load process C

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(e) Swap out B

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

(f) Load process D



Page Table Entries

0	0
1	1
2	2
3	3

Process A
page table

0	—
1	—
2	—

Process B
page table

0	7
1	8
2	9
3	10

Process C
page table

0	4
1	5
2	6
3	11
4	12


Process D
page table

13
14

Free frame
list

Figure 7.10 Data Structures for the Example of Figure 7.9 at Time Epoch (f)






Paging: Address Structure

- Each **logical address** in program is a **pair**
 - Page Number
 - Offset
 - E.g. Page 1 offset 5
- **Address Translation**
 - *Page number, offset*
 - To
 - *Frame number, offset*
- Done by **processor hardware**





Paging vs. Fixed Partitioning

- Paging has **small partition size** compared to fixed partitioning
- Paging does **not require partitions** to be **contiguous**
- Paging results into **small amount** of **internal fragmentation** (last page)





Paging: Example

- Consider the memory supporting 16 bit addresses,
 - address space is 64K
- Given page size 1024 bytes (1 K)
 - Number of pages?
 - $2^{16} / 2^{10} = 2^6$
 - 64 pages of 1K each
 - Address Structure?
 - 10 bits offset, 6 bits page number = 16 bit address





Paging: Example

- Logical Address = 31 bits
 - Address space = $2^1 * 2^{30}$ (2 G)
- Address space = 128M ($2^7 * 2^{20}$)
 - Logical Address = 27 bits
- Physical Address = 22 bits
 - Address space = $2^2 * 2^{20}$ (4 M)
- Address space = 16M ($2^4 * 2^{20}$)
 - Physical Address = 24 bits





Paging: Addresses

- Given physical address = 12 bits
- Logical address = 13 bits
- Page size = Frame size = 1 K
- Number of Pages?
- Number of Frames?
- Structure of Logical Address?
- Structure of Physical Address?






Paging: Addresses

- If page size and frame size are power of 2, then relative address and logical address are the same

Relative address = 1502


0000010111011110




(a) Partitioning

Logical address =
Page# = 1, Offset = 478

0000010111011110



(b) Paging
(page size = 1K)

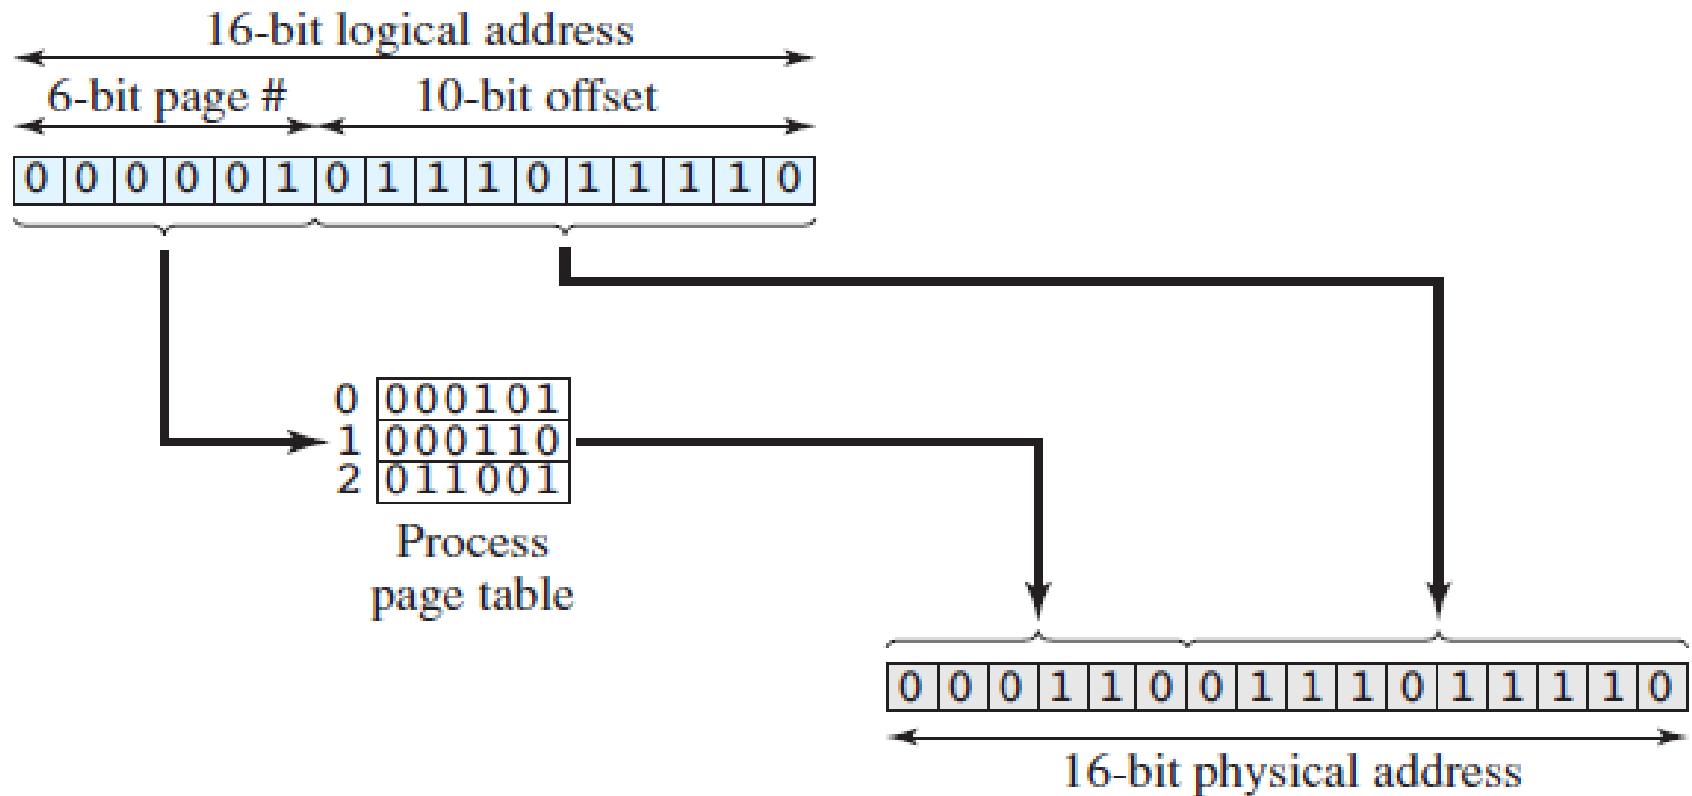


Paging: Address Translation

- Consider an address of $n+m$ bits, where leftmost n bits are the page number and the rightmost m bits are the offset
- Address translation Steps:
 - Extract the page number as the leftmost n bits
 - Use the page number as an index into process page table to find frame number, k
 - Generate physical address by appending frame number with offset



Paging: Address Translation





Paging: Summary

- Simple paging

- Divide main memory into many small equal size frames
- Divide process into frame-size pages
- Smaller processes require fewer pages
- Larger processes require more pages
- When a process is brought in, all of its pages are loaded into available frames and page table is setup





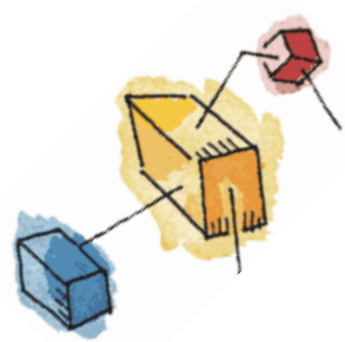
Segmentation


- A user program can be subdivided into segments
 - Segments may **vary in length**
 - There is a **maximum segment length**
- *Similar to dynamic partitioning due to unequal size segments*
- Addressing consist of two parts
 - a segment number and
 - an offset



Segmentation

- One segment table per process
- Segment table entry
 - Starting address of segment in main memory
 - Length of segment
- When a process enters **running state**,
 - Address of segment table is loaded into main memory



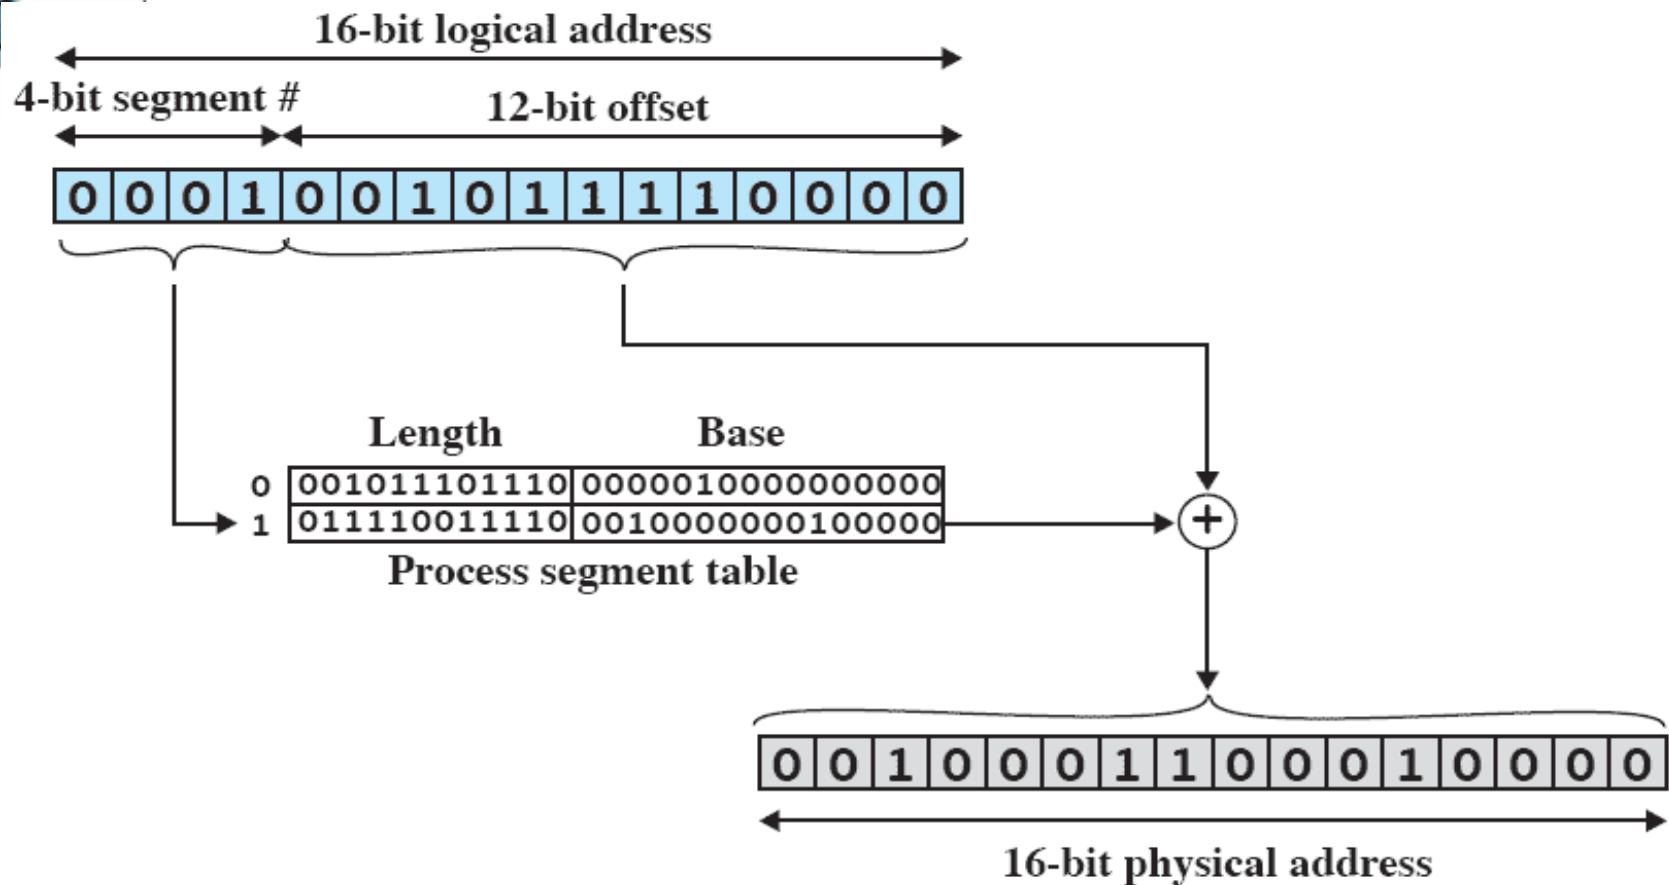


Segmentation: Address Translation

- **Address translation Steps:**
 - Extract the **segment number** as the leftmost n bits
 - Use the segment number as an **index** into process **segment table** to find starting physical address of segment
 - **Compare offset** with the length, if offset greater than length, address is invalid
 - Generate **physical address** by appending starting address with offset



Segmentation



(b) Segmentation

Figure 7.12 Examples of Logical-to-Physical Address Translation

Technique	Description	Strengths	Weaknesses
Fixed Partitioning	Main memory is divided into a number of static partitions at system generation time. A process may be loaded into a partition of equal or greater size.	Simple to implement; little operating system overhead.	Inefficient use of memory due to internal fragmentation; maximum number of active processes is fixed.
Dynamic Partitioning	Partitions are created dynamically, so that each process is loaded into a partition of exactly the same size as that process.	No internal fragmentation; more efficient use of main memory.	Inefficient use of processor due to the need for compaction to counter external fragmentation.
Simple Paging	Main memory is divided into a number of equal-size frames. Each process is divided into a number of equal-size pages of the same length as frames. A process is loaded by loading all of its pages into available, not necessarily contiguous, frames.	No external fragmentation.	A small amount of internal fragmentation.
Simple Segmentation	Each process is divided into a number of segments. A process is loaded by loading all of its segments into dynamic partitions that need not be	No internal fragmentation; improved memory utilization and reduced overhead compared to dynamic partitioning.	External fragmentation.