


*Operating Systems:
Internals and Design Principles, 6/E*
William Stallings

Chapter 4 Threads





Processes and Threads

- Processes have two characteristics:
 - **Resource ownership** - process includes a address space to hold the process image, can be assigned resource ownership
 - **Scheduling/execution** - follows an execution path that may be interleaved with other processes
- These two characteristics are treated independently by the operating system

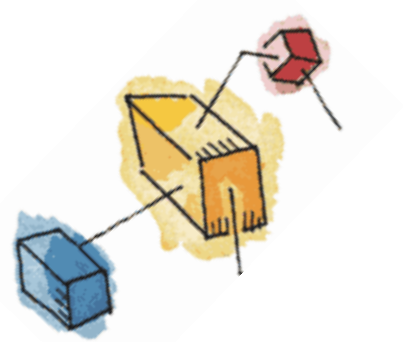




Processes and Threads

- The unit of dispatching is referred to as a **thread** or lightweight process
- The unit of resource ownership is referred to as a process or **task**





Multithreading

- The ability of an OS to support multiple, concurrent paths of execution within a single process.




Figure 4.1 Threads and Processes [ANDE97]



Single Thread Approaches

- **MS-DOS** supports a single user process and a single thread.
- Some **UNIX**, support multiple user processes but only support one thread per process

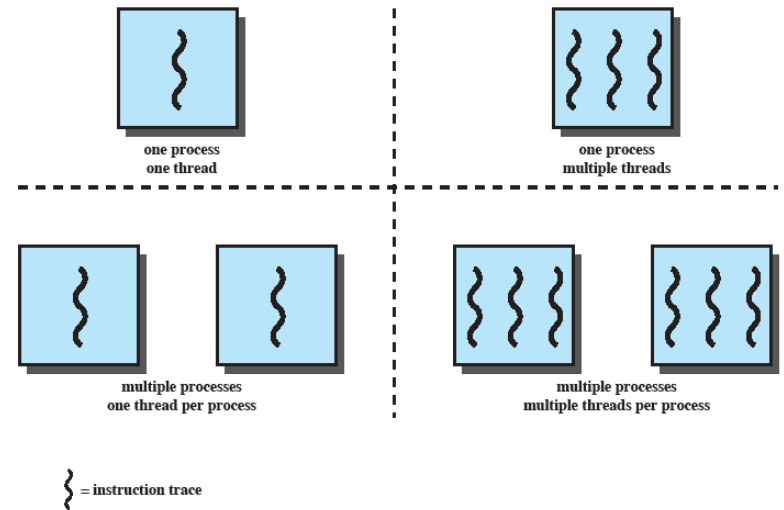


Figure 4.1 Threads and Processes [ANDE97]

Multithreading

- **Java run-time environment** is a single process with multiple threads
- Multiple processes *and* threads are found in **Windows, Solaris**, and many modern versions of **UNIX**


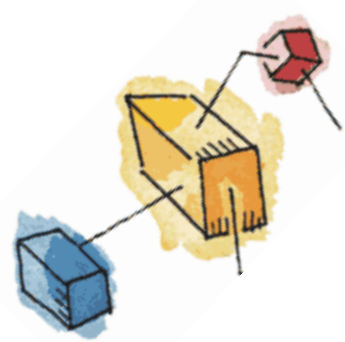


Figure 4.1 Threads and Processes [ANDE97]

Processes

- An **address space** which holds the process image
- Protected access to
 - Processors,
 - Other processes,
 - Files,
 - I/O resources





One or More Threads in Process

- Each thread has
 - An execution state (running, ready, etc.)
 - Saved thread context when not running
 - An execution stack
 - Some per-thread static storage for local variables
 - Access to the memory and resources of its process (all threads of a process share this)





Threads vs. processes

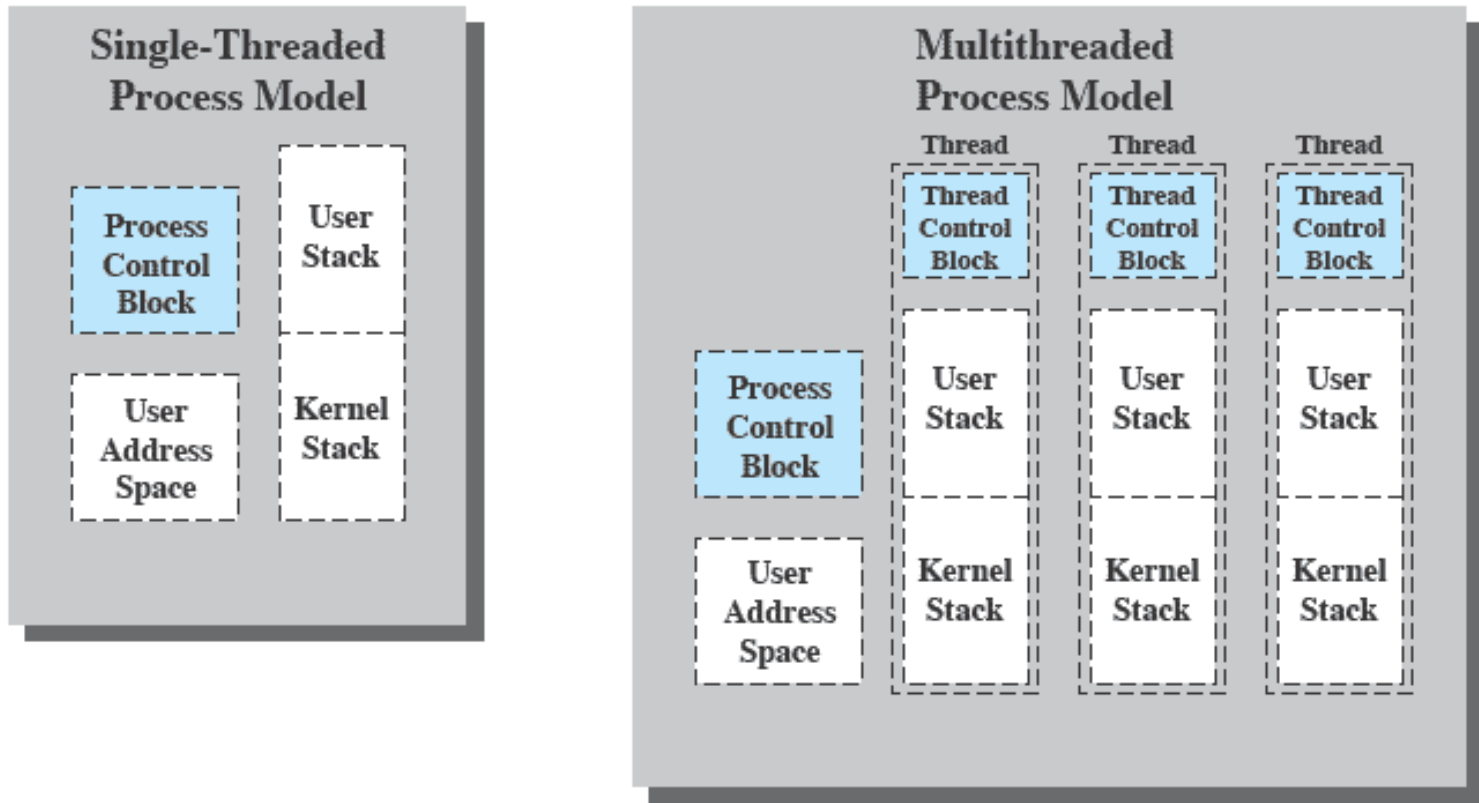


Figure 4.2 Single Threaded and Multithreaded Process Models





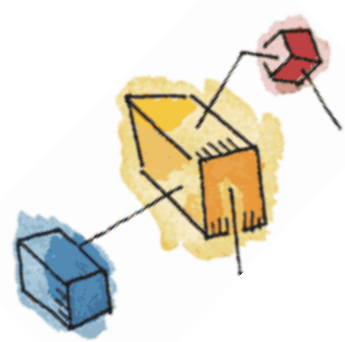
Benefits of Threads

- Takes **less time to create** a new thread than a process
- **Less time to terminate** a thread than a process
- **Switching** between two threads takes **less time** than switching processes
- Threads can **communicate** with each other
 - without invoking the kernel



Thread use in a Single-User System

- **Foreground and background work**
 - One thread for display, other for input etc
- **Asynchronous processing**
 - Periodic backup of main memory to disk
- **Speed of execution**
 - Multiple threads can run concurrently for multiprocessor systems
- **Modular program structure**
 - Easy design




A hand-drawn illustration in the top-left corner shows a large yellow box representing a process. Inside the box, there are several vertical lines representing threads. A blue box is connected to the left side of the yellow box, and a red box is connected to the top right. Lines also connect the threads inside the yellow box to these external boxes.

Threads: Issue

- Several actions that affect all of the threads in a process
 - The OS must manage these at the process level.
- **Examples:**
 - Suspending a process involves suspending all threads of the process
 - Termination of a process, terminates all threads within the process





Thread States and Operations

- **States:**
 - Running
 - Ready
 - Blocked
- **Operations:**
 - Spawn
 - Block
 - Unblock
 - Finish






Example: Remote Procedure Call


- Consider:
 - A program that performs two remote procedure calls (RPCs)
 - to two different hosts
 - to obtain a combined result.





RPC


Using Single Thread






(a) RPC Using Single Thread



RPC Using One Thread per Server



(b) RPC Using One Thread per Server (on a uniprocessor)

-  Blocked, waiting for response to RPC
-  Blocked, waiting for processor, which is in use by Thread B
-  Running

Multithreading on a Uniprocessor

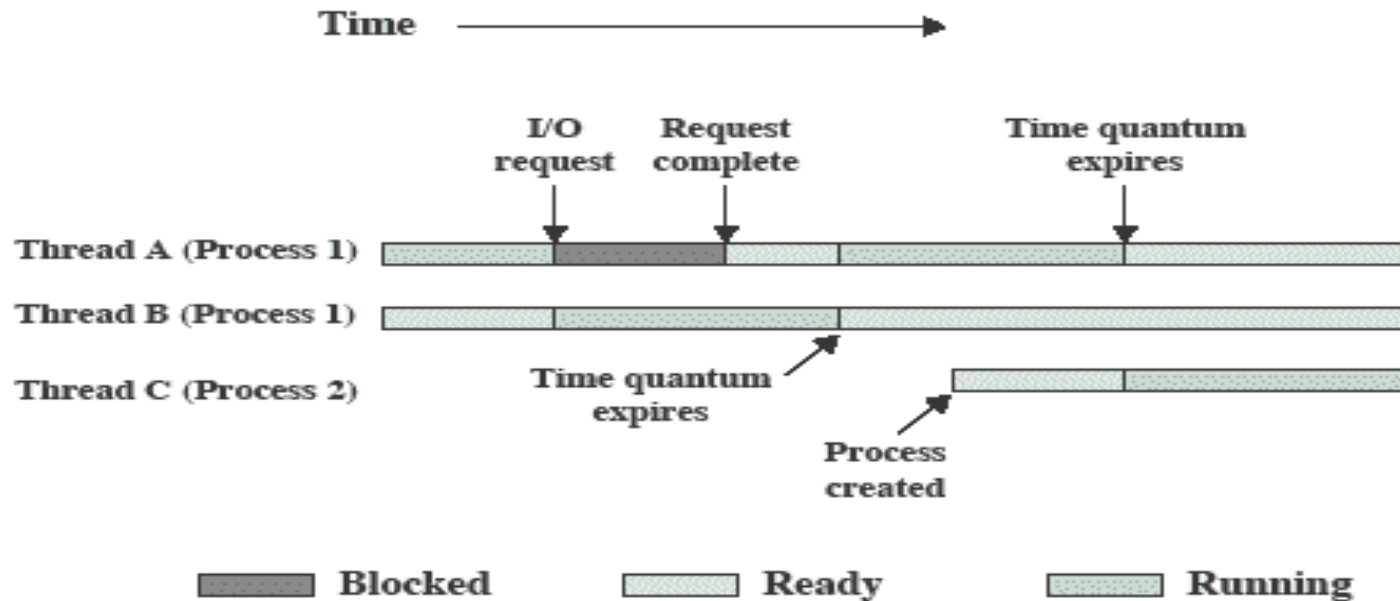


Figure 4.4 Multithreading Example on a Uniprocessor



Categories of Thread Implementation


- User Level Thread (ULT)
- Kernel level Thread (KLT)



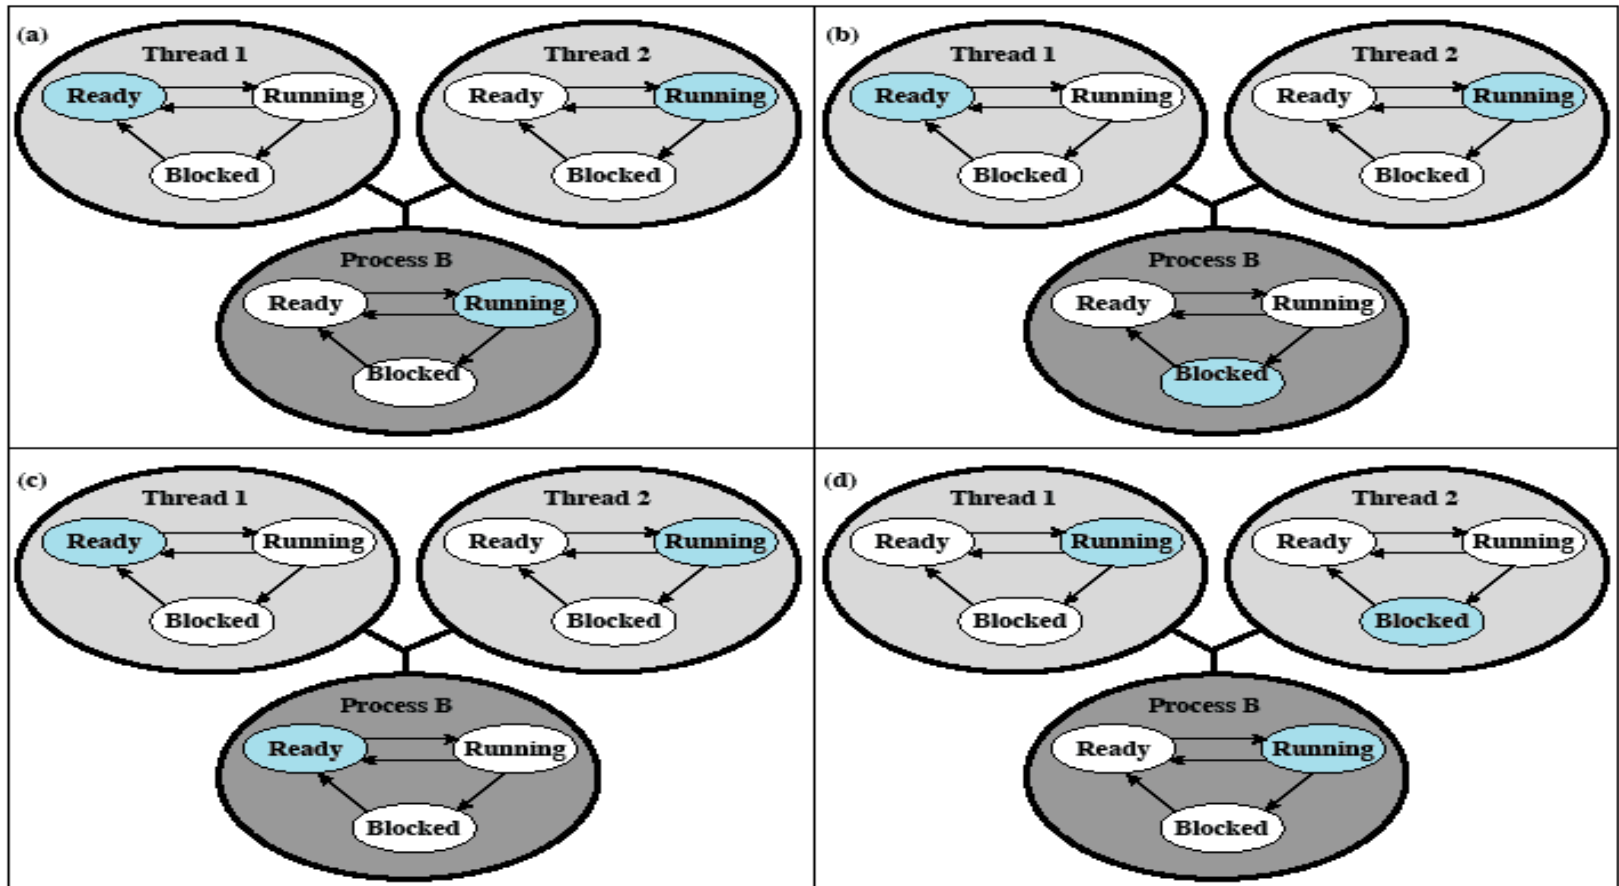


User-Level Threads

- All thread management is done by the application
- The kernel is not aware of the existence of threads
- Application begins with single thread
 - When the process is in running state, new thread can be spawned



Relationships between ULT Thread and Process States



Colored state
is current state


Figure 4.7 Examples of the Relationships Between User-Level Thread States and Process States



Relationships between ULT Thread and Process States

- a) Process B is **running**, using Thread 2
- b) Thread 2's application makes **system call** and blocks B
 - Kernel Switches to other process
 - Thread 2 still remains in **running state** (To maintain data structures) (control needs to return to Thread 2)
- c) Clock Interrupt
 - B is placed in **ready state**, thread state remains same
- d) Thread 2 needs **some action** to be performed by Thread 1





Advantages of ULT

- No need of kernel privileges for thread **switching** (no **mode switch**)
- **Scheduling** can be application specific
- ULT can run on **any OS**





Disadvantages of ULT

- Many system calls are **blocking**
 - Process is ultimately blocked
- Kernel assigns one process to only one processor at a time
 - **A single thread can execute at a time**
- Solution: **Jacketing**
 - Convert blocking system call to non blocking system call
 - Before requesting I/O, check whether it is busy (using jacketing routine)
 - If busy, then block the thread and transfer control to other thread






Kernel-Level Threads

- Kernel maintains context information for the process and the threads
 - No thread management done by application
- Scheduling is done on a thread basis
- Windows is an example of this approach

User
Space

Kernel
Space



(b) Pure kernel-level





Advantages of KLT

- The kernel can **simultaneously schedule** multiple threads from the same process on multiple processors.
- If one thread in a process is blocked, the kernel can **schedule another thread** of the same process.
- Kernel routines themselves can be multithreaded.





Disadvantage of KLT


- The transfer of control from one thread to another within the same process requires a **mode switch** to the kernel






Combined Approaches

- Thread creation done in the user space
- Bulk of scheduling and synchronization of threads by the application
- Multiple ULTs are mapped to smaller or equal number of KLTs
- Example is Solaris



(c) Combined





Relationship Between Thread and Processes

Table 4.2 Relationship Between Threads and Processes

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX

