

Big Data Analytics

Prepared By :- Prof . Shital Pathar

Question

- What is Data?
- What is information?
- What is knowledge/insights?

- Data ☐ Information
- Information ☐ Insights



Data contains value and knowledge

History


- **Flat file**
 - Text, comma separated file, program file
- **Database Management System**
 - Structure data
- **Relational Database Management System**
 - Codd at IBM in 1970
 - Oracle, MySQL, Microsoft SQL Server, PostgreSQL
- **Data warehousing**
 - Reporting and Data Analysis
 - Historical + Current + Future

Data Mining

- But to extract the knowledge data needs to be
 - Stored
 - Managed
 - And ANALYZED


Data Mining \approx Big Data \approx Predictive Analytics \approx Data Science

What is Big Data?



“A massive volume of both structured and unstructured data that is so large & complex it's difficult to process with traditional database management tools & software techniques.”

Big Data sources




Big Data is Everywhere!

- Lots of data is being collected and warehoused
 - E-commerce, E-shopping
 - Purchases at Department/ Grocery stores
 - Bank/Credit Card transactions
 - Social Network
 - Web Data



A Minute on the Internet in 2022



Data Units

Name	Equal to:	Size in Bytes		
Bit	1 bit	1/8		
Nibble	4 bits	1/2 (rare)		
Byte	8bits	1		
Kilobyte	1,024 bytes	1,024		
Megabyte	1,024 kilobytes	1,048,576		
Gigabyte	1,024 megabytes	1,073,741,824		
Terrabyte	1,024 gigabytes	1,099,511,627,776		
Petabyte	1,024 terrabytes	1,125,899,906,842,624		
Exabyte	1,024 petabytes	1,152,921,504,606,846,976		
Zettabyte	1,024 exabytes	1,180,591,620,717,411,303,424		
Yottabyte	1,024 zettabytes	1,208,925,819,614,629,174,706,176		

How much data?

- Google processes 20 PB a day (2008)
- Facebook has 2.5 PB of user data + 15 TB/day (2009)
- eBay has 6.5 PB of user data + 50 TB/day (2009)

Type of Data

1. Structured data Relational Data (Tables/Transaction)
2. Unstructured data Text Data (Web)
3. Semi-structured Data (XML)

What to do with these data?

- **Aggregation and Statistics**
 - Data warehouse and OLAP
- **Indexing, Searching, and Querying**
 - Keyword based search
 - Pattern matching
- **Knowledge discovery**
 - Data Mining
 - Statistical Modeling

Big Data: A definition

- Big data is a collection of data sets so large and complex that it becomes difficult to process using on-hand database management tools.
- The challenges include capture, cleaning, storage, search, sharing, analysis, and visualization.
- spot business trends, determine quality of research, prevent diseases, link legal citations, combat crime, and determine real-time roadway traffic conditions.

5 Vs of Big Data

1. Volume
2. Variety
3. Velocity
4. Veracity
5. Value



Big Data Analytics

- Examining large amount of data
- Appropriate information
- Identification of hidden patterns, unknown correlations
- Competitive advantage
- Better business decisions: strategic and operational
- Effective marketing, customer satisfaction, increased revenue

Data analytics vs Data analysis

- Most of us are at least somewhat knowledgeable about the stock market. Imagine that you are a newbie and that you want to start your trade with a profit. Now, describe your initial plan of action.
- As a new trader, you've probably researched sharemarket and trend records to get a sense of what's going on in the market. This technique includes **data analysis**.
- As a result of your newfound understanding of the stock pattern, you can now estimate the stock's future market price and purchase some shares. This serves as an example of a **data analytics** process.


Story of Big Data and Traditional System




Traditional System



Single Cook



Food Shelf



Traditional Processing
System



RDBMS


Traditional Scenario

Traditional Scenario:

2 orders per hour



Single Cook



Food Shelf

Traditional Scenario:

Data is generated at a steady rate and is structured in nature



Traditional Processing System



RDBMS

Failure of Traditional System

Scenario 2:

- They started taking Online orders
- 10 orders per hour




Single Cook
(Regular Computing System)

Food Shelf
(Data)

Big Data Scenario:

Heterogenous data is being generated at an alarming rate by multiple sources



Traditional Processing
System

RDBMS

Issue1 : Too many orders per hour

Solution : Hire Multiple Cooks



Food Shelf
(Data)



Scenario:

Multiple Cook cooking food

Issue:

Food Shelf becomes the BOTTLENECK



Scenario:

Multiple Processing Unit for data processing

Issue:

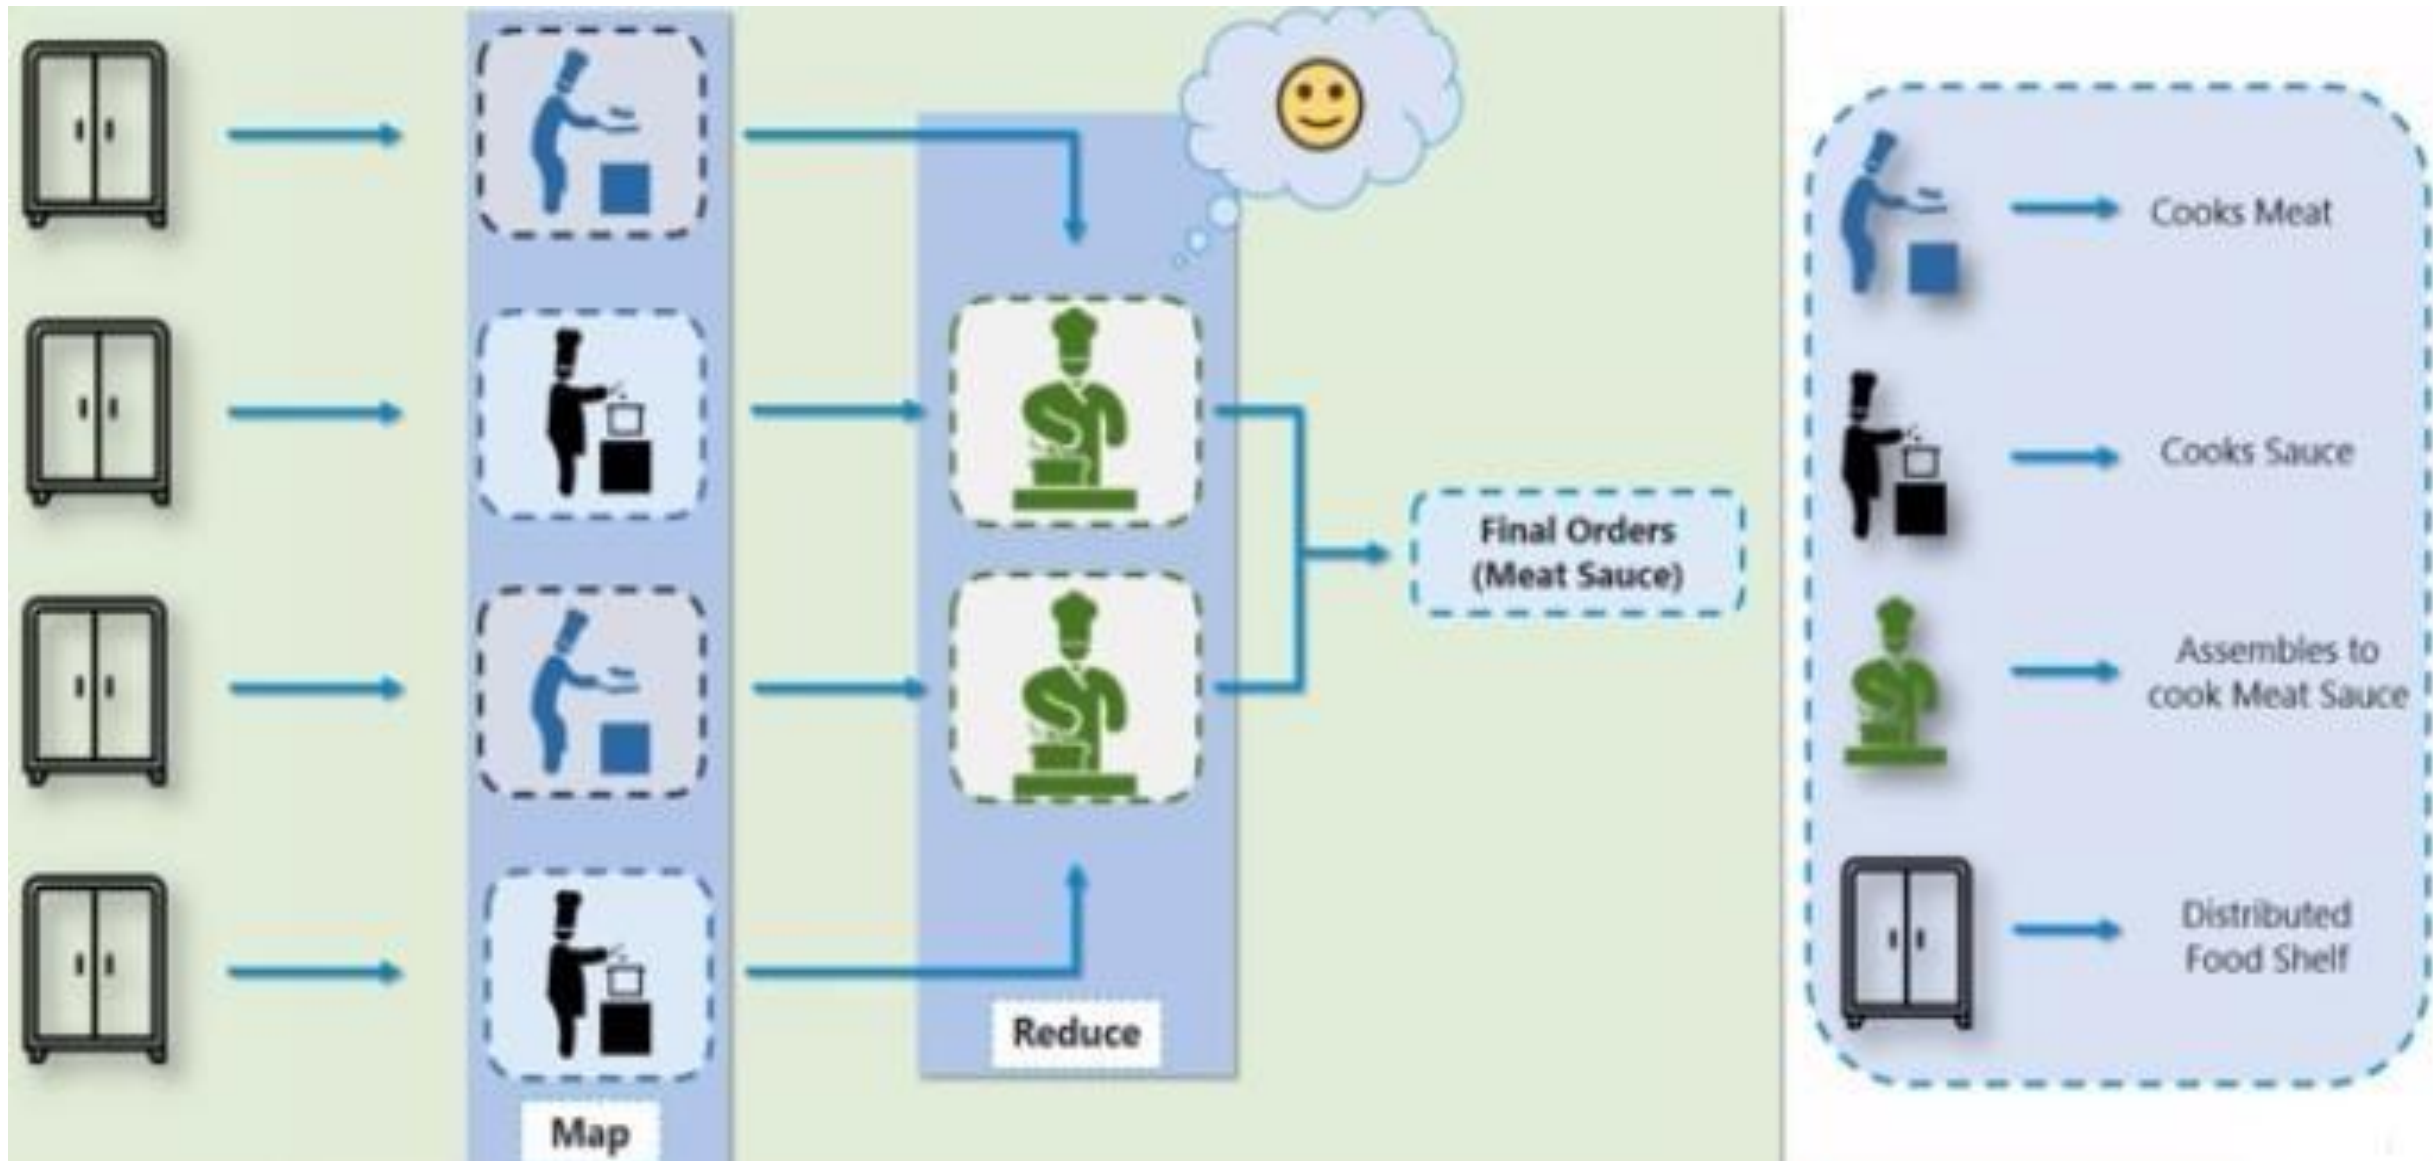
Bringing data to processing generated lots of Network overhead

Data Warehouse


Issue 2 : Food shelf becomes the bottleneck

Solution : Distributed and Parallel approach

Effective Solution



Need a Framework



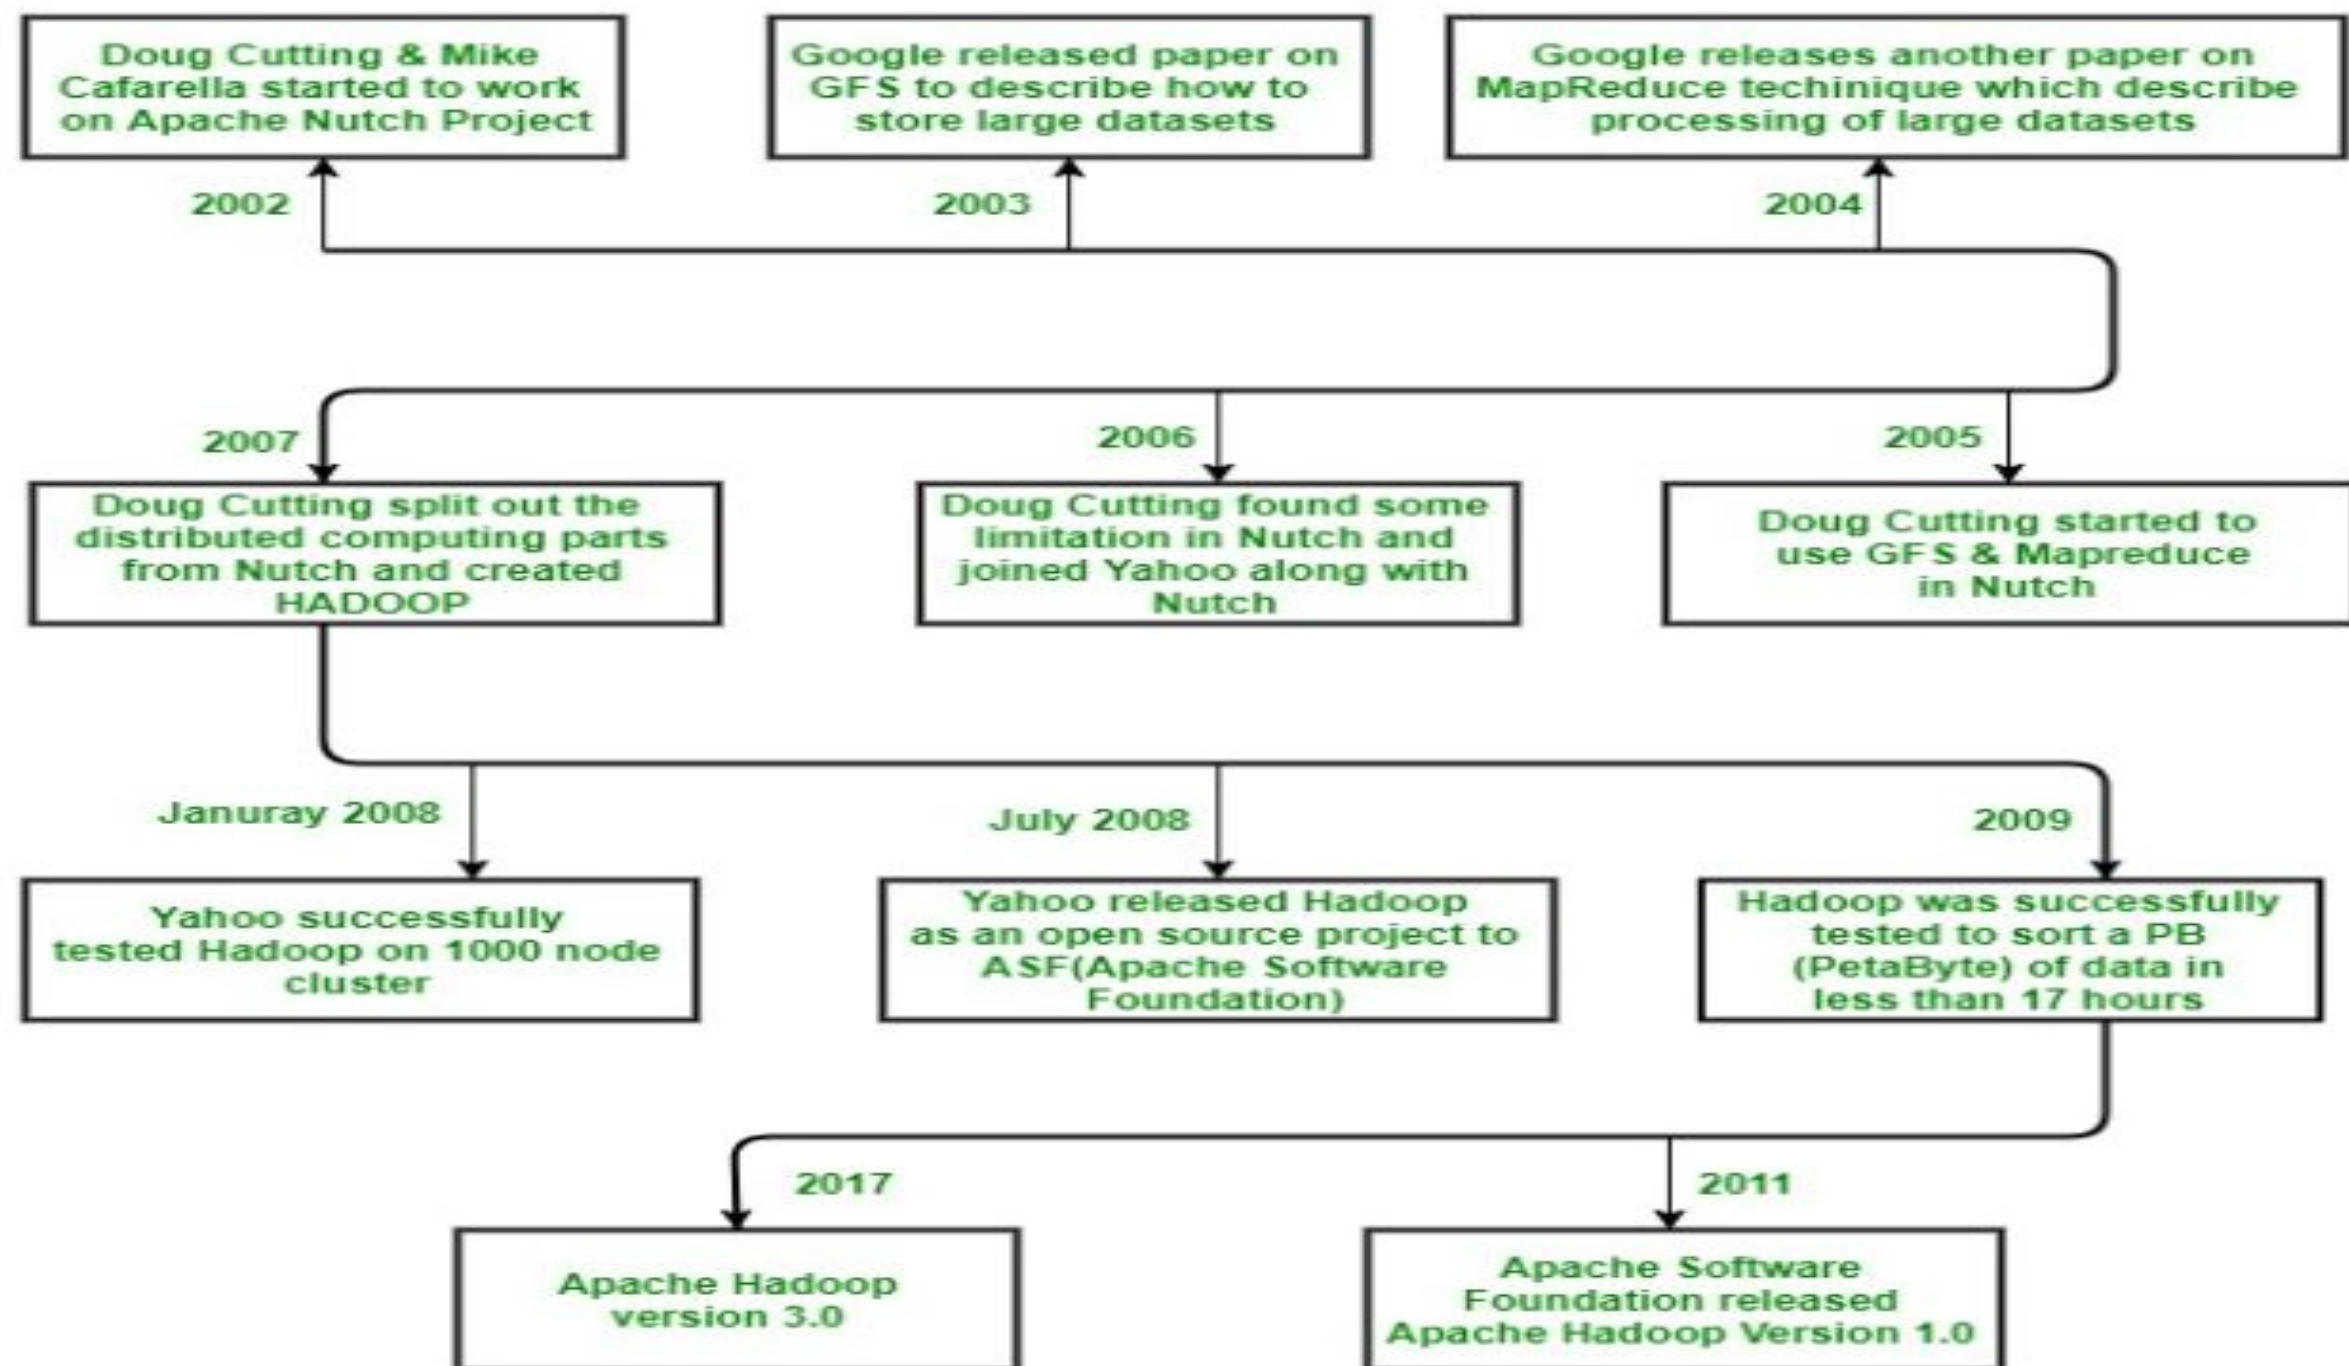
Apache Hadoop: Framework to process Big data

What is Hadoop?

- developed by **Doug cutting**
- Apache Hadoop is an **open-source** framework based on **Google's file system** that can deal with big data in a distributed environment. This distributed environment is built up of a cluster of machines that work closely together to give an impression of a single working machine.

- **Why the hadoop name is given?**

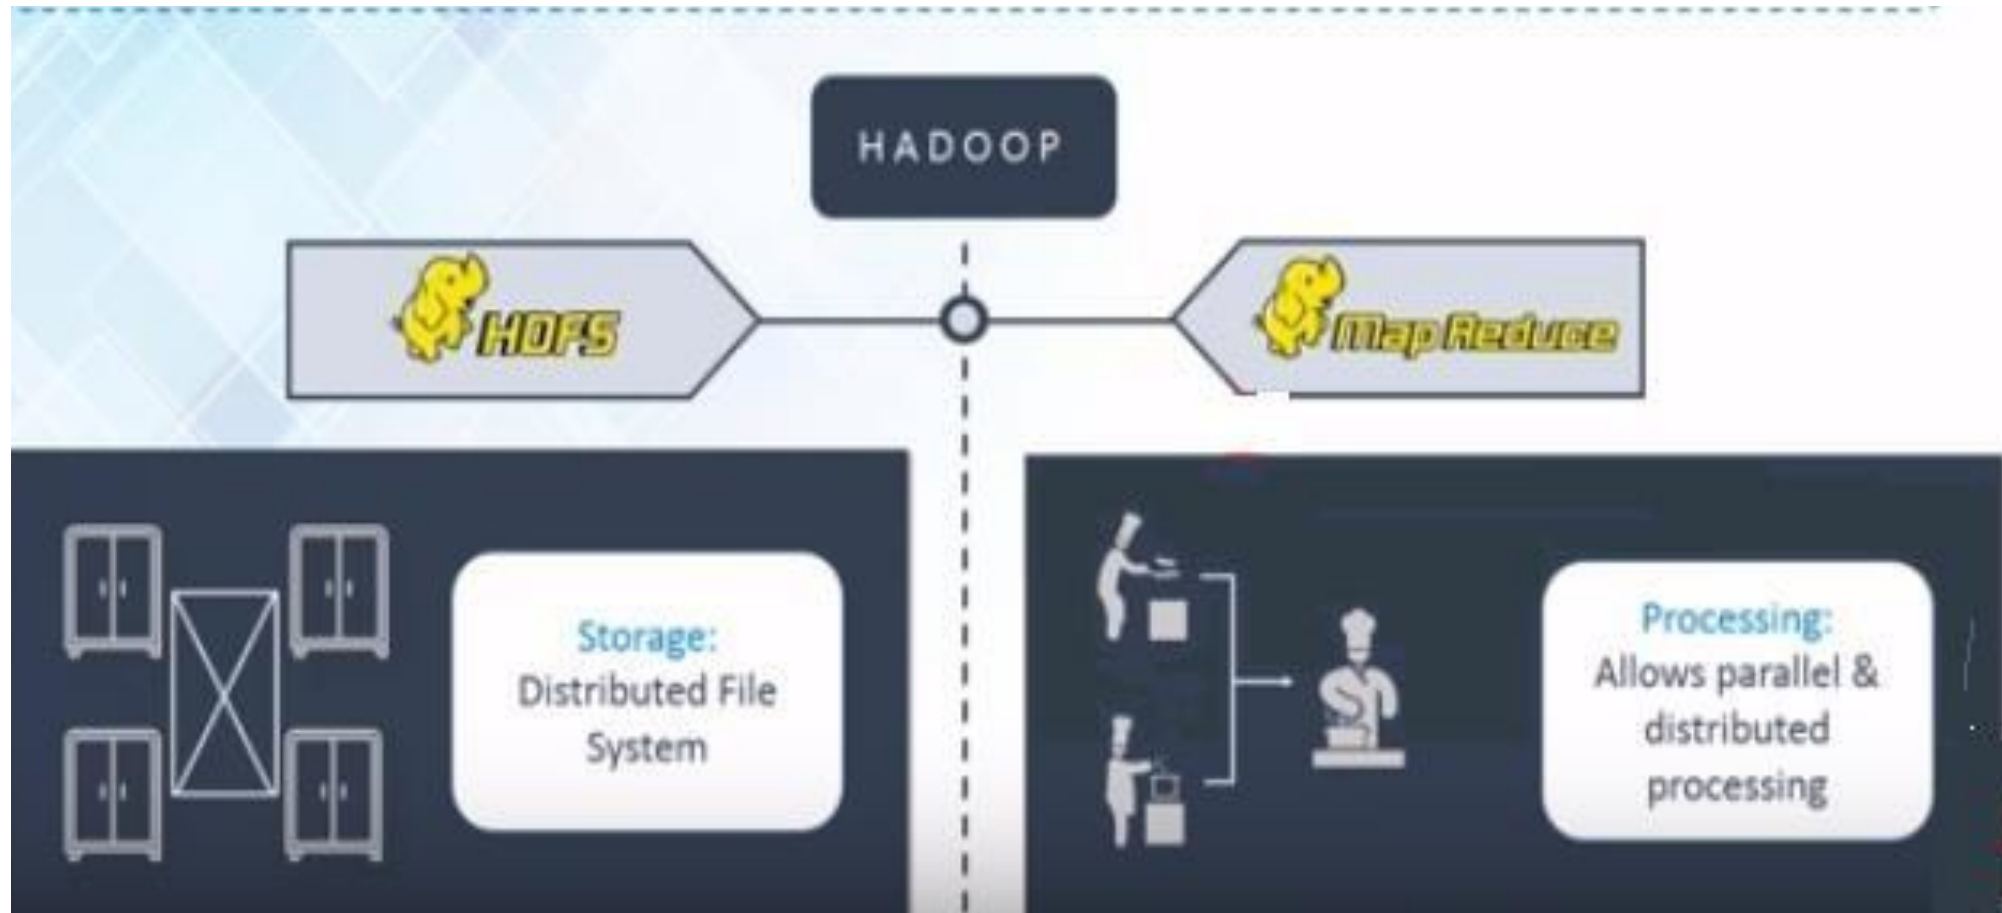




Apache Hadoop:

Framework to process Big data

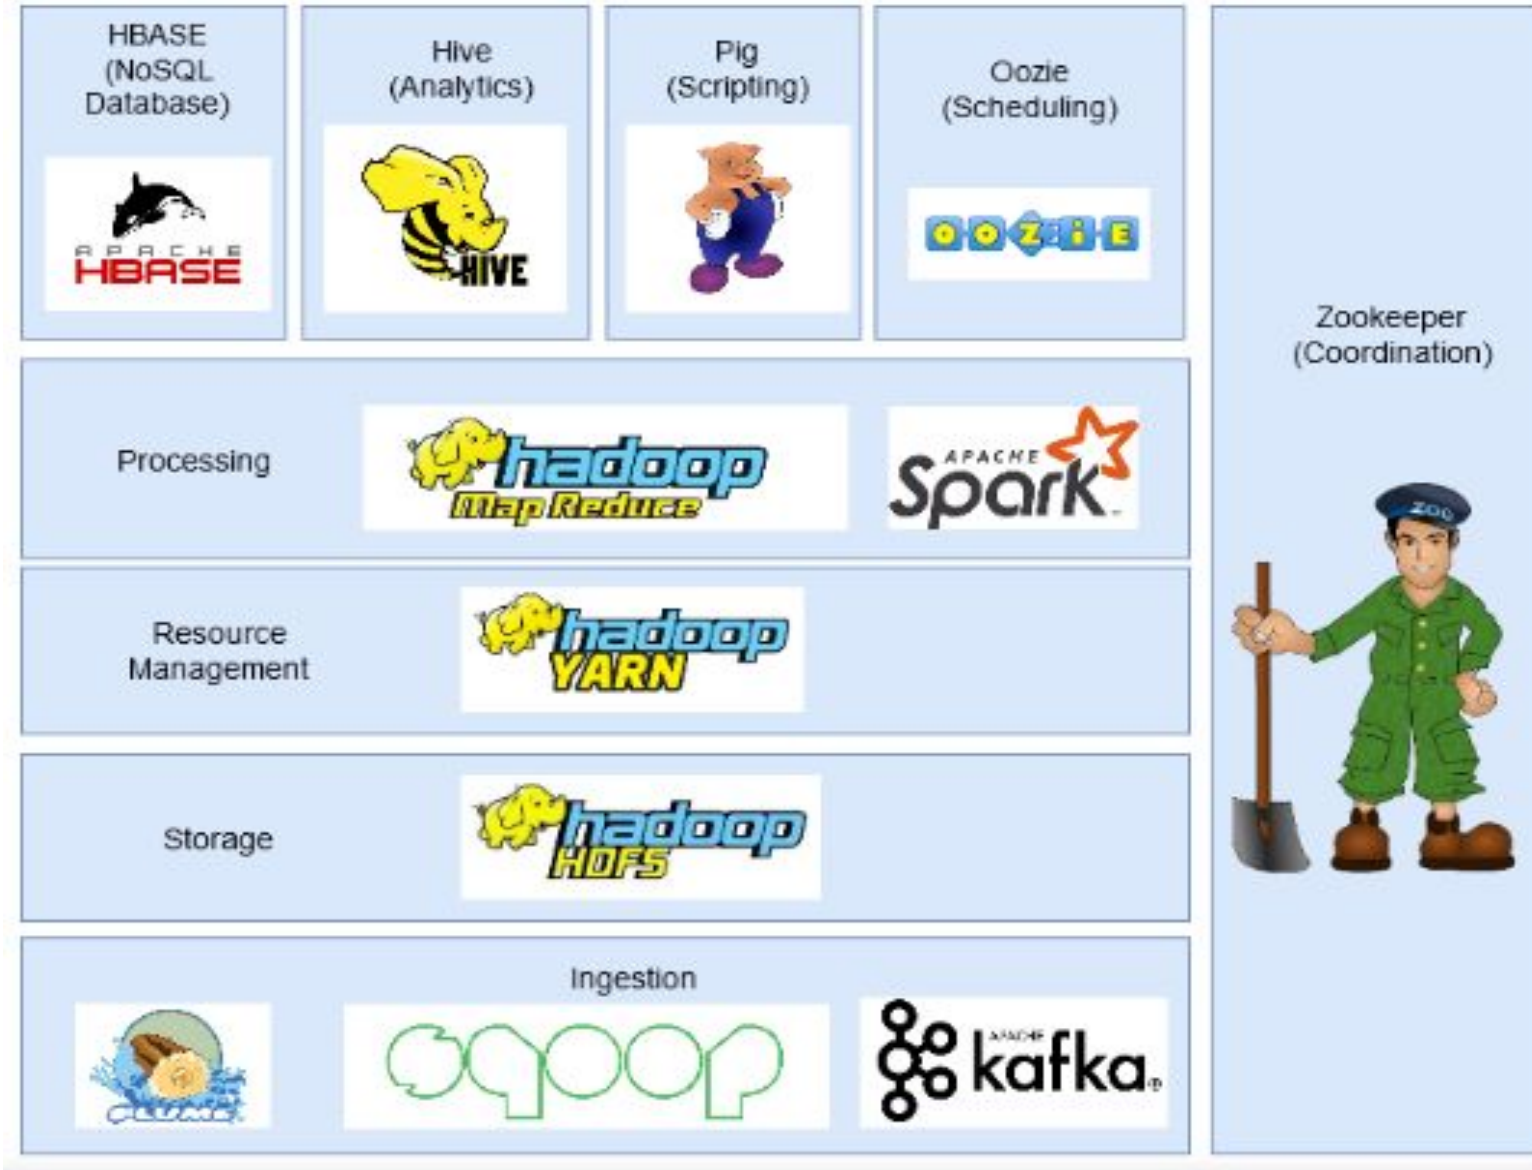
Hadoop is a **Framework** that allows us to **store and process** large data sets in **parallel** and **distributed** fashion.



Advantages of Hadoop Framework

- Hadoop is **highly scalable** because it handles data in a distributed manner
- Compared to vertical scaling in RDBMS, Hadoop offers **horizontal scaling**
- It creates and saves replicas of data making it **fault-tolerant**
- It is **economical** as all the nodes in the cluster are commodity hardware which is nothing but inexpensive machines
- Hadoop utilizes the **data locality concept** to process the data on the nodes on which they are stored rather than moving the data over the network thereby reducing traffic
- It can **handle any type of data**: structured, semi-structured, and unstructured. This is extremely important in today's time because most of our data (emails, Instagram, Twitter, IoT devices, etc.) has no defined format

Hadoop Ecosystem



Components of the Hadoop Ecosystem

- **HDFS (Hadoop Distributed File System)**

HDFS is the primary or major component of Hadoop ecosystem and is responsible for storing large data sets of structured or unstructured data across various nodes

- **MapReduce**

By making the use of distributed and parallel algorithms, MapReduce makes it possible to carry over the processing logic and helps to write applications which transform big data sets into a manageable one.

- **YARN**

Yet Another Resource Negotiator, as the name implies, YARN is the one who helps to manage the resources across the clusters. In short, it performs scheduling and resource allocation for the Hadoop System.

- **HBase**

It's a NoSQL database which supports all kinds of data and thus capable of handling anything of Hadoop Database.

- **Pig**

Pig does the work of executing commands and in the background, all the activities of MapReduce are taken care of. After the processing, pig stores the result in HDFS.

- **Apache Spark**

Apache Spark is an essential product from the Apache software foundation, and it is considered as a powerful data processing engine. Spark is empowering the big data applications around the world. It all started with the increasing needs of enterprises and where MapReduce is unable to handle them.

- **Hive**

HIVE performs reading and writing of large data sets. However, its query language is called as HQL

- **Sqoop**

a. Sqoop is a front-end interface that enables in moving bulk data from Hadoop to relational databases and into variously structured data marts.

- **Flume**

Flume collects, aggregates, and moves large sets of data from its origin and sends it back to HDFS. It works as a fault-tolerant mechanism. It helps in transmitting data from a source into a Hadoop environment.

- **Kafta**

Apache Kafka is a distributed streaming system that is emerging as the preferred solution for integrating real-time data from multiple stream-producing sources and making that data available to multiple stream-consuming systems concurrently

- **Oozie**

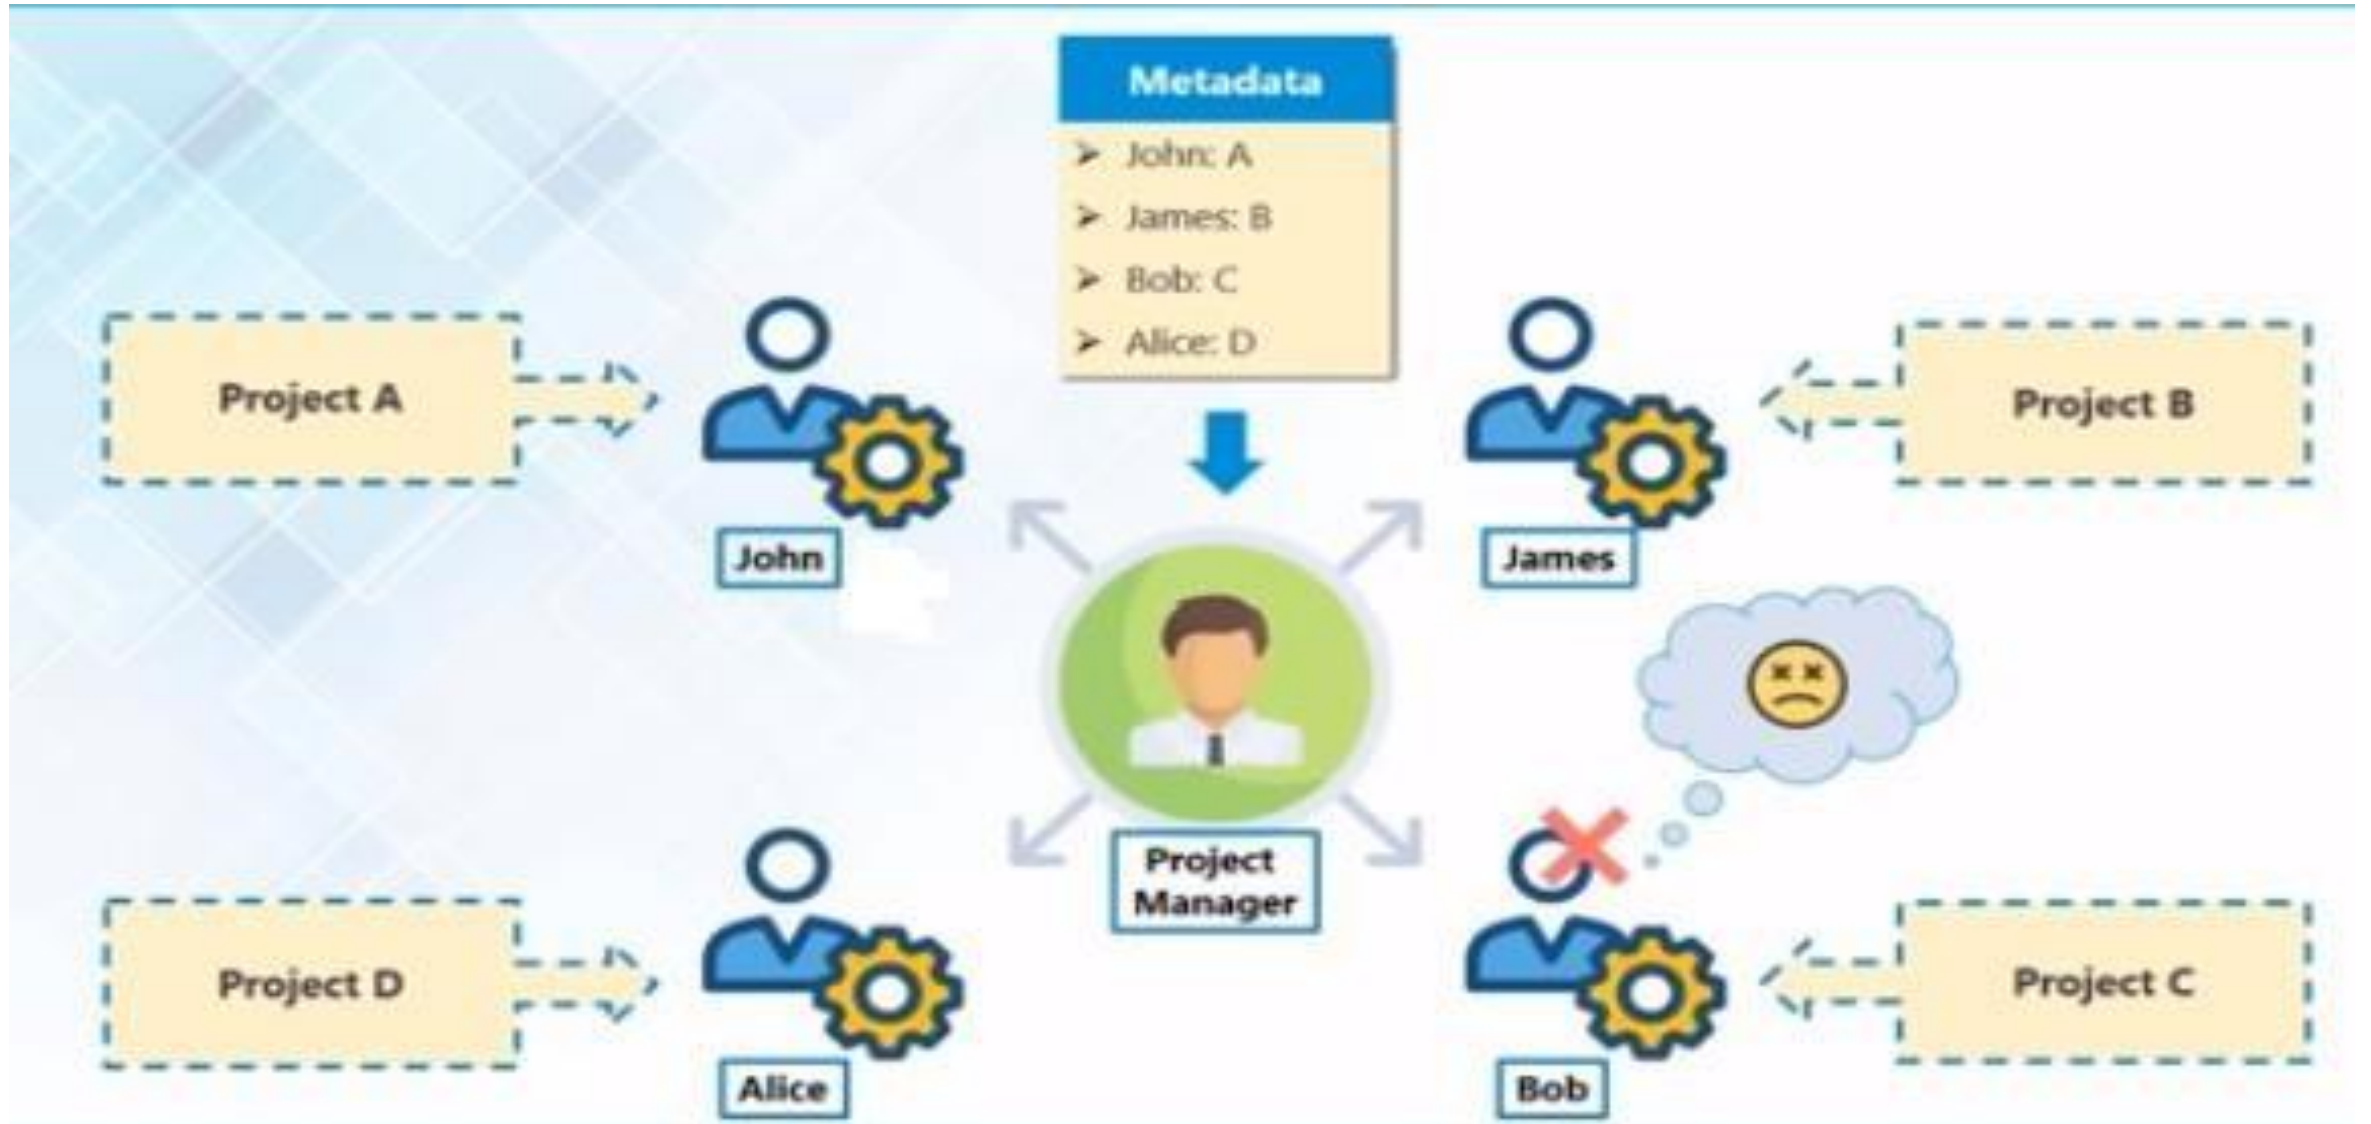
Apache Ooze is a tool in which all sort of programs can be pipelined in a required manner to work in Hadoop's distributed environment. Oozie works as a scheduler system to run and manage Hadoop jobs.

- **Zookeeper**

- Apache Zookeeper is an open-source project designed to coordinate multiple services in the Hadoop ecosystem. Organizing and maintaining a service in a distributed environment is a complicated task. Zookeeper solves this problem with its simple APIs and Architecture. Zookeeper allows developers to focus on core applications instead of concentrating on a distributed environment of the application.

Hadoop: Master/Slave Architecture

Hadoop: Master/Slave Architecture



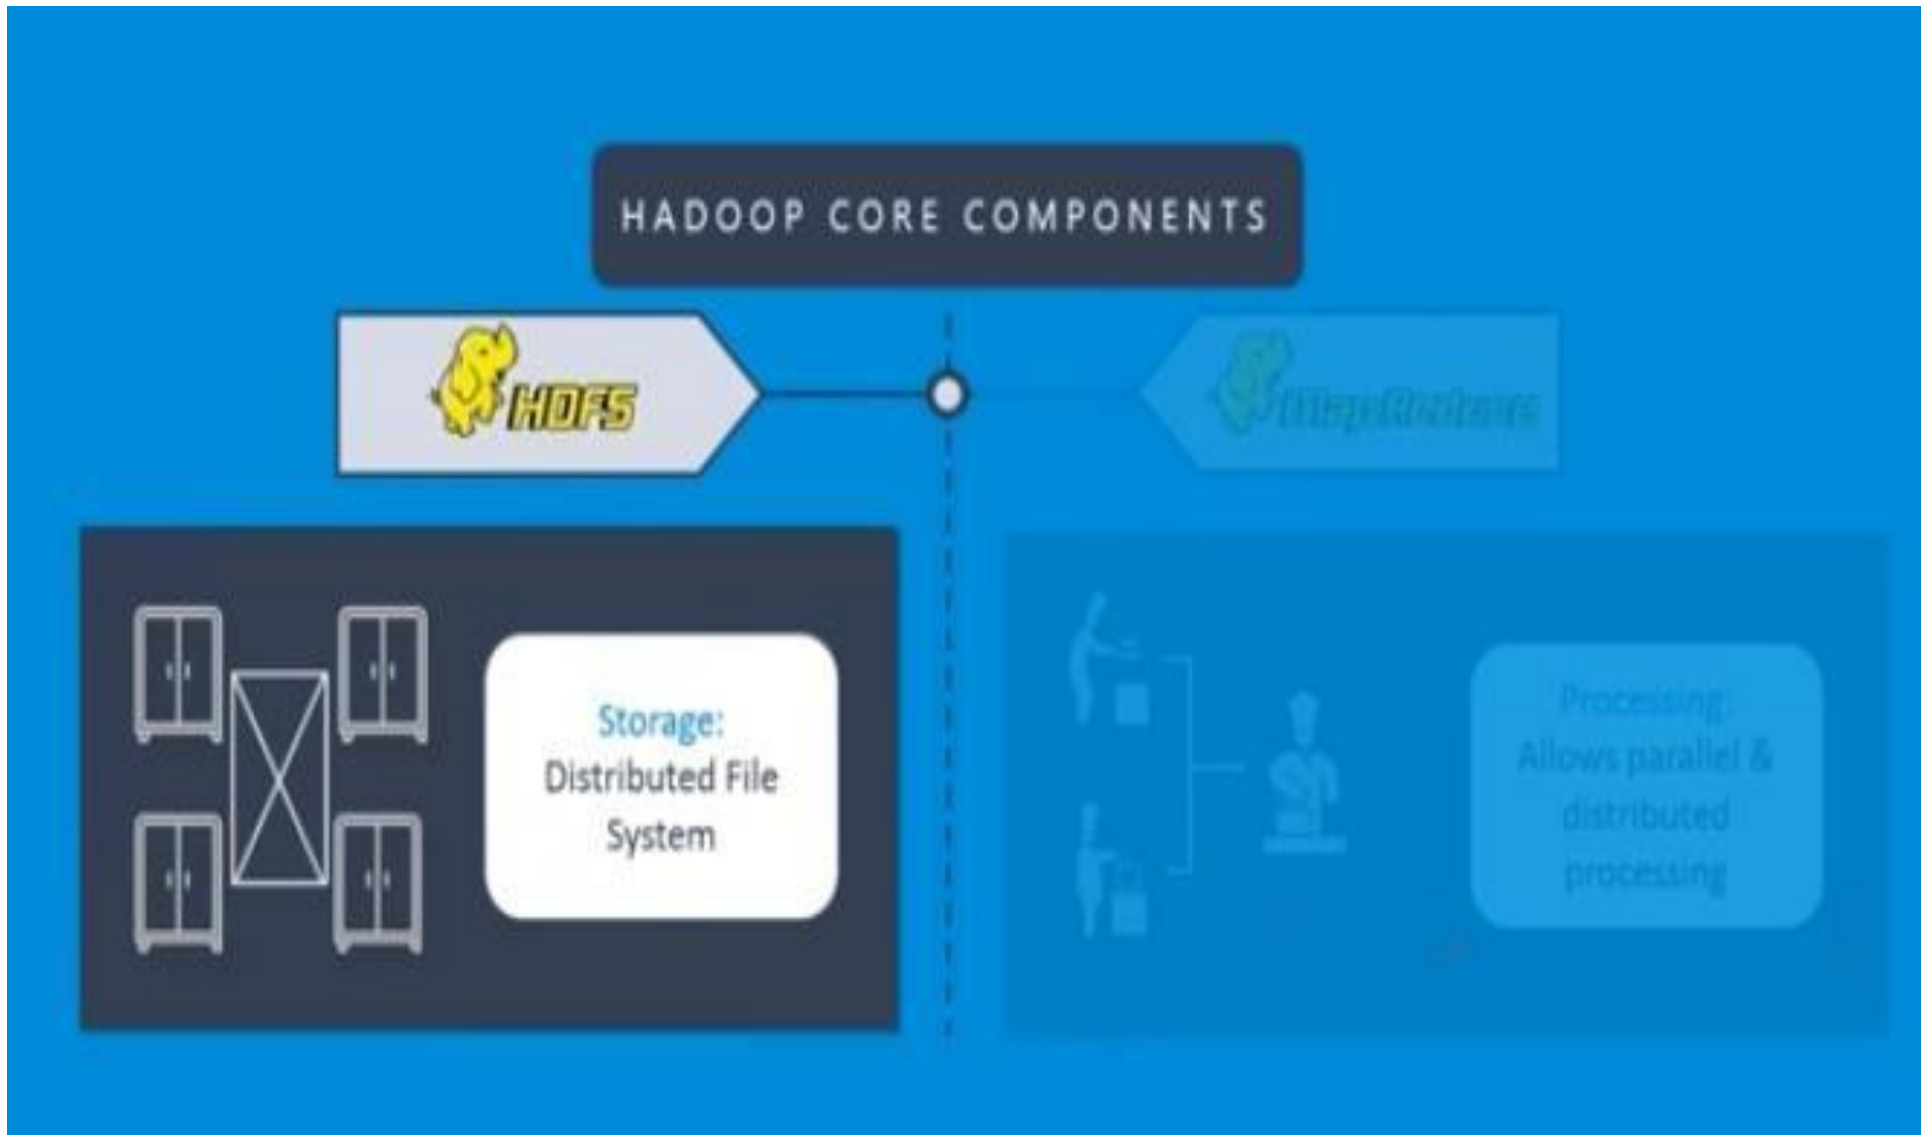
Hadoop: Master/Slave Architecture




Hadoop: Master/Slave Architecture



HDFS



“Moving Computation is Cheaper than Moving Data”



HDFS Core Components

1

Name Node


2

Data Node

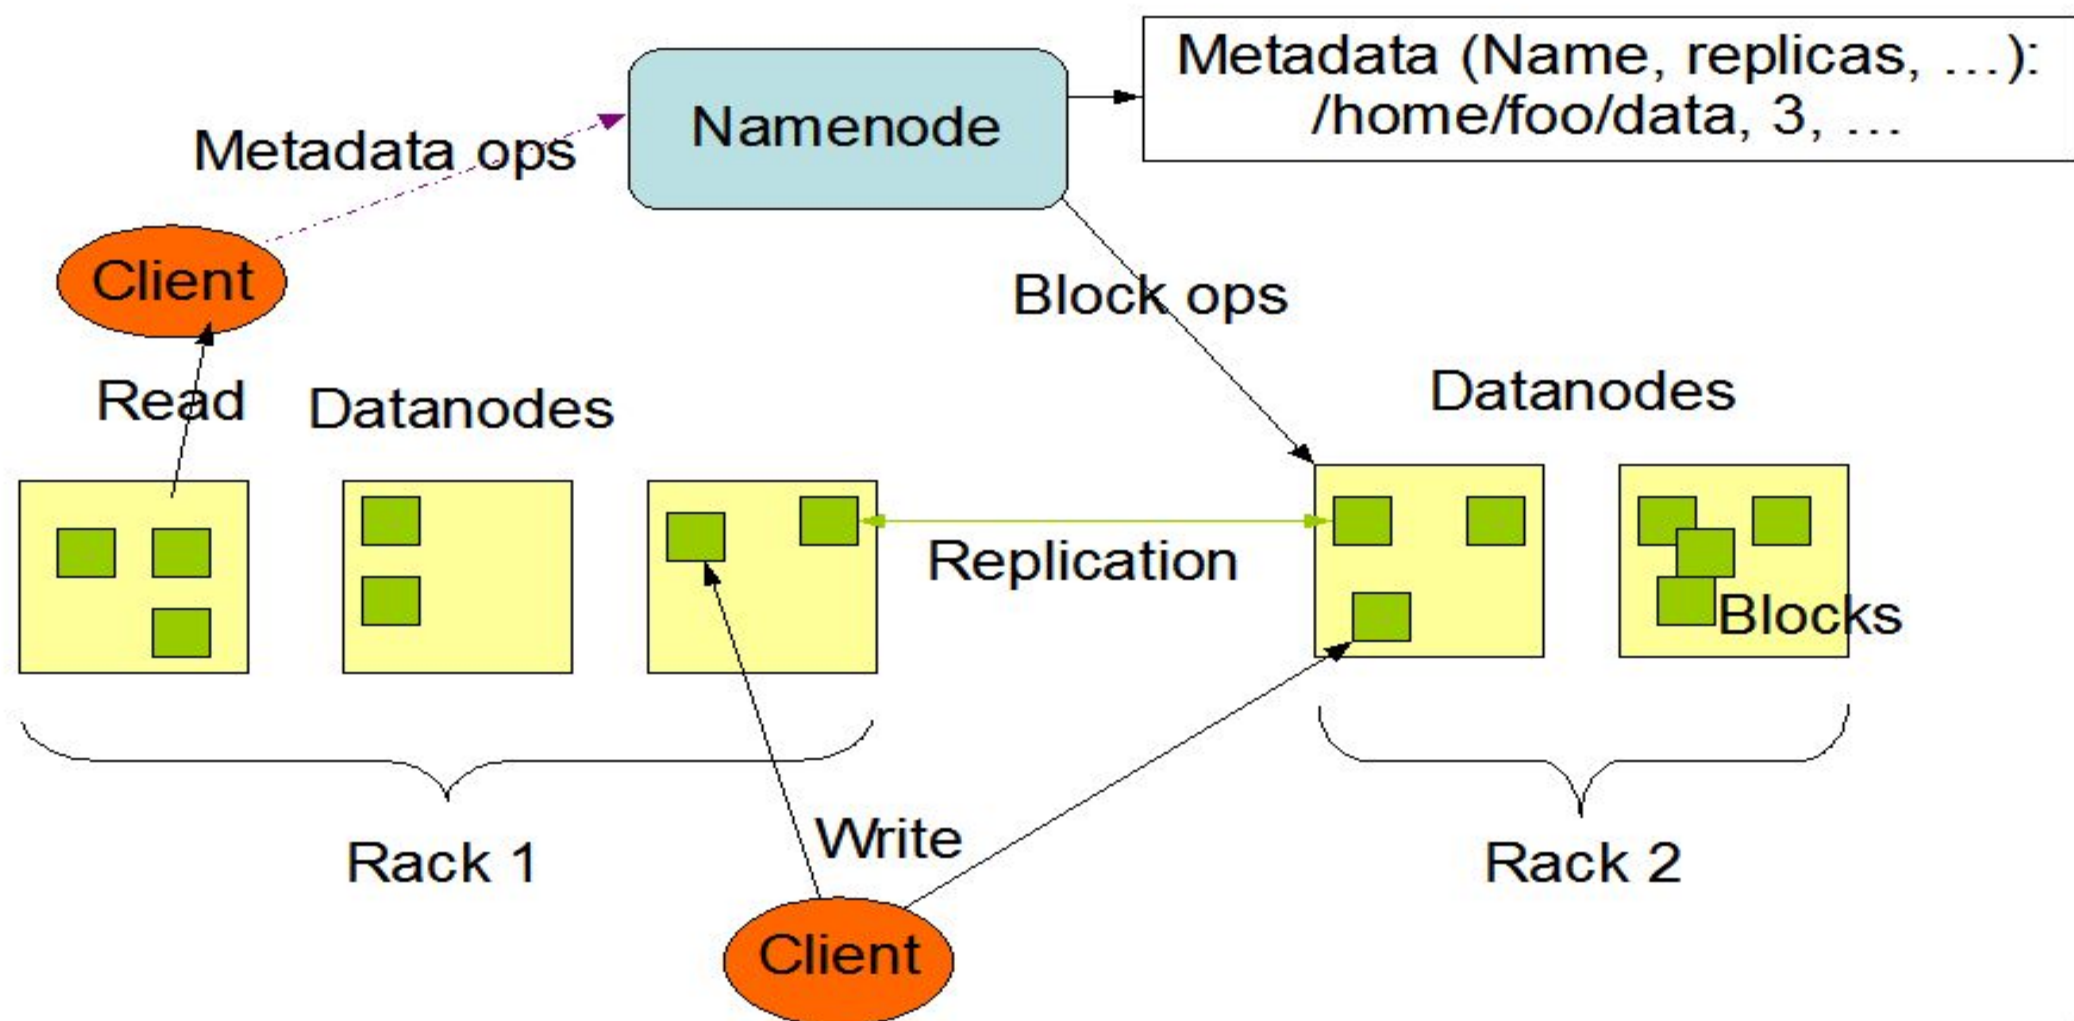
3

**Secondary
Name Node**

NameNode and DataNode



HDFS Architecture



Name Node

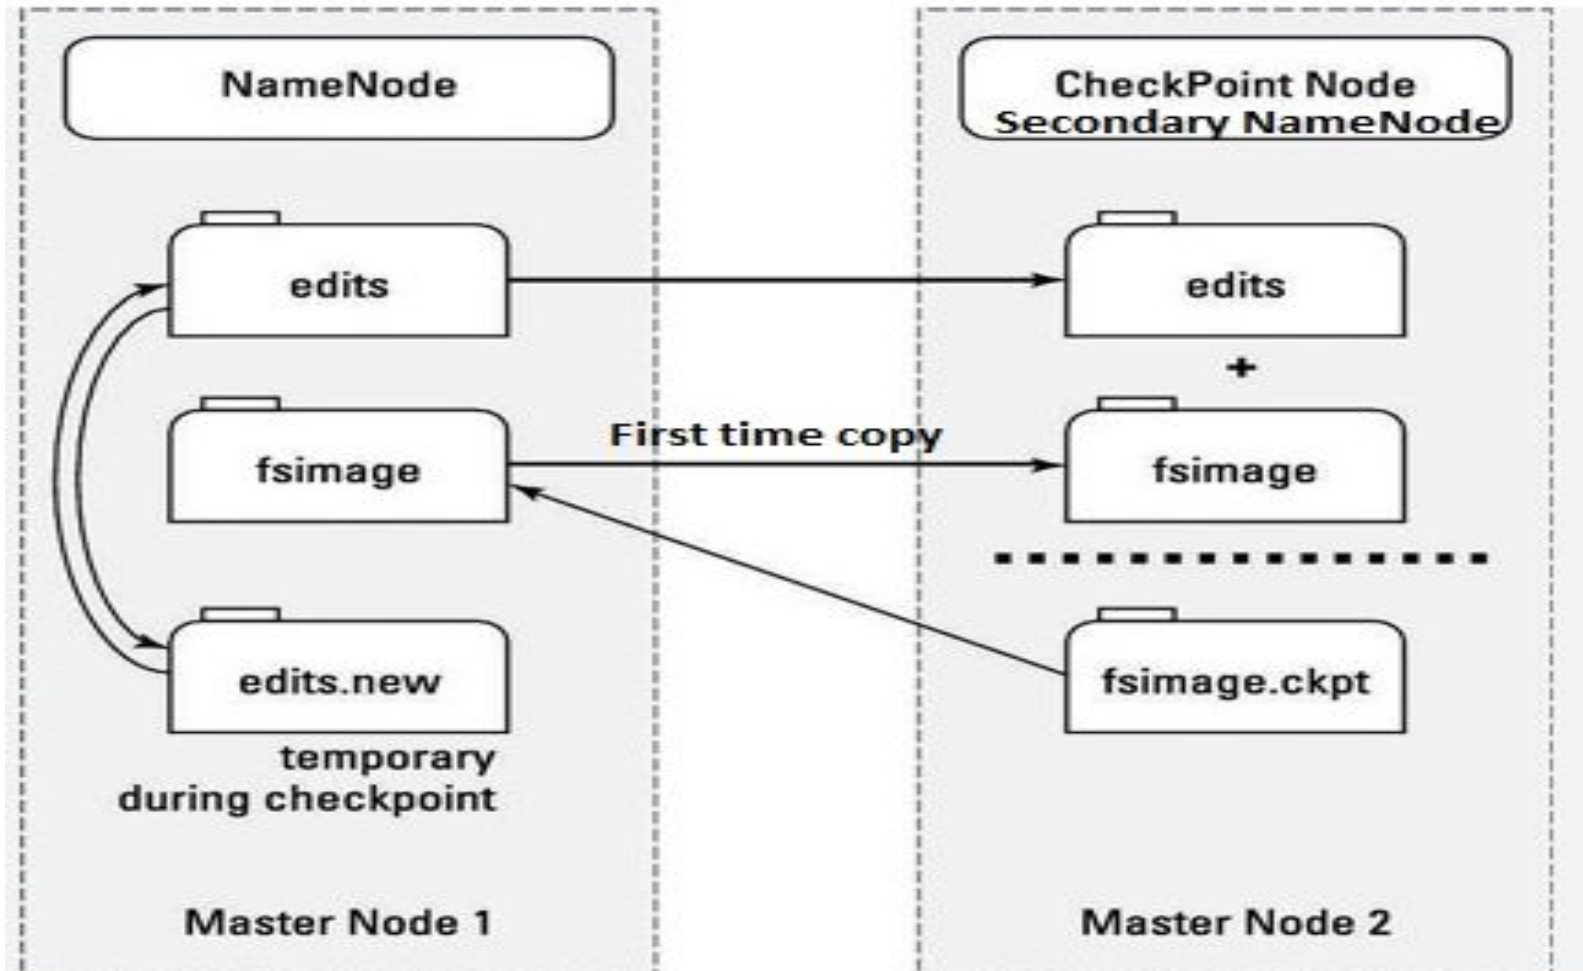
- Maintains and manages Datanodes
- Records metadata i.e . information about the data blocks
(e.g.location of the block stored, the size of file, permissions, hierarchy etc.
- Received heartbeat
- store block report from all data nodes

DataNode

- **Slave** daemons
- stores actual data
- serves read and write requests from the clients

Secondary Namenode & Checkpointing

editLog(in RAM) stores recent changes in RAM,
fsImage(in Disk) stores all changes in Disk



Checkpointing

- Checkpointing is the process of combining FsImage with Editlogs.
- Secondary Namenode take over the responsibility of checkpointing ,therefore making namenode more available.
- prevents EditLogs from getting too large.
- checkpointing happens periodically(default : 1 hour)

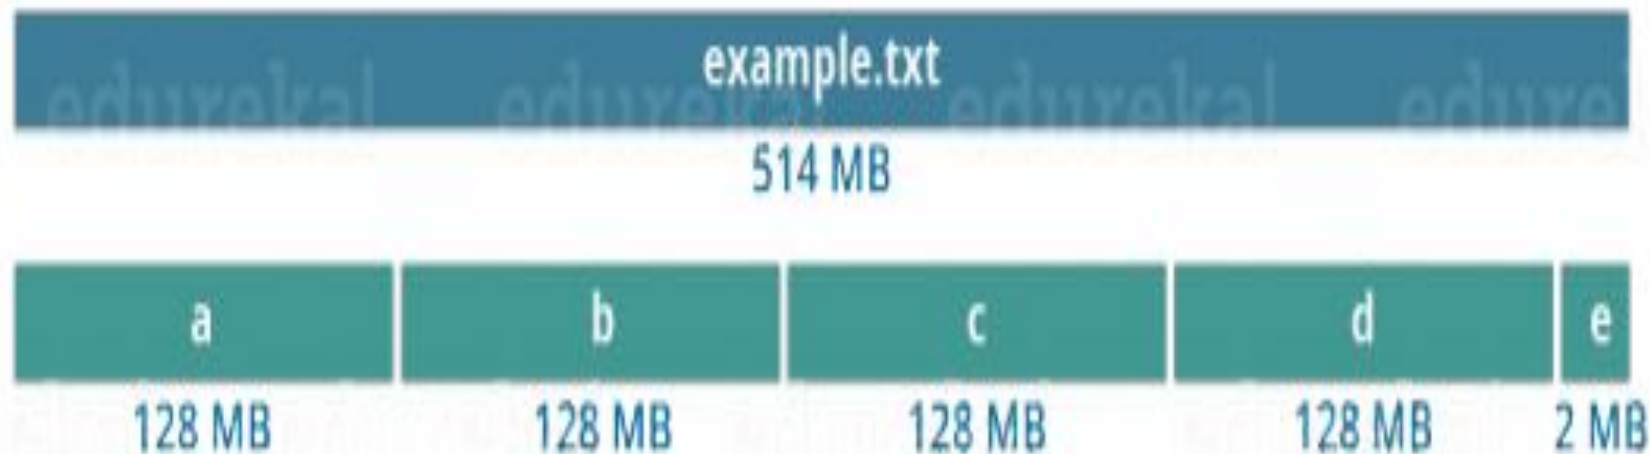
How the data is actually stored in Datanodes? HDFS Data blocks

HDFS

- A Hadoop **cluster** is a collection of computers, known as nodes, that are networked together to perform parallel computations on big data sets.
- Divide the large file into pieces called **blocks**.
- It uses **rackId** to identify the datanode from rack .
- A **rack** is the collection of datanodes within the cluster.

HDFS Data Blocks

- Blocks are the nothing but the smallest continuous location on your hard drive where data is stored.
- In general, in any of the File System, you store the data as a collection of blocks. Similarly, HDFS stores each file as blocks which are scattered throughout the Apache Hadoop cluster.
- The default size of each block is **128 MB** in Apache Hadoop 2.x (64 MB in Apache Hadoop 1.x) which you can configure as per your requirement.

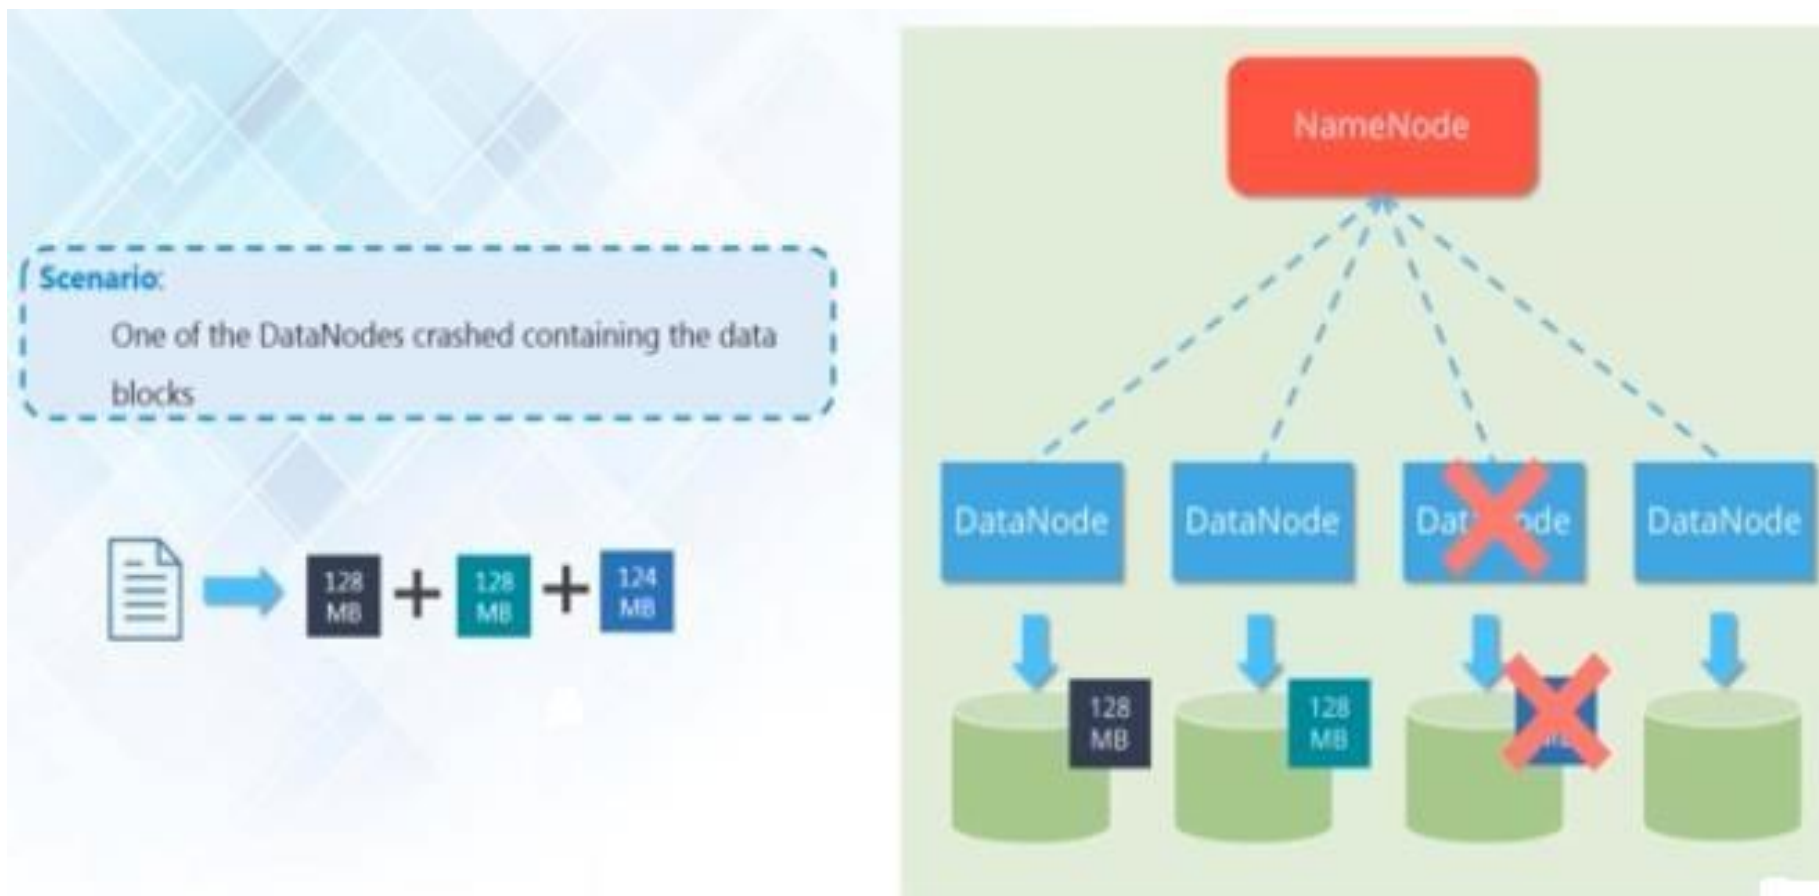


Advantages of HDFS

- If data is more than the total size of a Hadoop cluster, more computers can be added to cluster.
- If a huge data file is stored on a single machine and it takes 4 seconds to process that file, in Hadoop cluster that data file can be divided into cluster and processing time can be reduced.

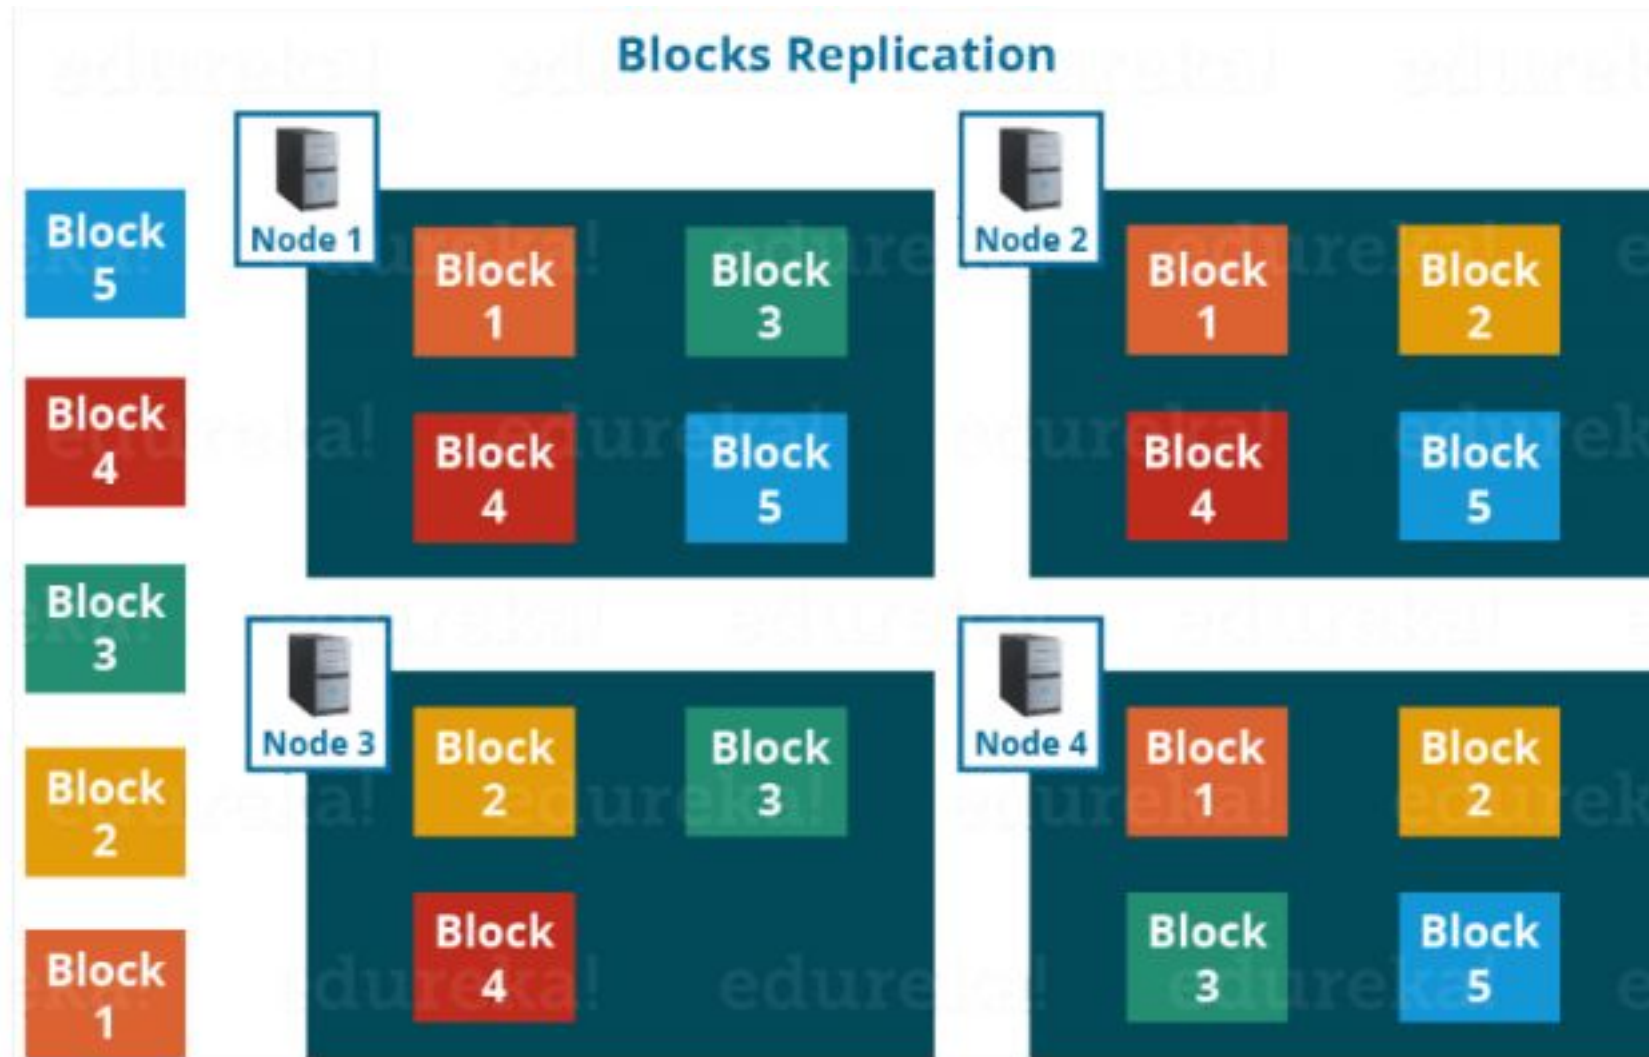
Fault Tolerance:
How Hadoop cope up with Datanode failure?

Fault Tolerance



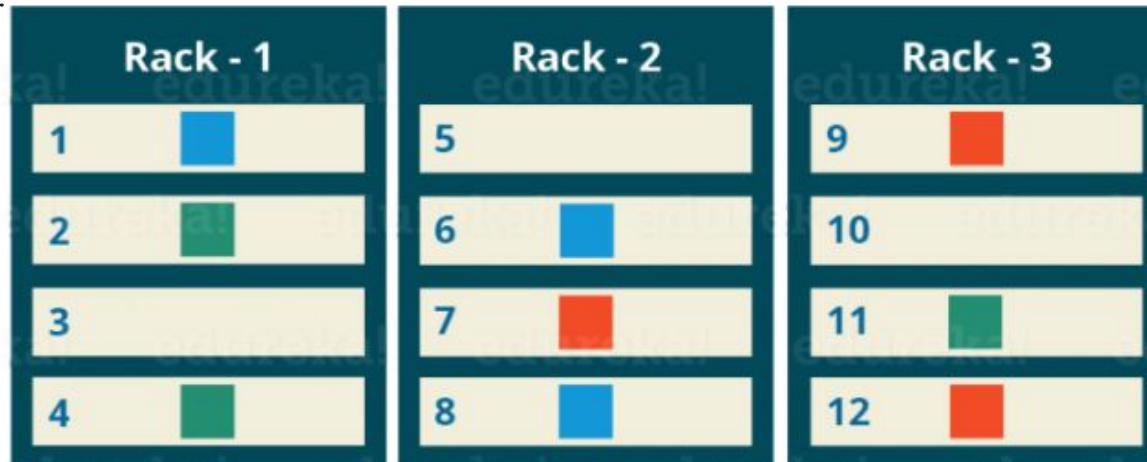
Solution : Replication factor

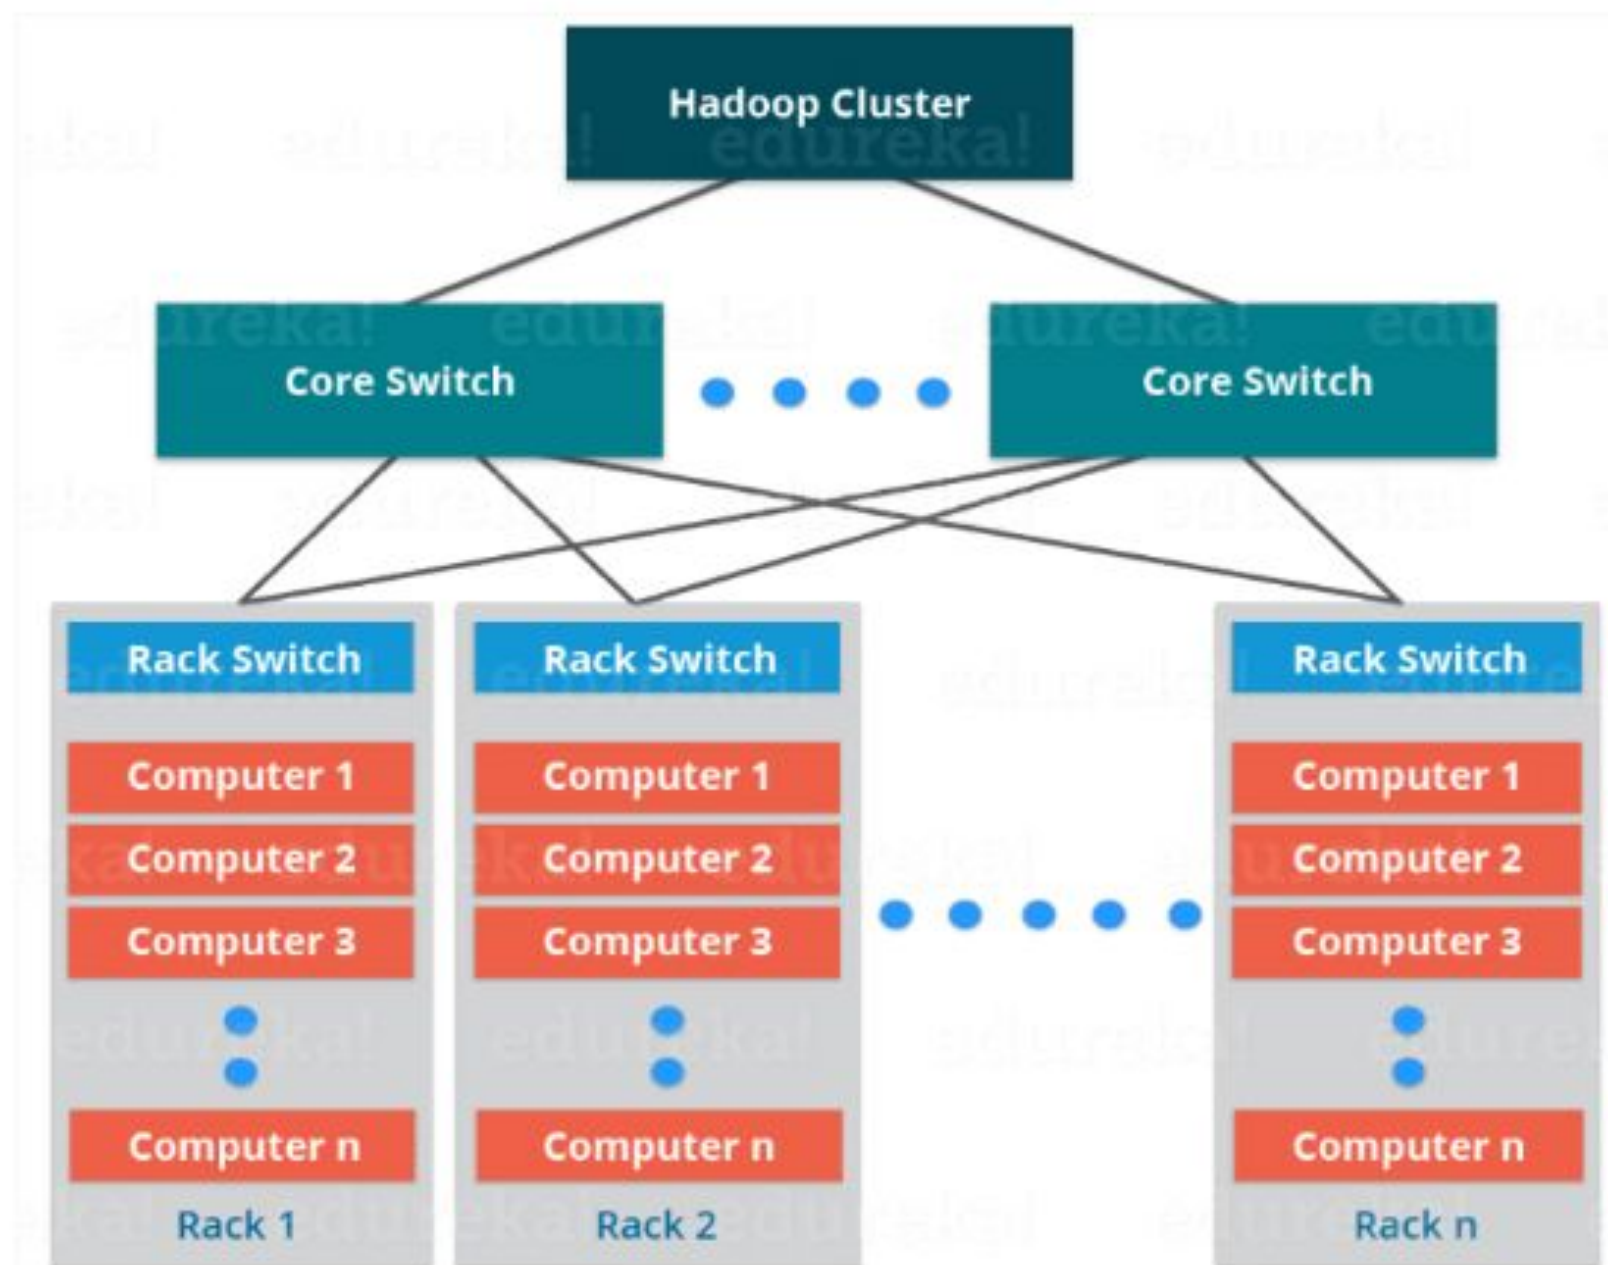
Fault Tolerance: Replication Factor



Rack Awareness

- Hadoop components are rack-aware. For example, HDFS block placement will use rack awareness for fault tolerance by placing one block replica on a different rack. This provides data availability in the event of a network switch failure or partition within the cluster.
- Hadoop master daemons obtain the rack id of the cluster slaves by invoking either an external script or java class as specified by configuration files. Using either the java class or external script for topology, output must adhere to the java **org.apache.hadoop.net.DNSToSwitchMapping** interface.
- The interface expects a one-to-one correspondence to be maintained and the topology information in the format of `‘/myrack/myhost’`, where `‘/’` is the topology delimiter, `‘myrack’` is the rack identifier, and `‘myhost’` is the individual host. Assuming a single /24 subnet per rack, one could use the format of `‘/192.168.100.0/192.168.100.5’` as a unique rack-host topology mapping.





HDFS Write Mechanism

Suppose a situation where an HDFS client, wants to write a file named “example.txt” of size 248 MB.


Assume that the system block size is configured for 128 MB (default). So, the client will be dividing the file “example.txt” into 2 blocks – one of 128 MB (Block A) and the other of 120 MB (block B).

- Suppose, the NameNode provided following lists of IP addresses to the client:
 - For Block A, list A = {IP of DataNode 1, IP of DataNode 4, IP of DataNode 6}
 - For Block B, set B = {IP of DataNode 3, IP of DataNode 7, IP of DataNode 9}

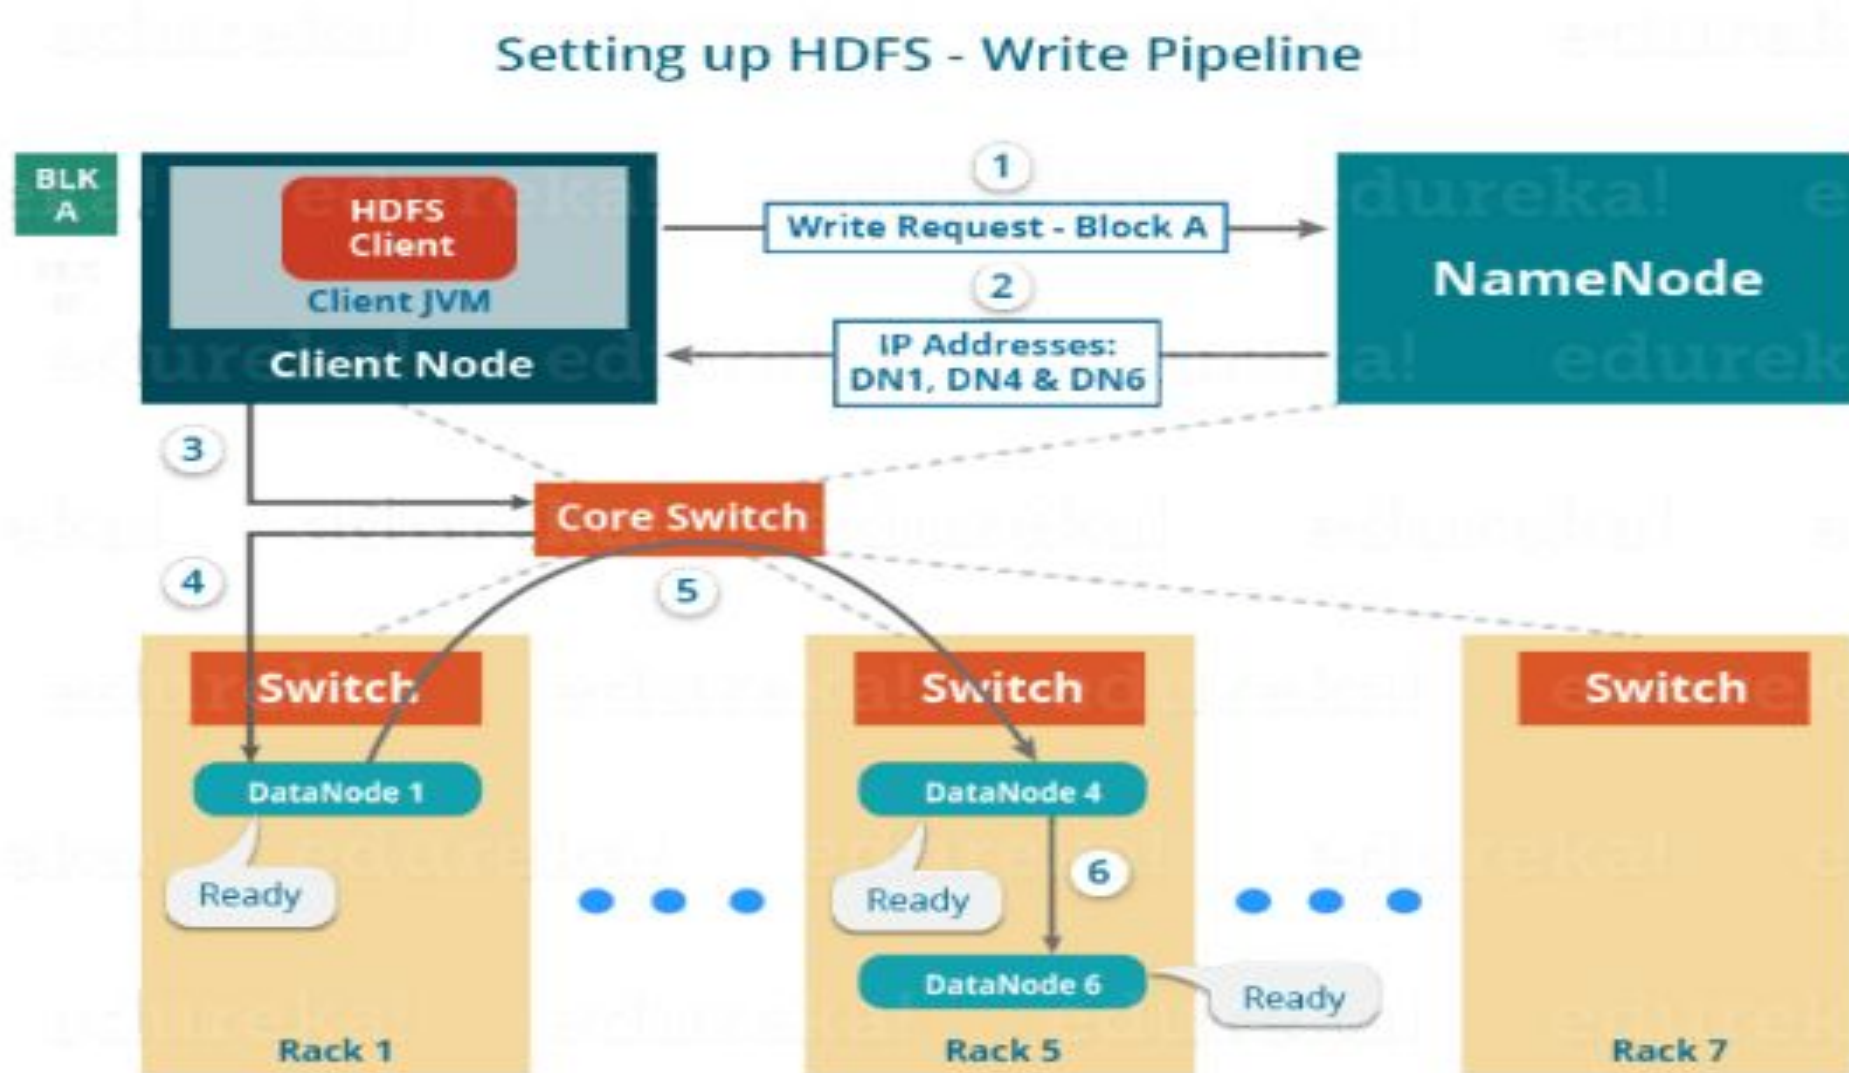


HDFS Write Mechanism

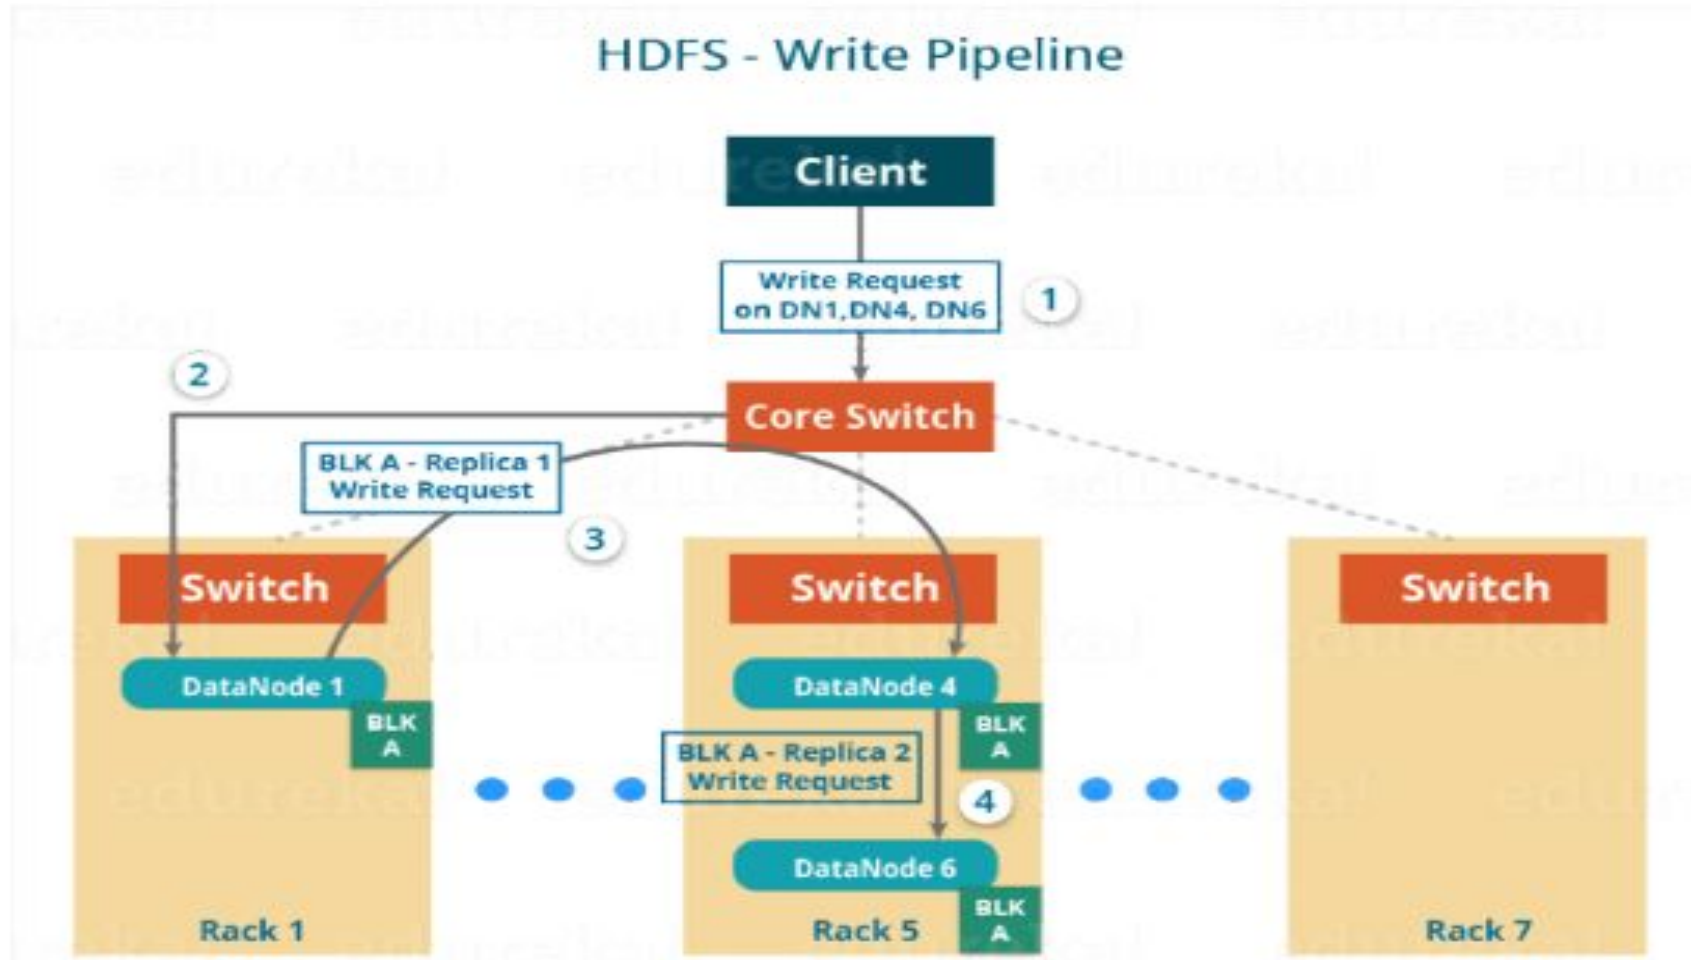
1. Set up of Pipeline
2. Data streaming and replication
3. Shutdown of Pipeline (Acknowledgement stage)



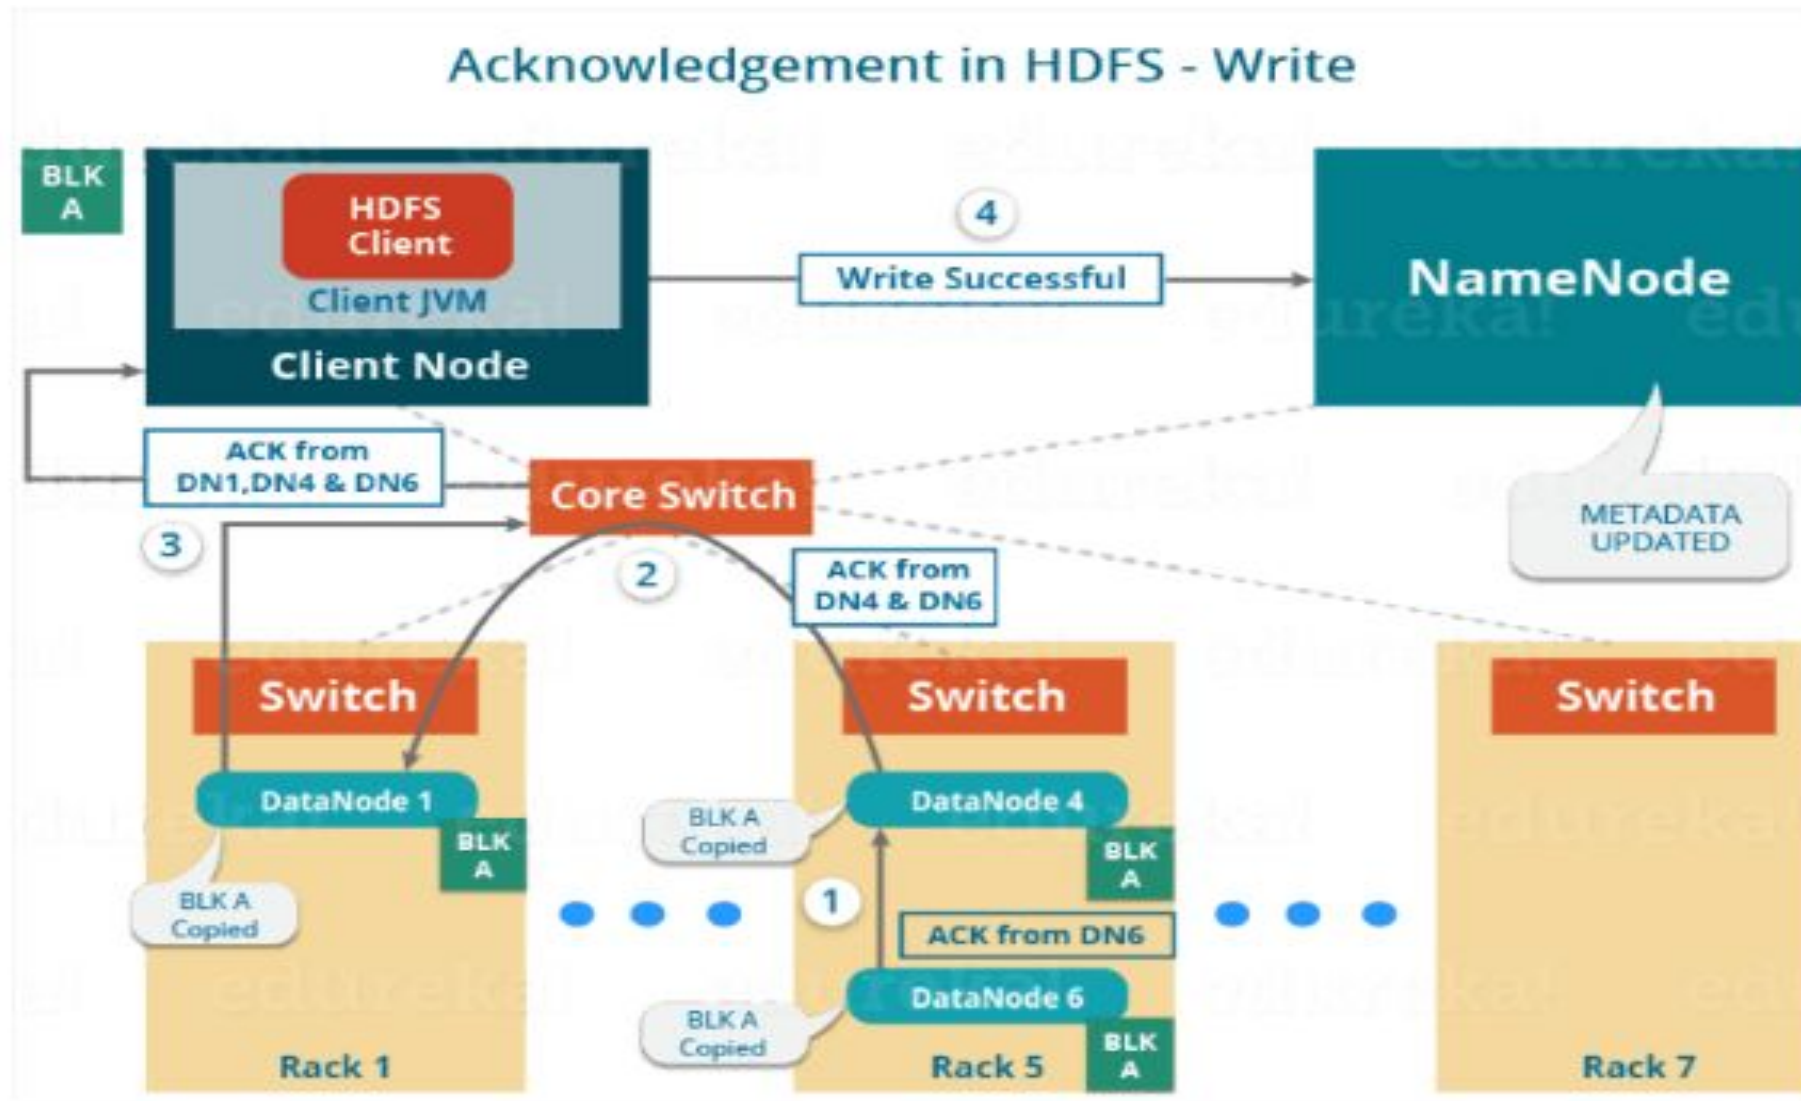
HDFS Write Mechanism – Pipeline Setup



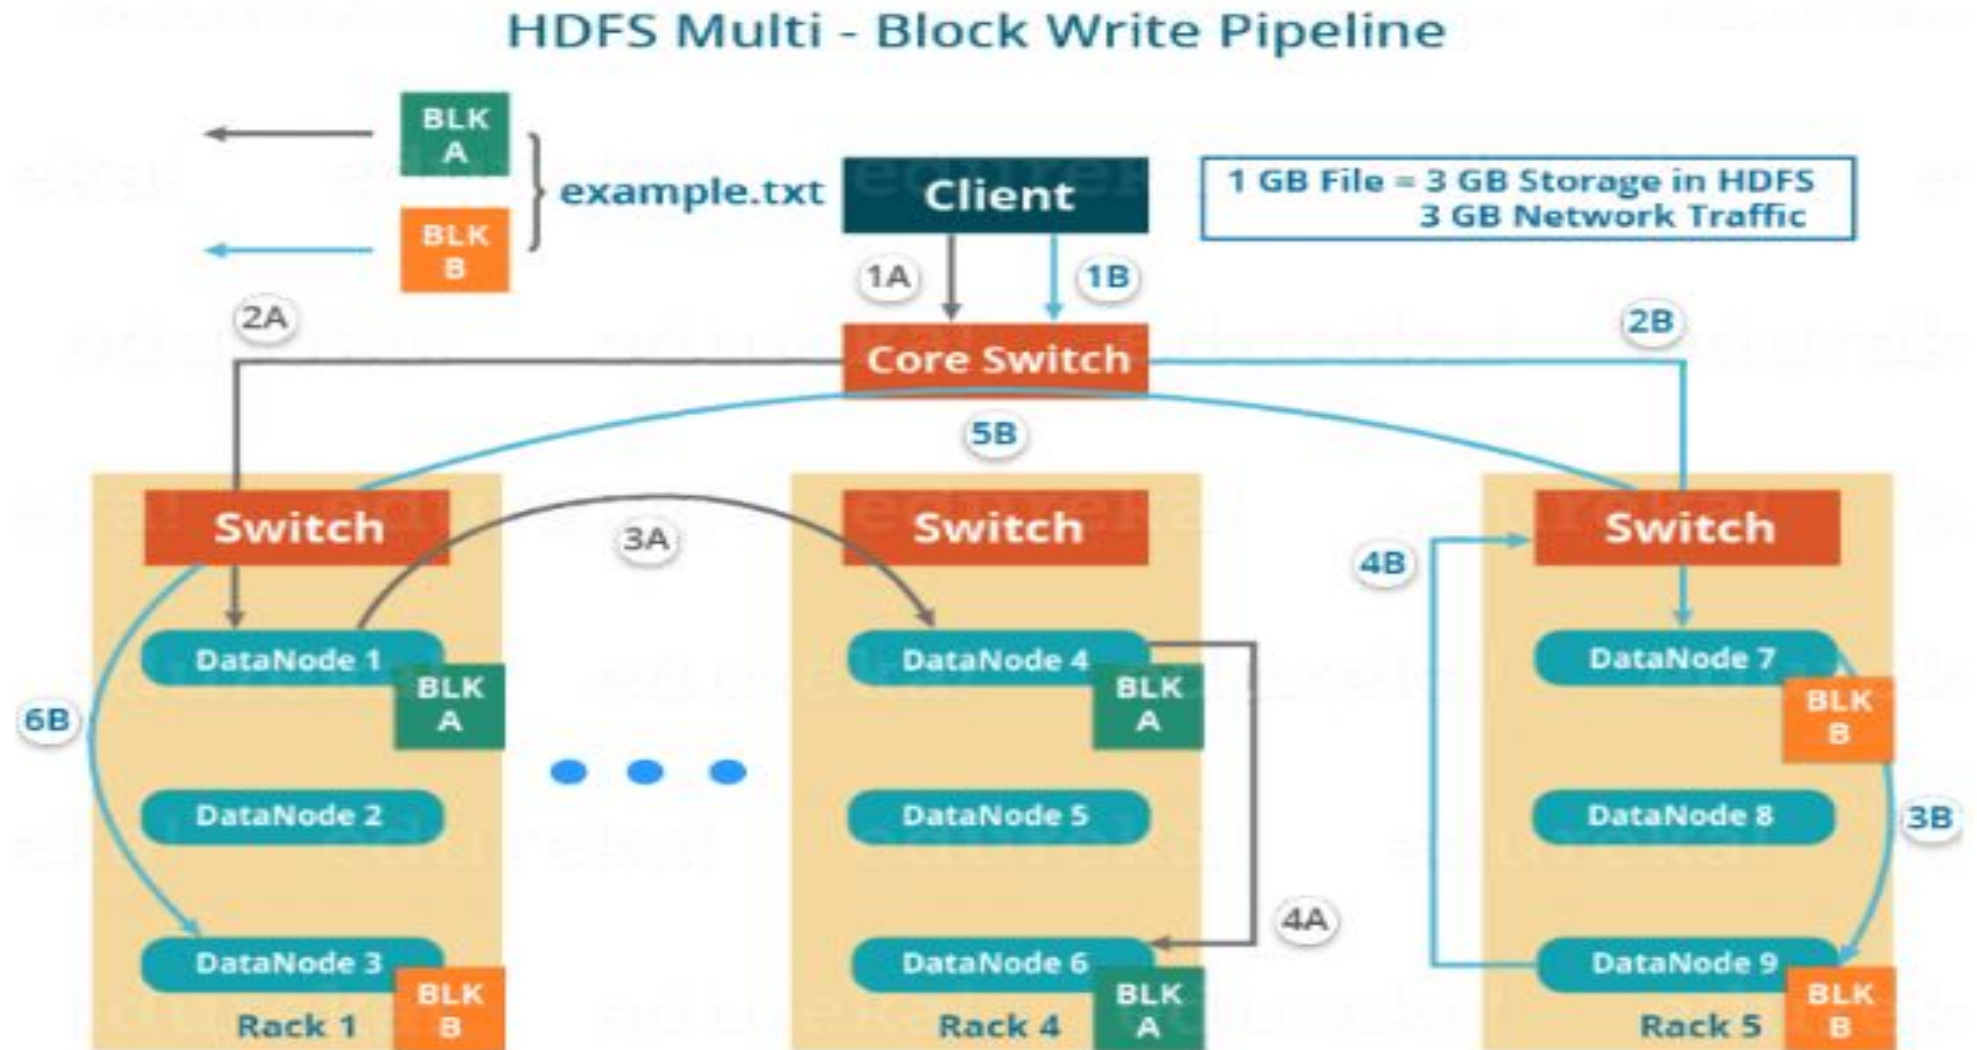
HDFS Write Mechanism –Data Streaming



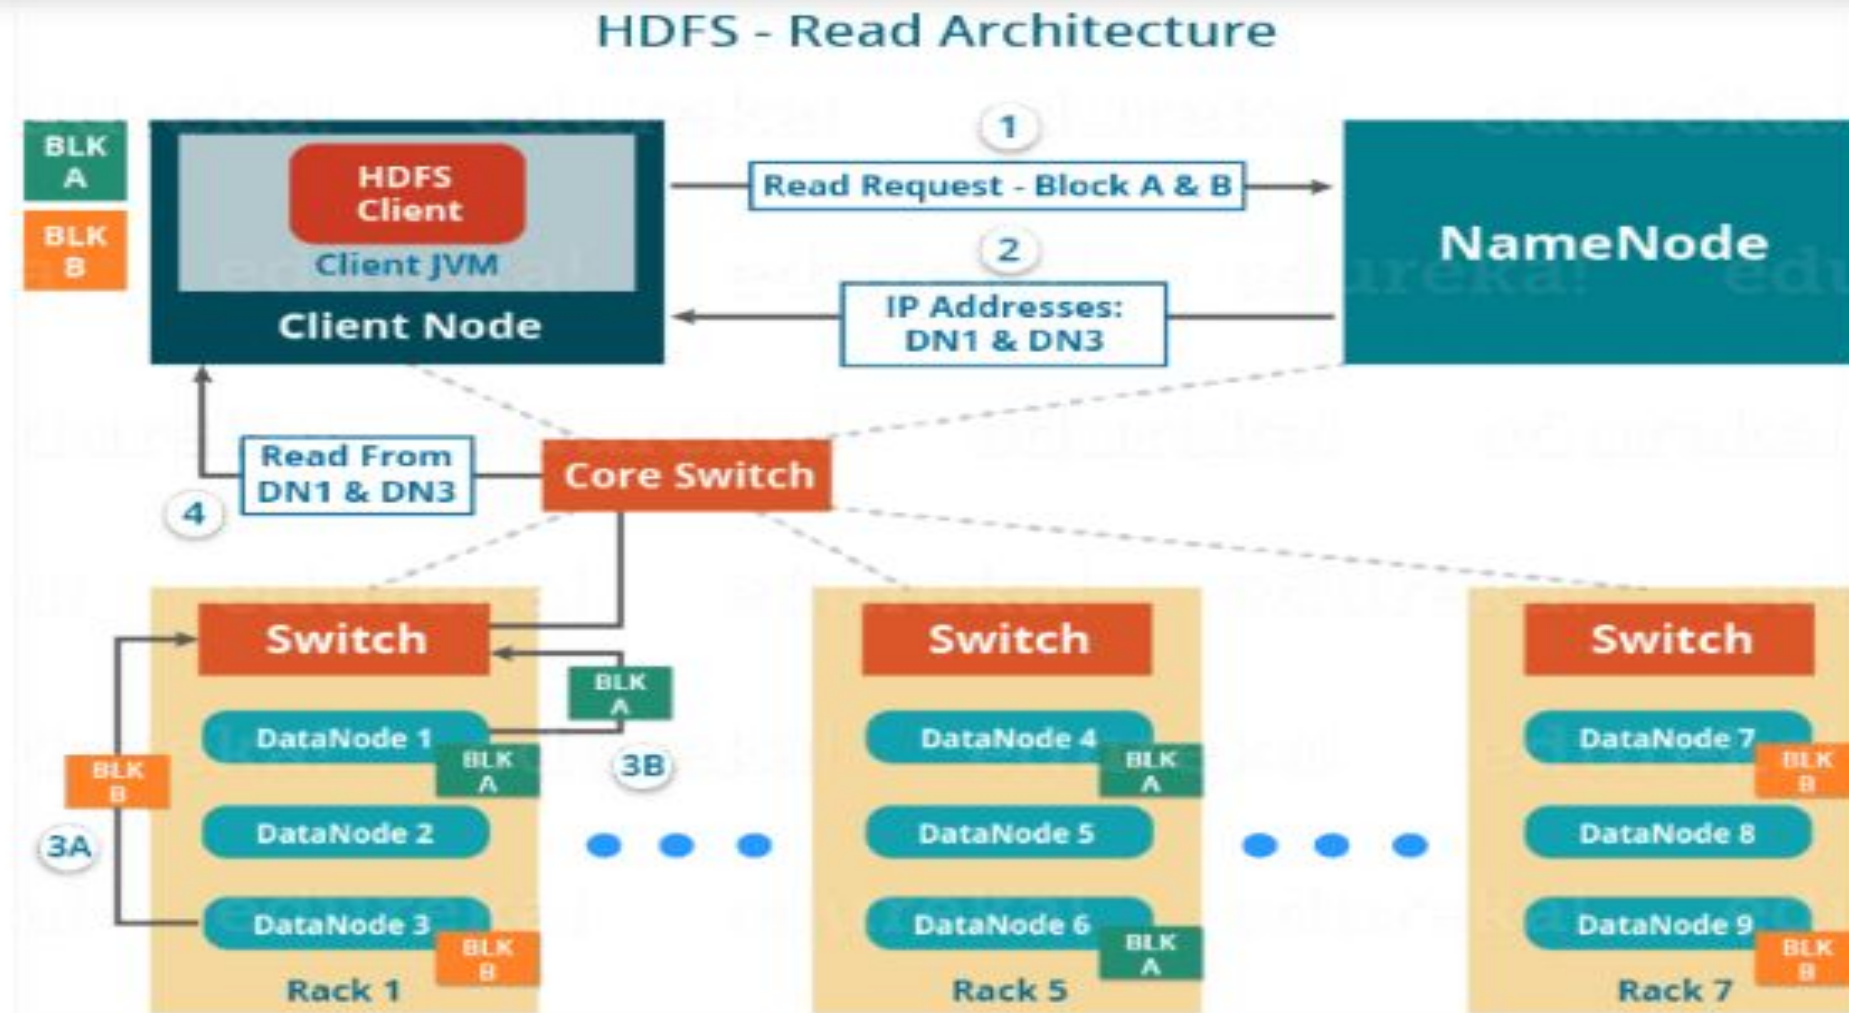
HDFS Write Mechanism – Acknowledgment



HDFS Multi-block Write Mechanism



HDFS Read (Anatomy of File Read)



Replica Selection

- To minimize global bandwidth consumption and read latency, HDFS tries to satisfy a read request from a replica that is **closest** to the reader.
- If there exists a replica on the **same rack** as the reader node, then that replica is preferred to satisfy the read request
- If HDFS cluster spans multiple data centers, then a replica that is resident in the **local data center** is preferred over any remote replica.

Assumptions and Goals of HDFS

- Hardware Failure
- Streaming Data Access
- Large Data Sets
- Simple Coherency Model

(HDFS applications need a write-once-read-many access model for files)

- “Moving Computation is Cheaper than Moving Data”
- Portability Across Heterogeneous Hardware and Software Platforms

The Communication Protocols

- All HDFS communication protocols are layered on top of the **TCP/IP protocol**.
- A client establishes a connection to a configurable TCP port on the NameNode machine.
- It talks the ClientProtocol with the NameNode.
- The DataNodes talk to the NameNode using the DataNode Protocol. A Remote Procedure Call (RPC) abstraction wraps both the Client Protocol and the DataNode Protocol.
- By design, the NameNode never initiates any RPCs. Instead, it **only responds to RPC** requests issued by DataNodes or clients.

Types of Failures in HDFS(Providing Robustness)

1. **DataNode Failure**

Data Disk Failure, Heartbeats (Solution- Re-Replication)

Data Integrity - the Datanode arrived is corrupted(Solution- checksum)

2. **Network Failure**

Cluster Rebalancing-In the event of a sudden high demand for a particular file, a scheme might dynamically create additional replicas and rebalance other data in the cluster.

3. **Namenode Failure**

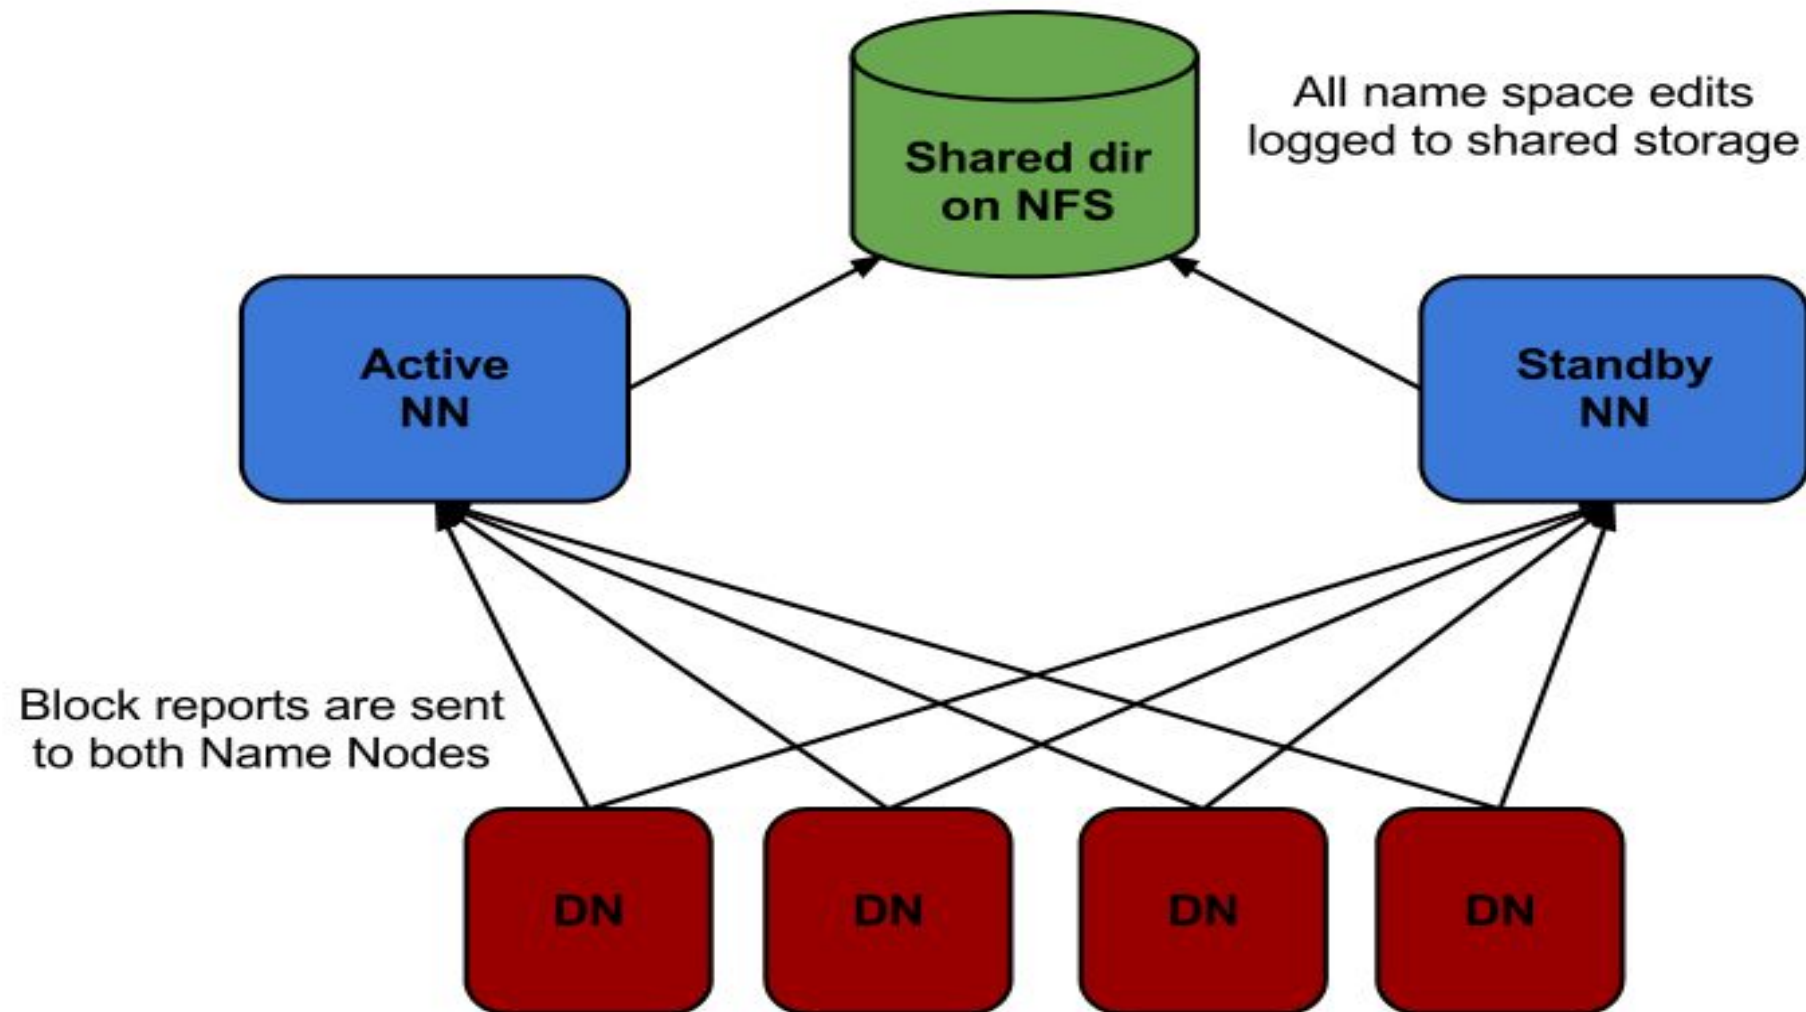
The FsImage and the EditLog are central data structures of HDFS. A corruption of these files can cause the HDFS instance to be non-functional.

Solution

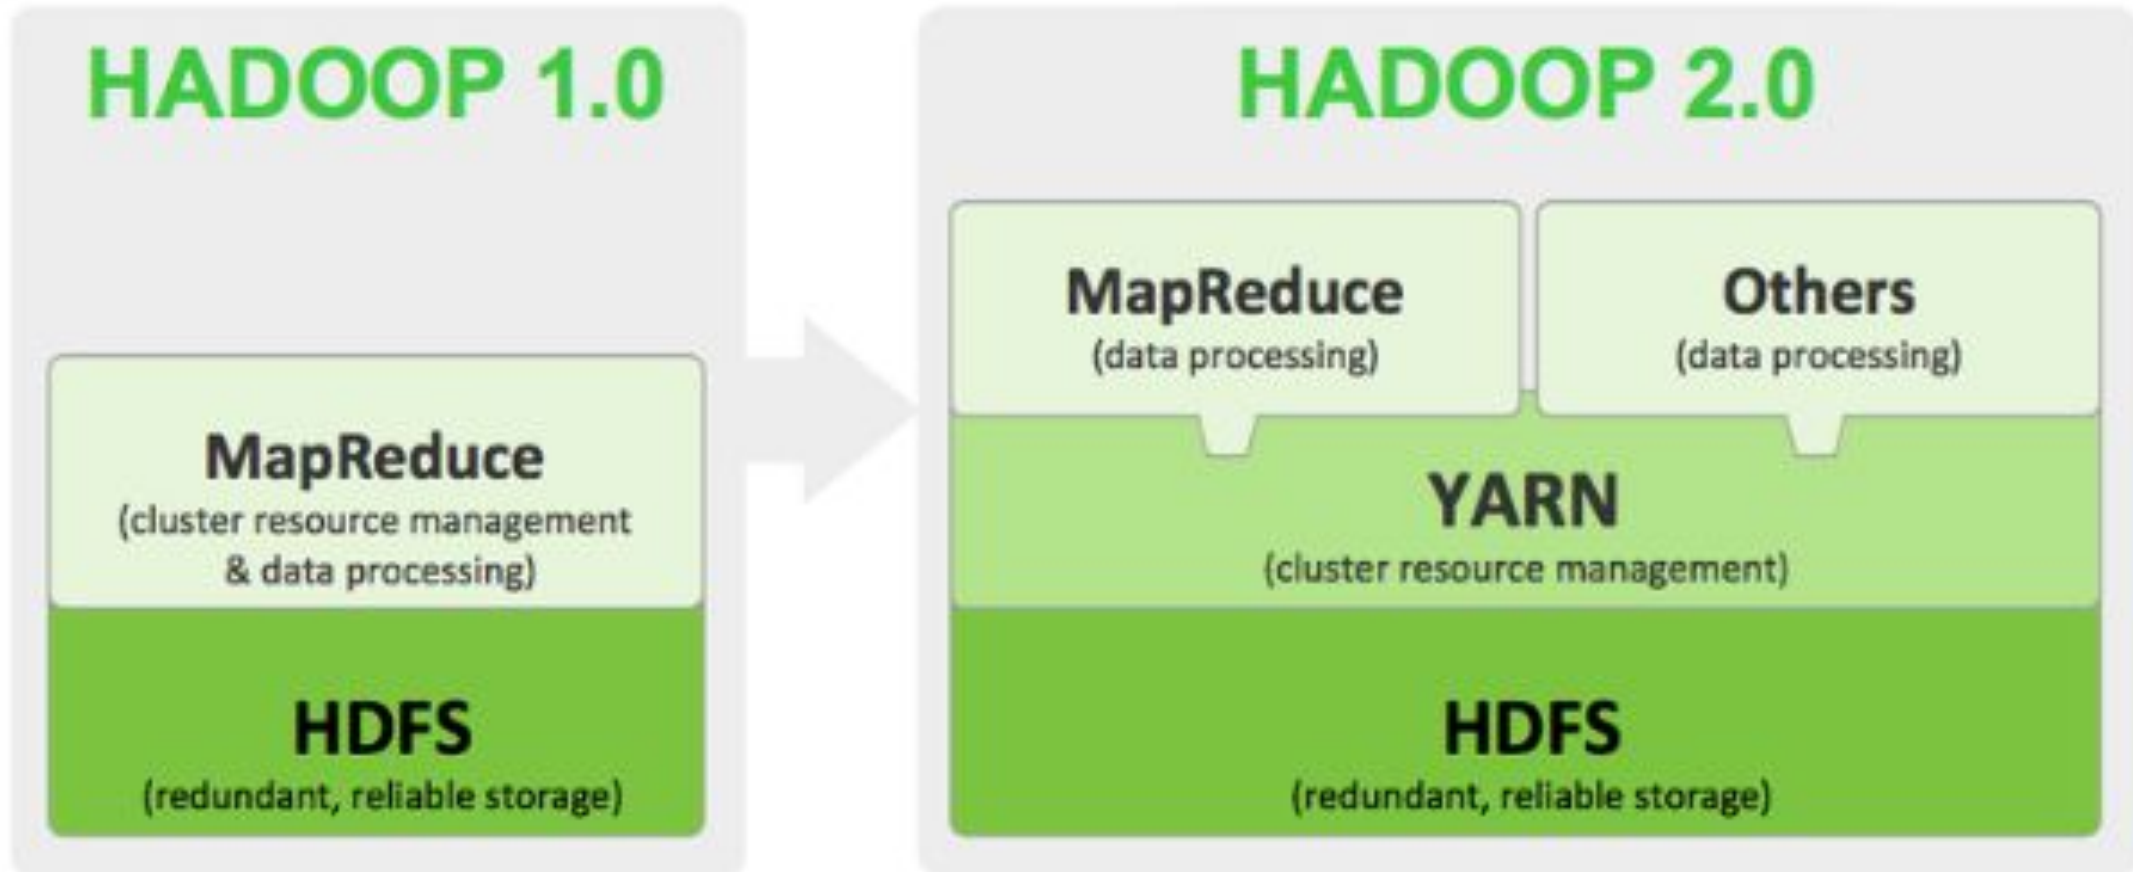
- To synchronous updating of multiple copies of the FsImage and EditLog
- To enable High Availability using multiple NameNodes

<https://hadoop.apache.org/docs/r2.10.2/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>

Multiple Name node(Hadoop 2.x)




Hadoop Version



Erasure Coding(Hadoop 3.x)

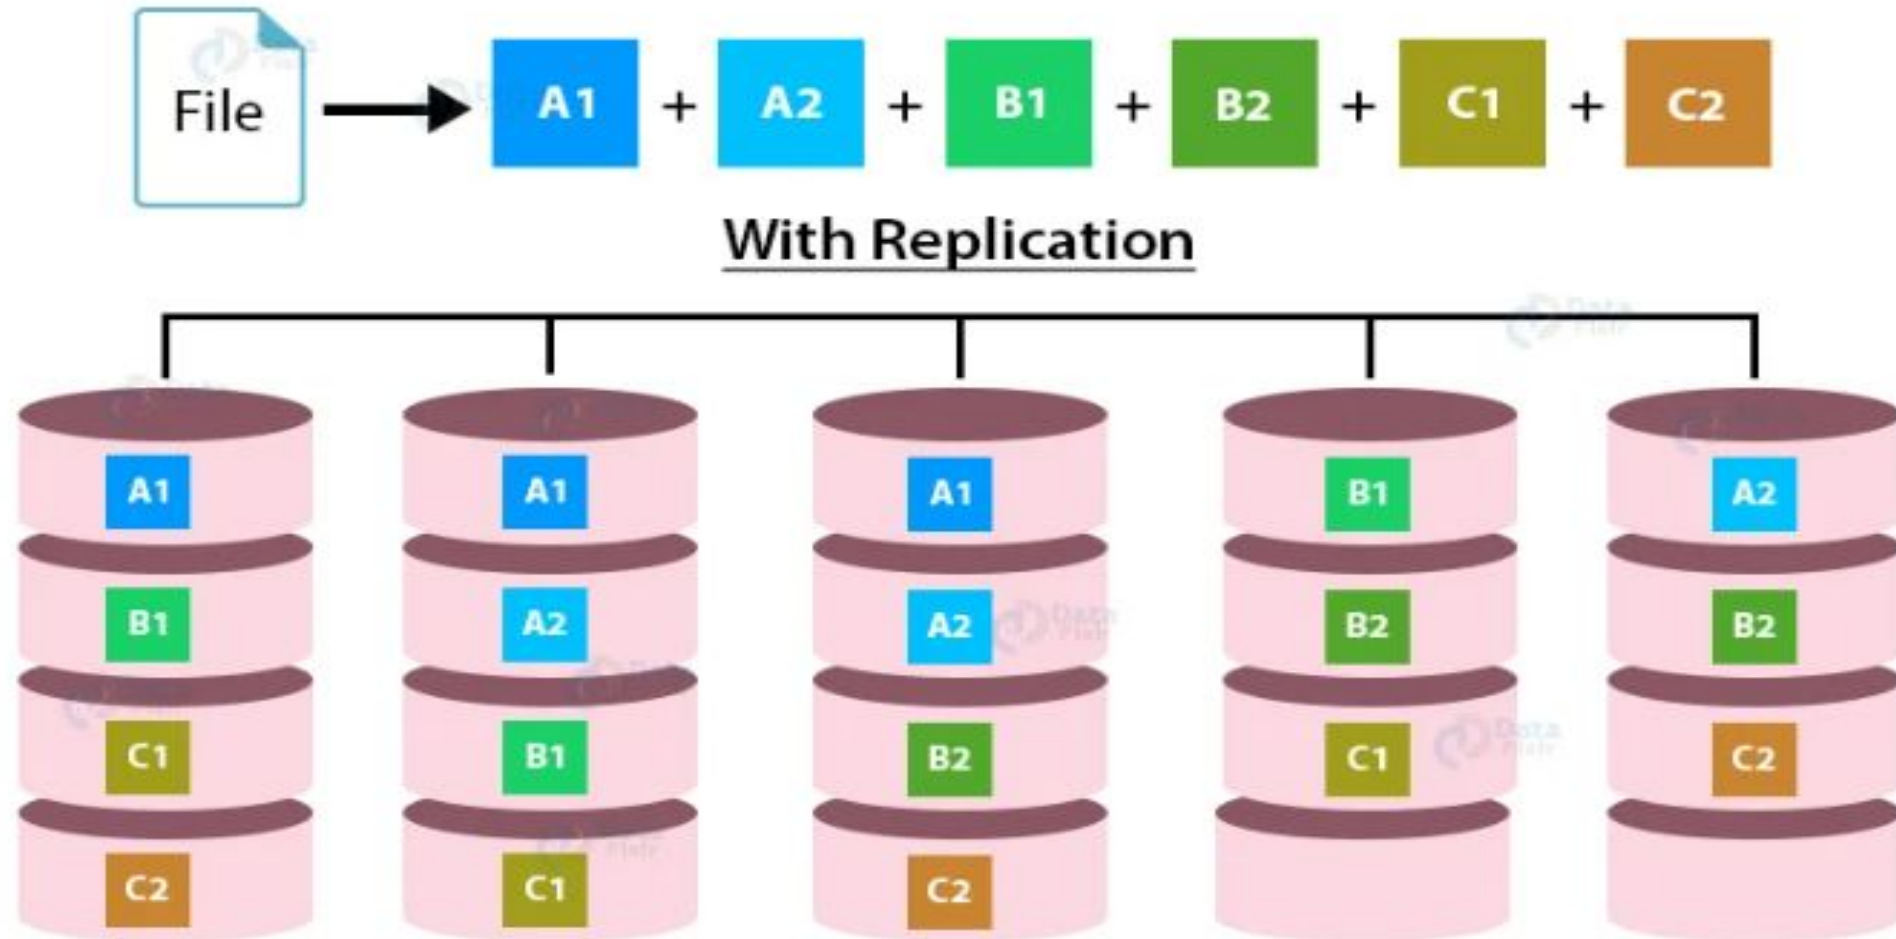
Block Replication



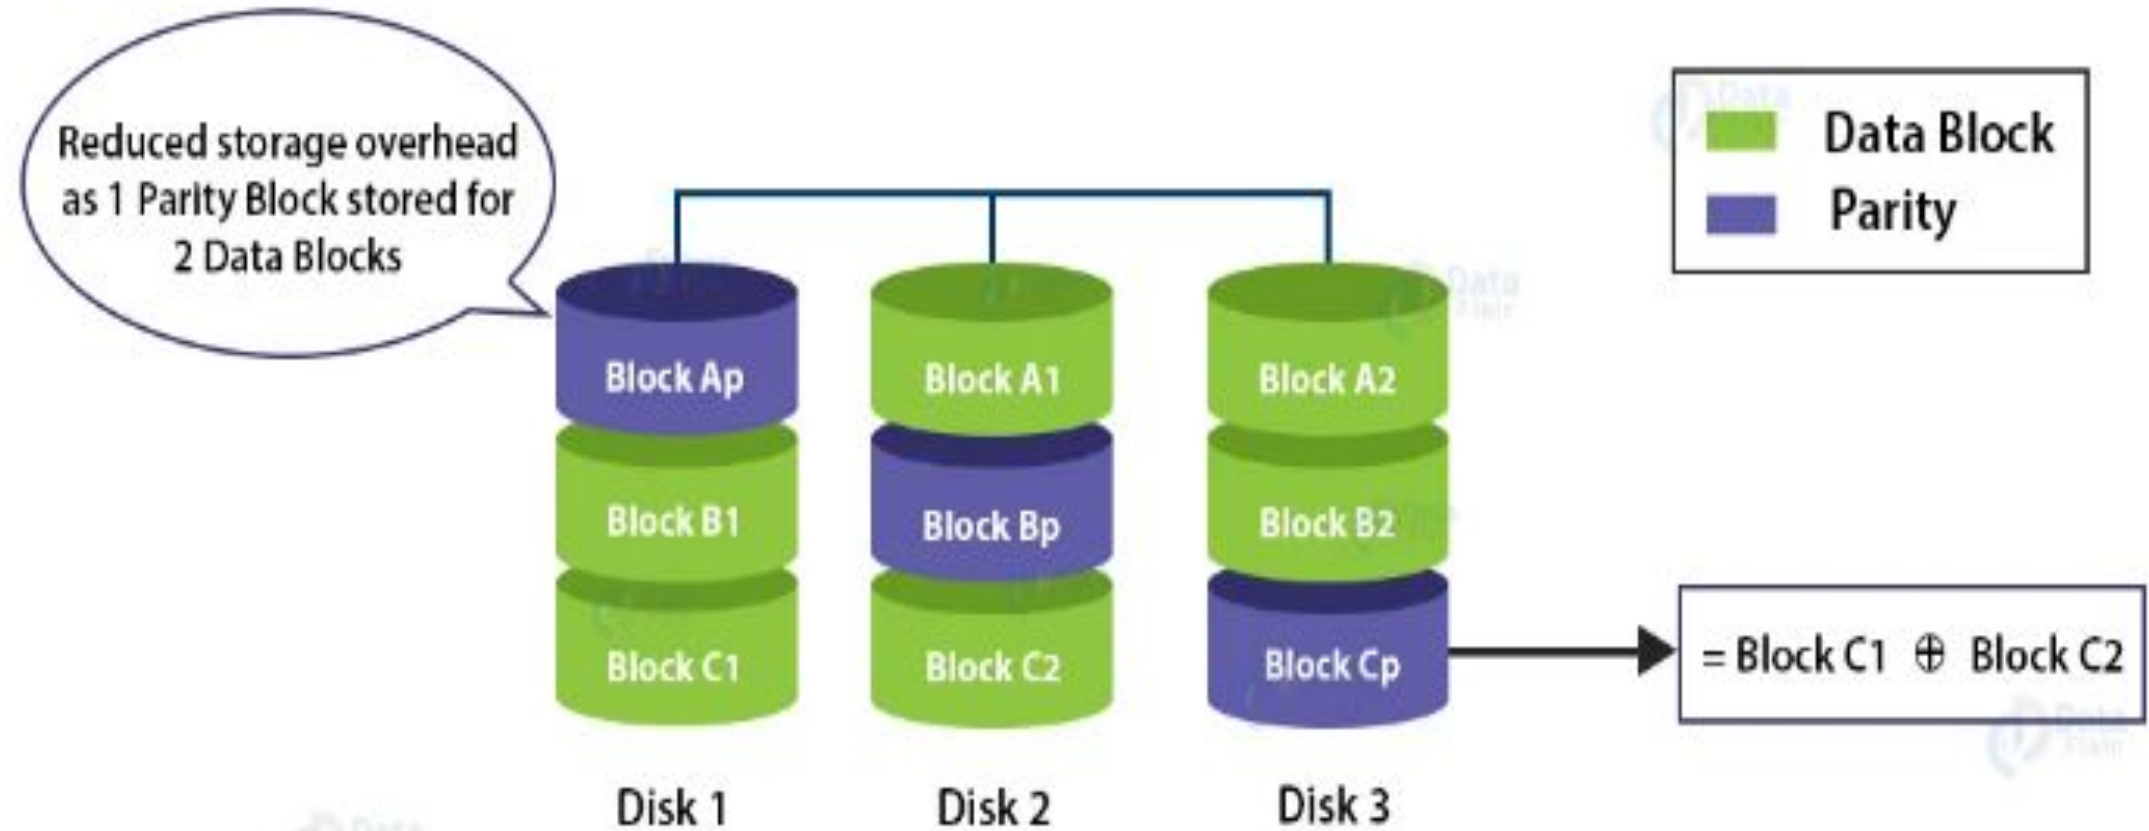
Space Used = $128 \times 3 = 384$ Mb

Therefore, 200% Storage Overhead

6 blocks with replication factor 3= 18 blocks



Erasure Coding



Erasure Coding(Cont.)

- In storage systems, the **Redundant Array of Inexpensive Disks (RAID)** uses Erasure Coding.
- RAID implements Erasure Coding by **striping**, that is, dividing logically sequential data such as file into smaller units (bit, byte, or block) and storing consecutive units on different disks.
- For each strip of the original dataset, a certain number of parity cells are calculated based on the **Erasure Coding algorithm** and stored, the process known as **encoding**.
- The error in any striping cell can be recovered from the calculation based on the remaining data and parity cells; the process known as **decoding**.
- Thus using Erasure Coding in HDFS improves storage efficiency while providing the same level of fault tolerance and data durability as traditional replication-based HDFS deployment.

Difference between hadoop 1.x,2.x,3.x

Key	Hadoop 1.x	Hadoop 2.x	Hadoop 3.x
Resource Management	Map reduce	YARN	YARN
Minimum supported Java version	JAVA 6	JAVA 7	JAVA 8
Windows OS Support	no	yes	yes
Fault Tolerance	Replication	Replication	Erasure coding
Storage Scheme	3x Replication	3x Replication	Erasure coding
Storage Overhead	200% of HDFS is consumed.	200% of HDFS is consumed.	50% of HDFS is consumed.
Scalability	4000 nodes in a cluster.	10000 nodes in a cluster.	more than 10000 nodes in a cluster.

References

- Book
 - Big Data and Analytics – Seema Acharya and Subhashini C – Wiley India
 - <https://hadoop.apache.org/docs/r2.10.2/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>