

Lab - 2

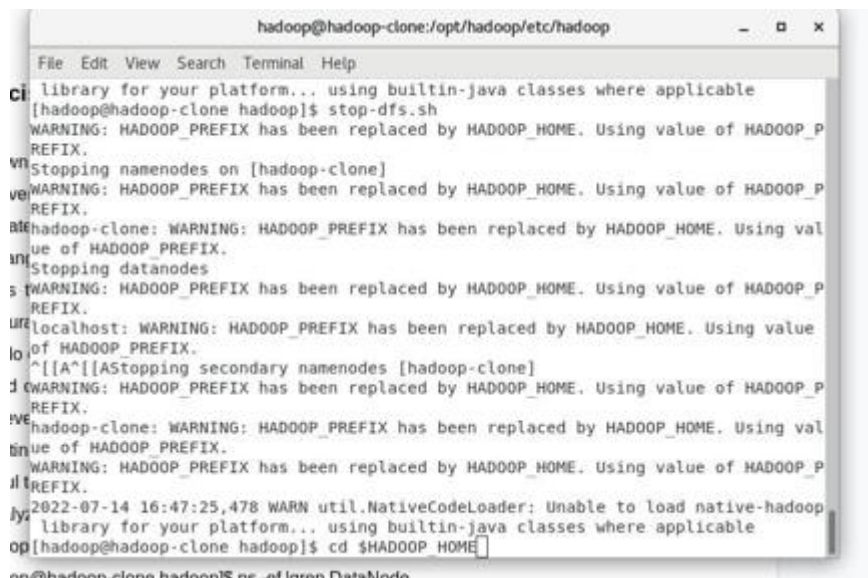
Aim: "Configuring and experiencing Hadoop Single node cluster. Start/stop script and monitoring."

Rahil Acharya

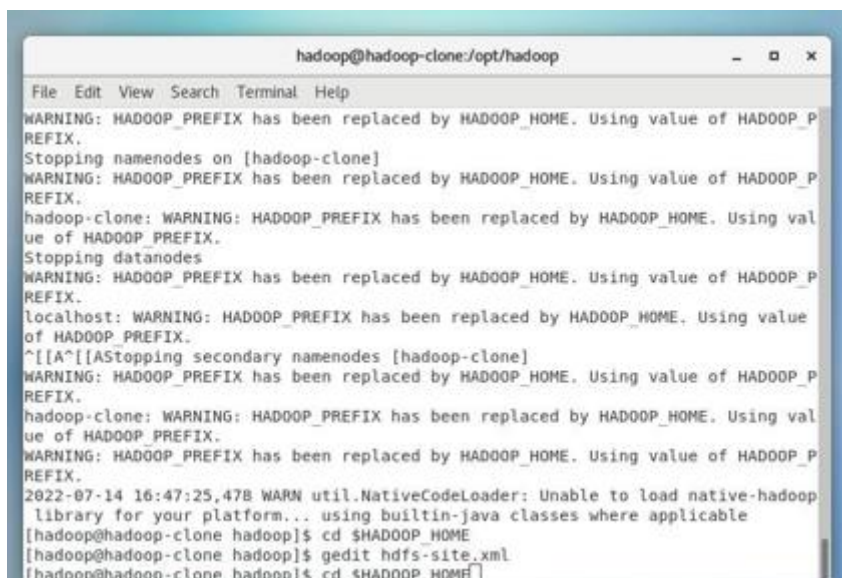
CE 004

20CEUOD004

Change the URL of HDFS Web UI to localhost:51234 instead of localhost:50070 and then access the web UI using new URL. [Hint: refer configuration files] [Know that traditionally configuration based servers need to be restarted after any server configuration change.]



```
hadoop@hadoop-clone:/opt/hadoop/etc/hadoop
File Edit View Search Terminal Help
ci library for your platform... using builtin-java classes where applicable
[hadoop@hadoop-clone hadoop]$ stop-dfs.sh
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
vn Stopping namenodes on [hadoop-clone]
ve WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
at hadoop-clone: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using val
ue of HADOOP_PREFIX.
in Stopping datanodes
s WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
ur localhost: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value
of HADOOP_PREFIX.
lo ^[[A^[[AStopping secondary namenodes [hadoop-clone]
d WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
ve hadoop-clone: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using val
ue of HADOOP_PREFIX.
in WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
il 2022-07-14 16:47:25,478 WARN util.NativeCodeLoader: Unable to load native-hadoop
ly library for your platform... using builtin-java classes where applicable
op [hadoop@hadoop-clone hadoop]$ cd $HADOOP_HOME
nn@hadoop-clone hadoop$ nc -l 50070
```



```
hadoop@hadoop-clone:/opt/hadoop
File Edit View Search Terminal Help
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
Stopping namenodes on [hadoop-clone]
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
hadoop-clone: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using val
ue of HADOOP_PREFIX.
Stopping datanodes
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
localhost: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value
of HADOOP_PREFIX.
^[[A^[[AStopping secondary namenodes [hadoop-clone]
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
hadoop-clone: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using val
ue of HADOOP_PREFIX.
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_P
PREFIX.
2022-07-14 16:47:25,478 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
[hadoop@hadoop-clone hadoop]$ cd $HADOOP_HOME
[hadoop@hadoop-clone hadoop]$ gedit hdfs-site.xml
[hadoop@hadoop-clone hadoop]$ cd $HADOOP_HOME
```

Overview 'hadoop-clone:9000' (active)

Started:	Thu Jul 14 16:50:43 +0530 2022
Version:	3.2.0, re97acb3bd8f3befd27418996fa5d4b50bf2e17bf
Compiled:	Tue Jan 08 11:38:00 +0530 2019 by sunilg from branch-3.2.0
Cluster ID:	CID-a0711dac-8474-44f9-8368-d68c1ff4c838
Block Pool ID:	BP-1266388418-127.0.0.1-1560101340487

Summary

Security is off.

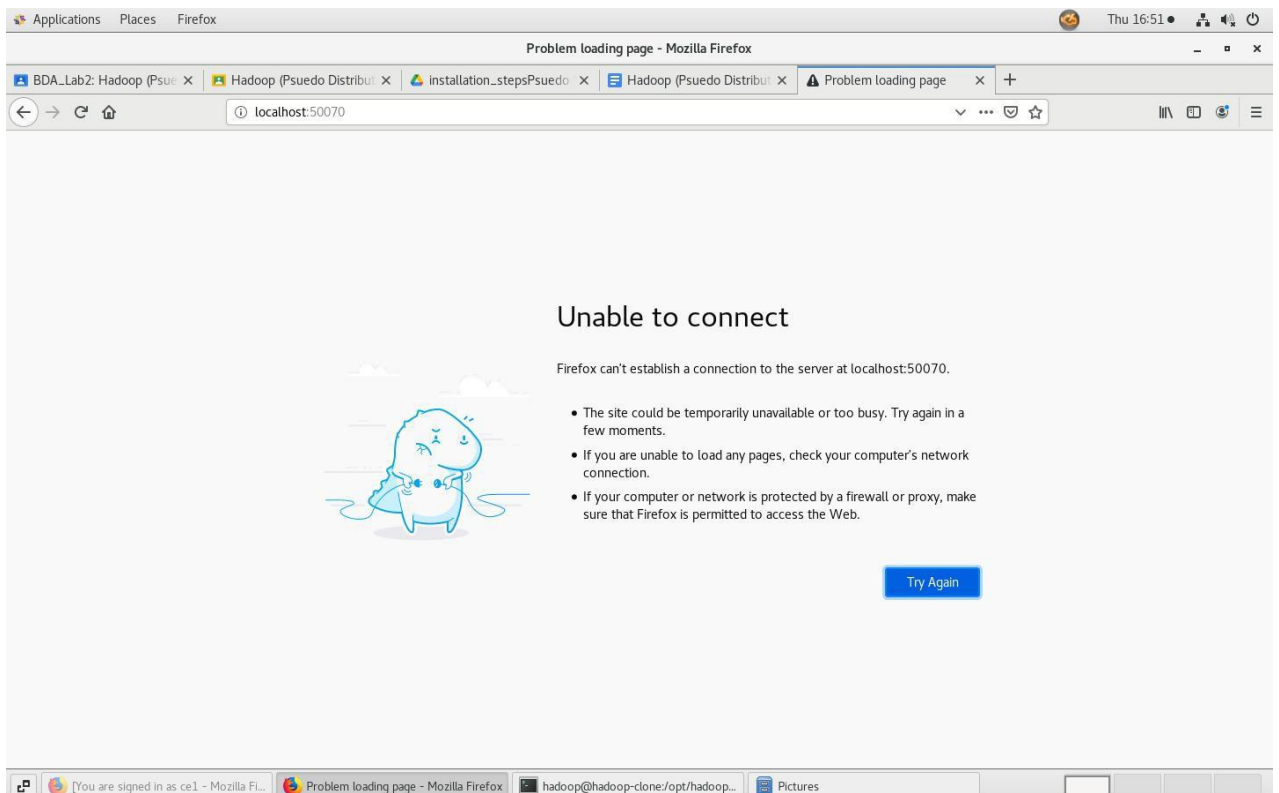
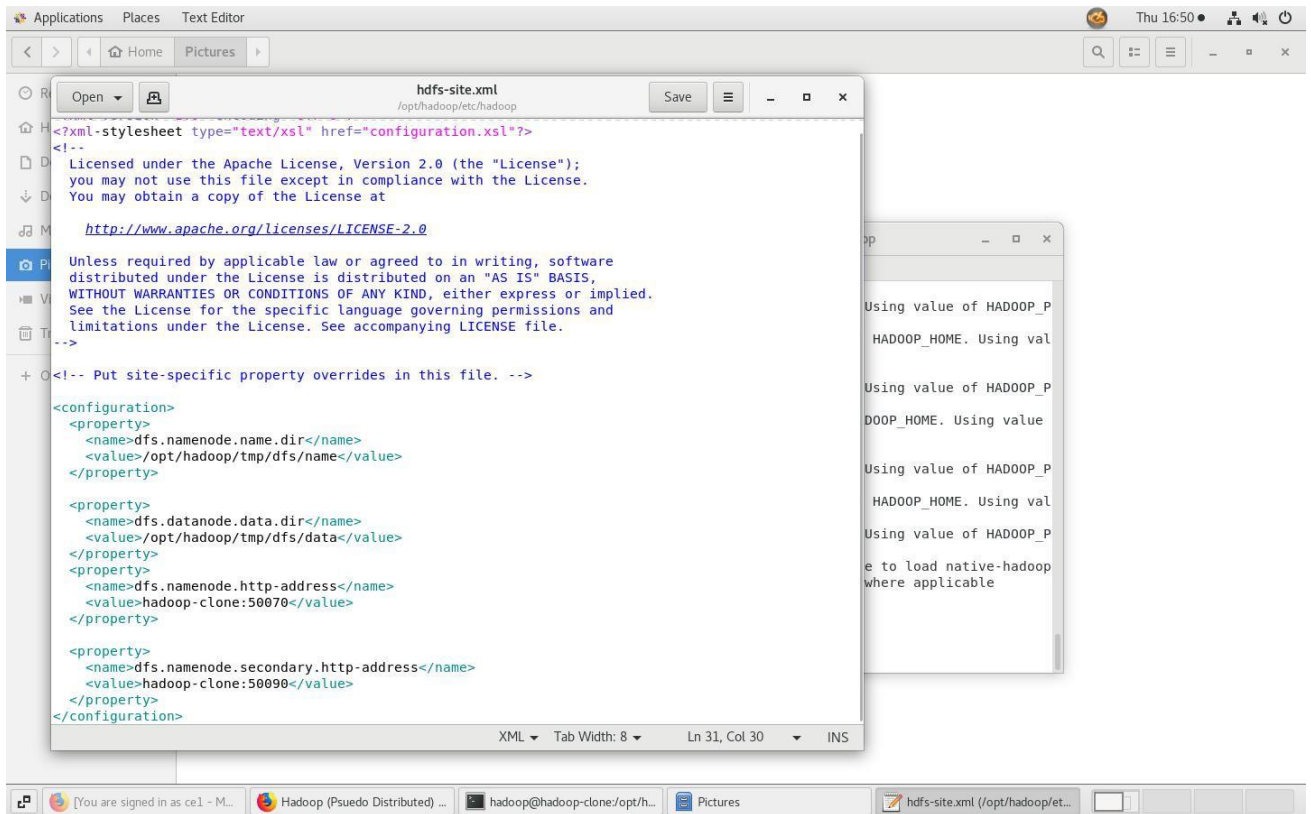
Safe mode is ON. The reported blocks 5 has reached the threshold 0.9990 of total blocks 5. The number of live datanodes 1 has reached the minimum number 0. In safe mode extension, Safe mode will be turned off automatically in 20 seconds.

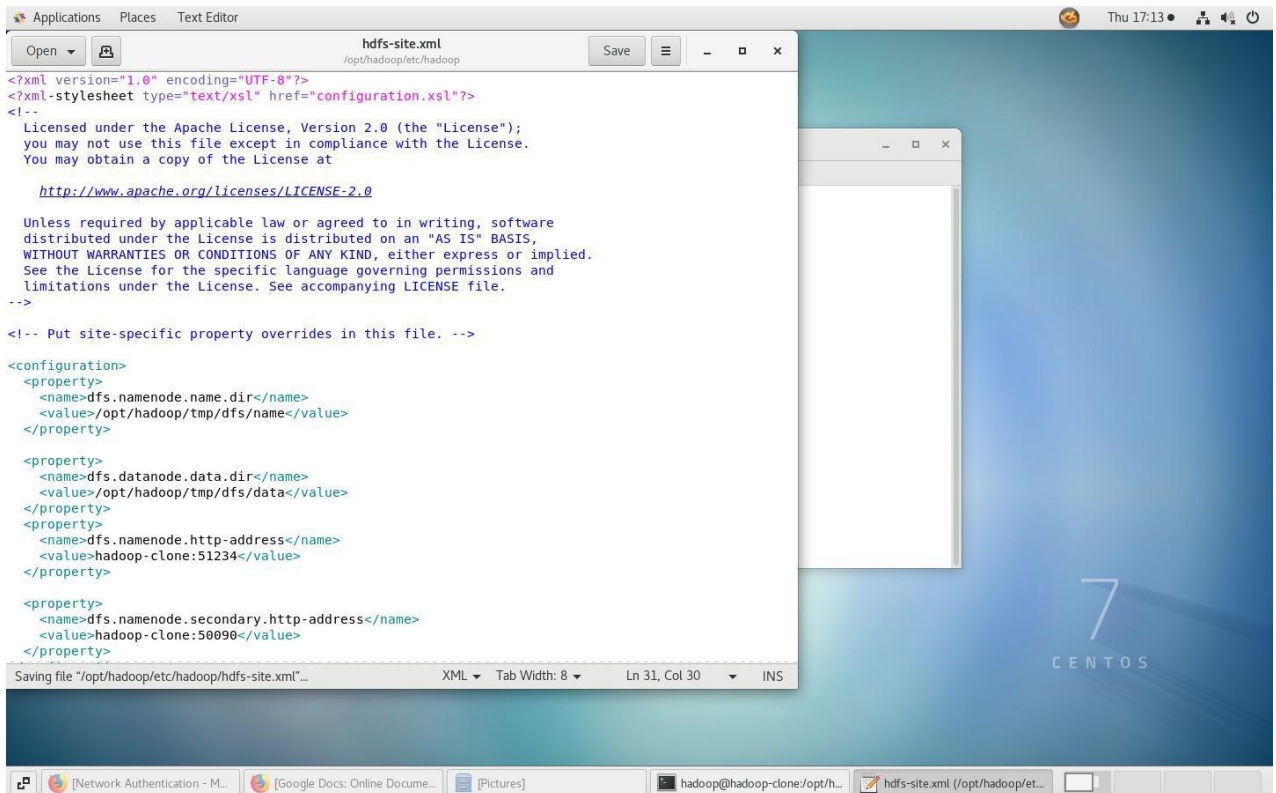
42 files and directories, 5 blocks (5 replicated blocks, 0 erasure coded block groups) = 47 total filesystem object(s).

Heap Memory used 84.03 MB of 423.5 MB Heap Memory. Max Heap Memory is 3.42 GB.

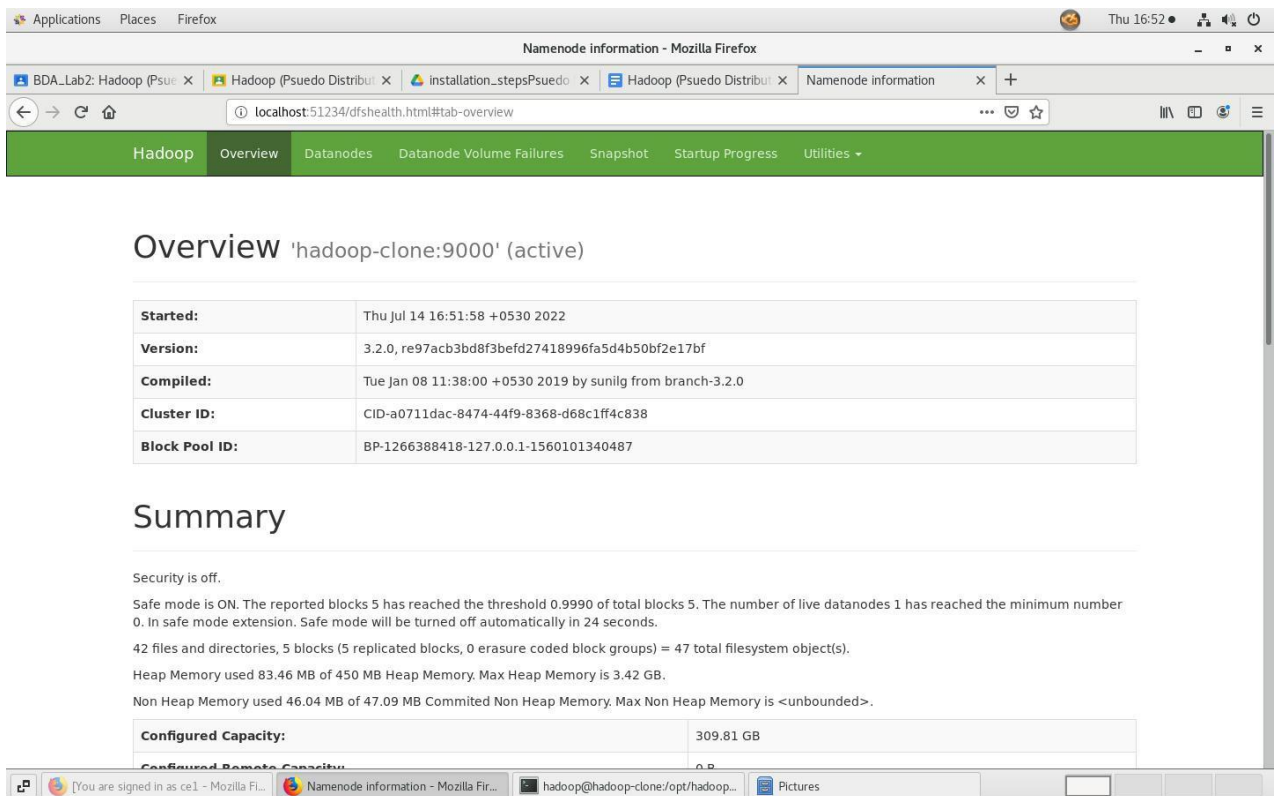
Non Heap Memory used 46.69 MB of 48.03 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	309.81 GB
----------------------	-----------

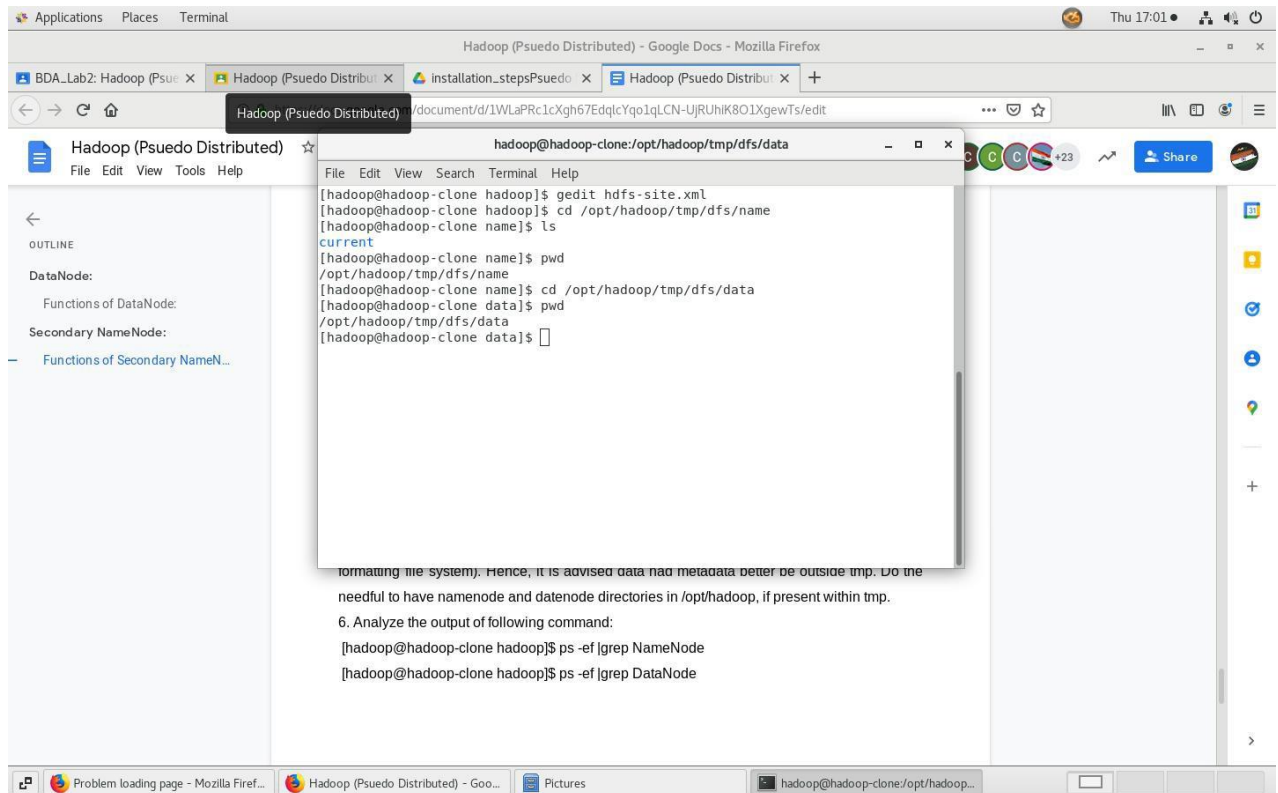




Undo changes done in task1.



Find out what is current location of namenode and datanode storedirectories. Normally whenever cleanup is required, tmp folder is one to be emptied up (not to confuse with formatting file system). Hence, it is advised data and metadata better be outside tmp. Do the needful to have namenode and datanode directories in /opt/hadoop, if present within tmp.



The screenshot shows a terminal window titled "hadoop@hadoop-clone:/opt/hadoop/tmp/dfs/data". The terminal output is as follows:

```
[hadoop@hadoop-clone hadoop]$ gedit hdfs-site.xml
[hadoop@hadoop-clone hadoop]$ cd /opt/hadoop/tmp/dfs/name
[hadoop@hadoop-clone name]$ ls
current
[hadoop@hadoop-clone name]$ pwd
/opt/hadoop/tmp/dfs/name
[hadoop@hadoop-clone name]$ cd /opt/hadoop/tmp/dfs/data
[hadoop@hadoop-clone data]$ pwd
/opt/hadoop/tmp/dfs/data
[hadoop@hadoop-clone data]$
```

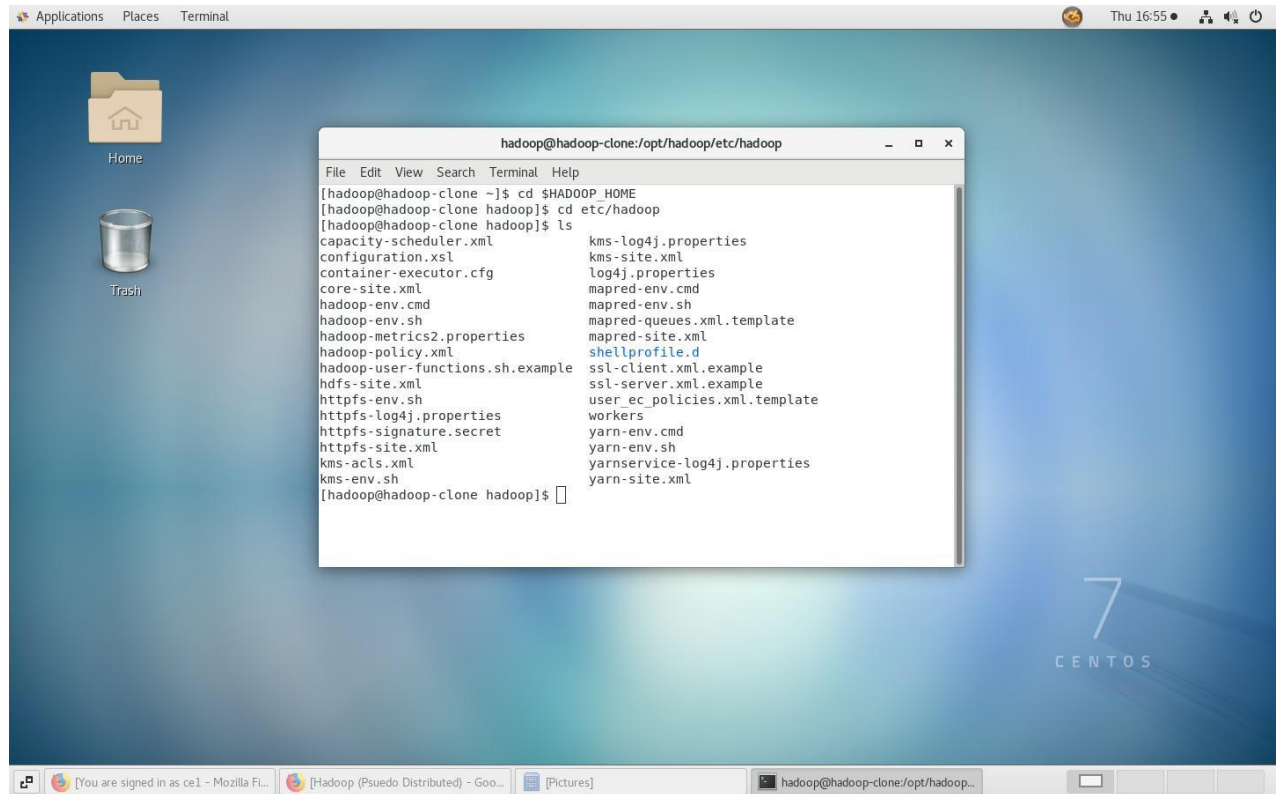
Below the terminal window, there is a text block that reads:

formatting file system. Hence, it is advised data and metadata better be outside tmp. Do the needful to have namenode and datanode directories in /opt/hadoop, if present within tmp.

6. Analyze the output of following command:

```
[hadoop@hadoop-clone hadoop]$ ps -ef |grep NameNode
[hadoop@hadoop-clone hadoop]$ ps -ef |grep DataNode
```

Analyze the output of following command: [hadoop@hadoop-clone hadoop]\$ ps -ef |grep NameNode[hadoop@hadoop-clone hadoop]\$ ps -ef |grep DataNode



The screenshot shows a CentOS 7 desktop environment. A terminal window is open, displaying the following commands and output:

```
hadoop@hadoop-clone:~/opt/hadoop/etc/hadoop
File Edit View Search Terminal Help
[hadoop@hadoop-clone ~]$ cd $HADOOP_HOME
[hadoop@hadoop-clone hadoop]$ cd etc/hadoop
[hadoop@hadoop-clone hadoop]$ ls
capacity-scheduler.xml      kms-log4j.properties
configuration.xml           kms-site.xml
container-executor.cfg      log4j.properties
core-site.xml               mapred-env.cmd
hadoop-env.cmd              mapred-env.sh
hadoop-env.sh               mapred-queues.xml.template
hadoop-metrics2.properties mapred-site.xml
hadoop-policy.xml           shellprofile.d
hadoop-user-functions.sh.example ssl-client.xml.example
hdfs-site.xml               ssl-server.xml.example
https-env.sh                user_ec_policies.xml.template
https-log4j.properties      workers
https-signature.secret      yarn-env.cmd
kms-acls.xml                yarn-env.sh
kms-env.sh                  yarnservice-log4j.properties
                             yarn-site.xml
[hadoop@hadoop-clone hadoop]$
```

The desktop background is blue with a large white number '7' and the word 'CENTOS' below it. The taskbar at the bottom shows several open applications: a terminal window, a file manager window, a web browser window, and a terminal window titled 'hadoop@hadoop-clone/opt/hadoop...'. The system clock in the top right corner shows 'Thu 16:55'.

Applications Places Terminal

Thu 16:53

Hadoop (Pseudo Distributed) Hadoop (Pseudo Distributed) installation_stepsPs

File Edit View Search Terminal Help

hadoop@hadoop-clone:/opt/hadoop/etc/hadoop

hadoop-clone: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
Stopping datanodes
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
localhost: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
Stopping secondary namenodes [hadoop-clone]
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
hadoop-clone: WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
2022-07-14 16:52:24,504 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[hadoop@hadoop-clone hadoop]\$ ps -ef |grep NameNode
hadoop 30117 11951 0 16:53 pts/0 00:00:00 grep --color=auto NameNode
[hadoop@hadoop-clone hadoop]\$ ps -ef |grep NameNode
hadoop 30143 11951 0 16:53 pts/0 00:00:00 grep --color=auto NameNode
[hadoop@hadoop-clone hadoop]\$ ps -ef |grep NameNode
hadoop 30256 11951 0 16:53 pts/0 00:00:00 grep --color=auto NameNode
[hadoop@hadoop-clone hadoop]\$

https://docs.google.com/document/d/1WLaPRC...

Hadoop (Pseudo Distributed)

File Edit View Tools Help

OUTLINE

DataNode:

Functions of DataNode:

Secondary NameNode:

Functions of Secondary NameNode:

Exercise:

1. Download and install hadoop using version specific navigation
2. Locate the log directory and refer to the log file
3. Change the URL of HDFS Web access the web UI using new configuration based servers need
4. Undo changes done in task1.
5. Find out what is current location of namenode and datanode storage directories. Normally whenever cleanup is required, tmp folder is one to be emptied up (not to confuse with formatting file system). Hence, it is advised data nad metadata better be outside tmp. Do the needful to have namenode and datenode directories in /opt/hadoop, if present within tmp.
6. Analyze the output of following command:
[hadoop@hadoop-clone hadoop]\$ ps -ef |grep NameNode
[hadoop@hadoop-clone hadoop]\$ ps -ef |grep DataNode